

Image Processing I

Computer Vision
Fall 2019
Columbia University

Homework 1, Problem 2b

- The TAs posted a hint on Piazza titled “Homework 1 Problem 2b”
- To maximize your partial credit on problems, show your work and justifications.

Homework 2

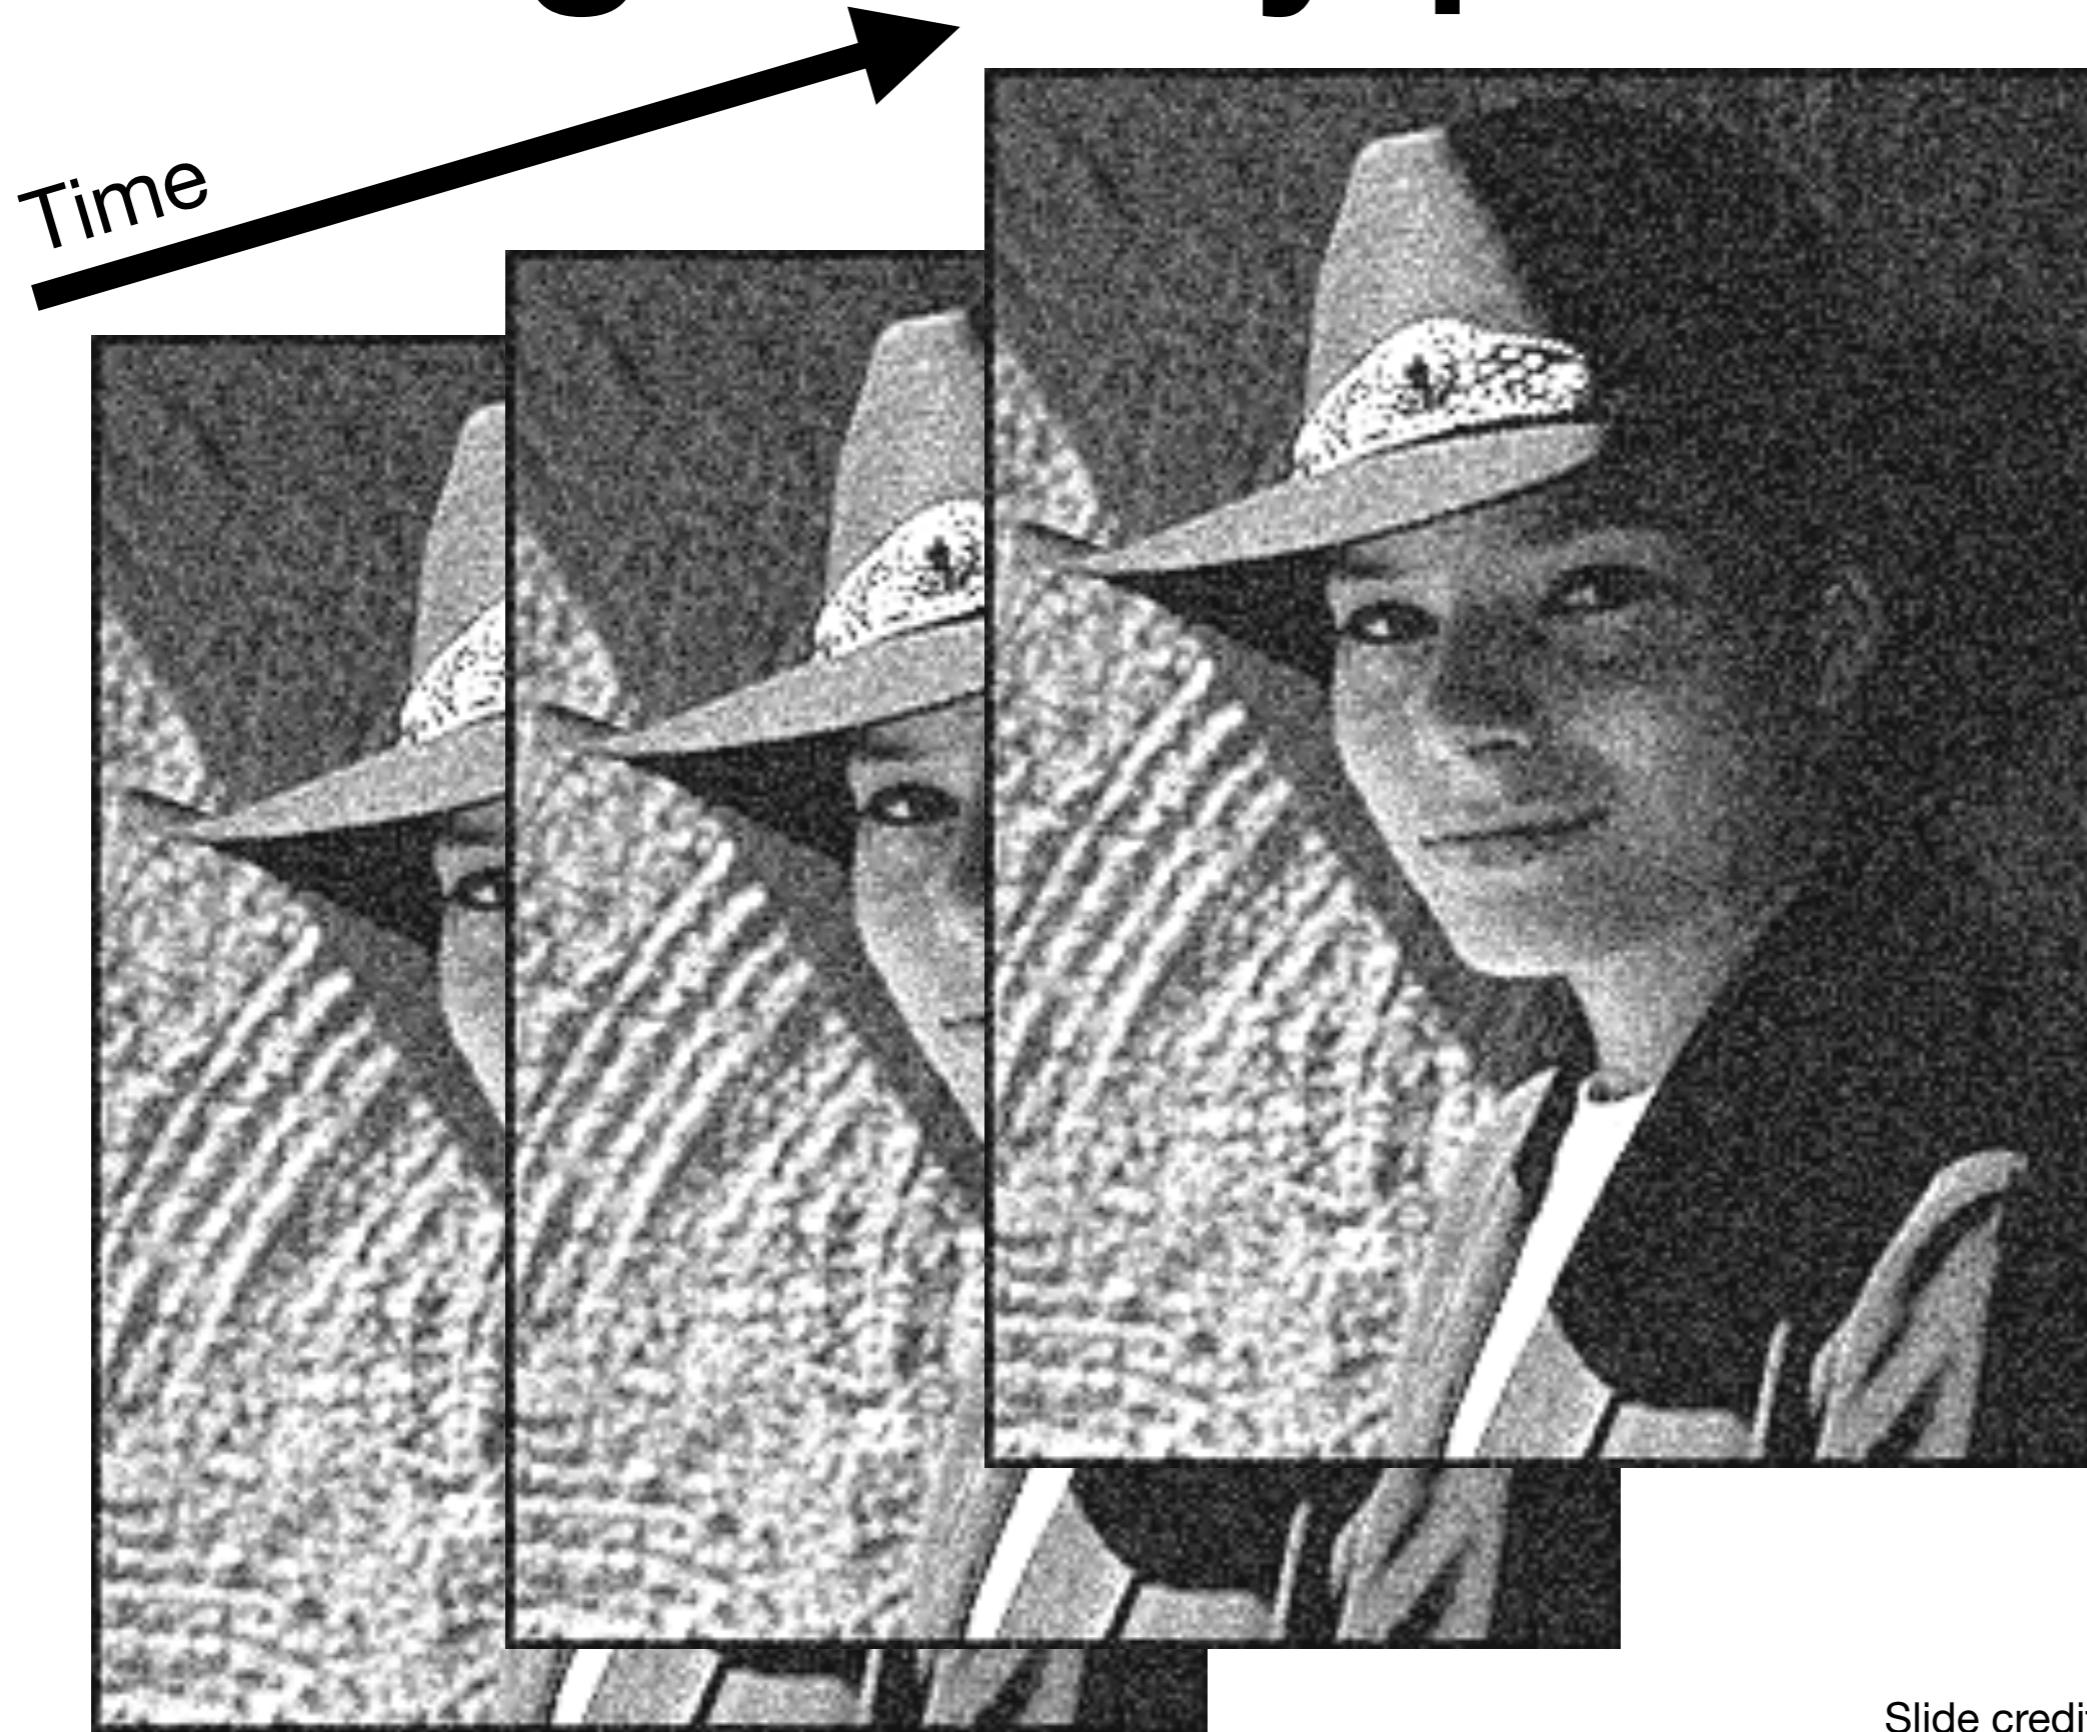
- Homework 2 will be released soon (if not already)
- Covers lecture today and Thursday

Image denoising



Slide credit: S. Lazebnik

Average many photos!



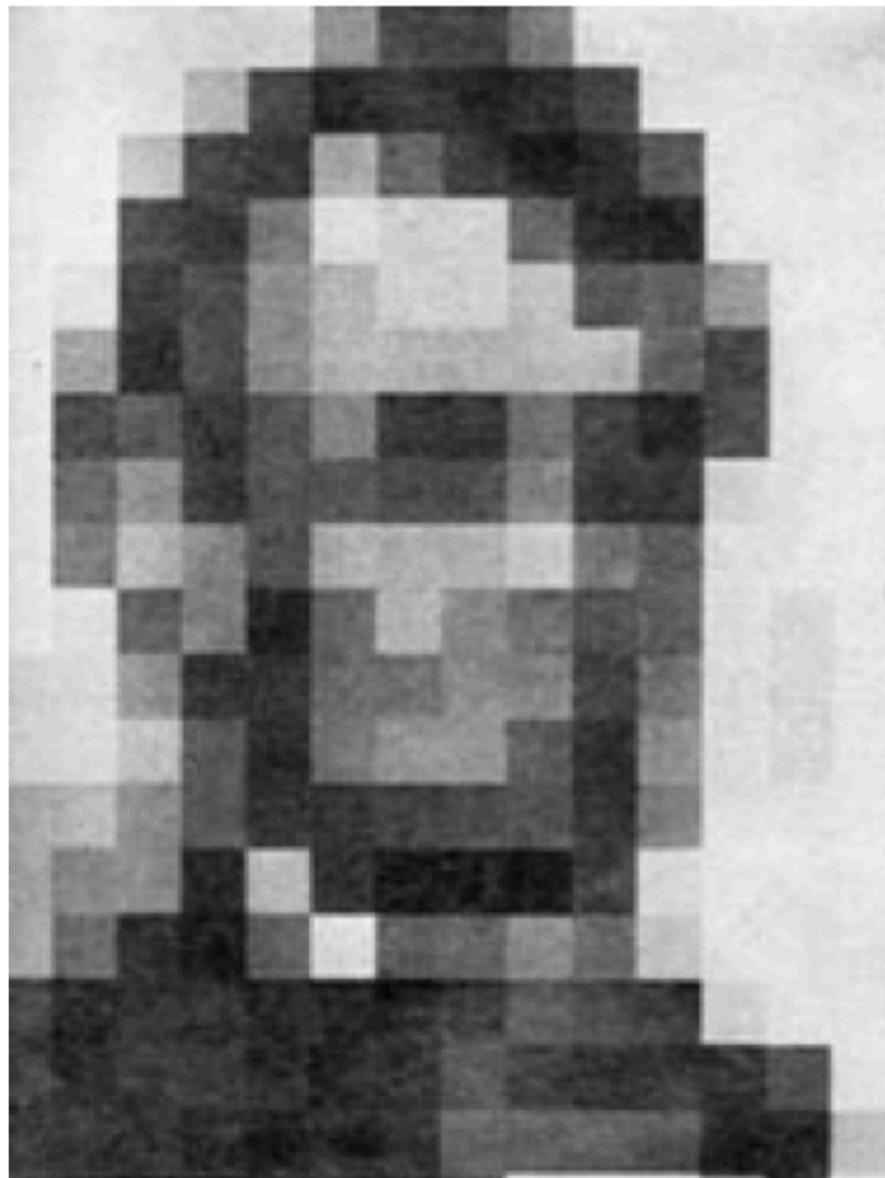
Slide credit: S. Lazebnik

What if just one?



Slide credit: S. Lazebnik

Reminder: Images as Functions



234	7	89	7	98	98	7	9	7	7	5
43	7	0	123	4	13	454	23	5	87	
67	5	76	4	3	56	67	87	65	45	
97	0	6	3	6	25	7	3	587	8	
78	5	54	7	876	71	54	76	9	75	
45	81	67	78	78	5	4	75	86	8	
5	4	3	35	8	256	6	4	3	36	
7	6	64	3	4	7	77	76	4	54	
64	35	46	46	64	56	7	56	4	7	
75	464	576	75	75	75	57	64	75	75	

$$F[x, y]$$

Moving Average

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

0									

Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

	0	10								

Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

0	10	20							

Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

0	10	20	30						

Moving Average

$$F[x, y]$$

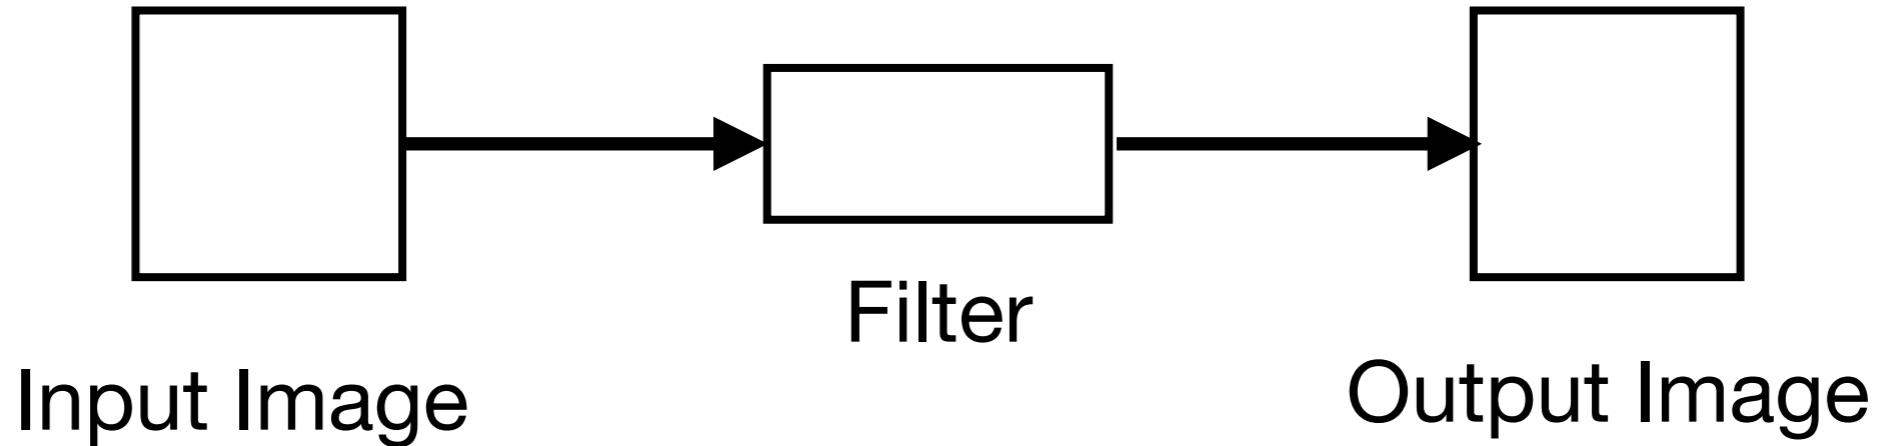
Moving Average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	
0	0	0	90	90	90	90	90	0	0	
0	0	0	90	90	90	90	90	0	0	
0	0	0	90	90	90	90	90	0	0	
0	0	0	90	0	90	90	90	0	0	
0	0	0	90	90	90	90	90	0	0	
0	0	0	0	0	0	0	0	0	0	
0	0	90	0	0	0	0	0	0	0	
0	0	0	0	0	0	0	0	0	0	

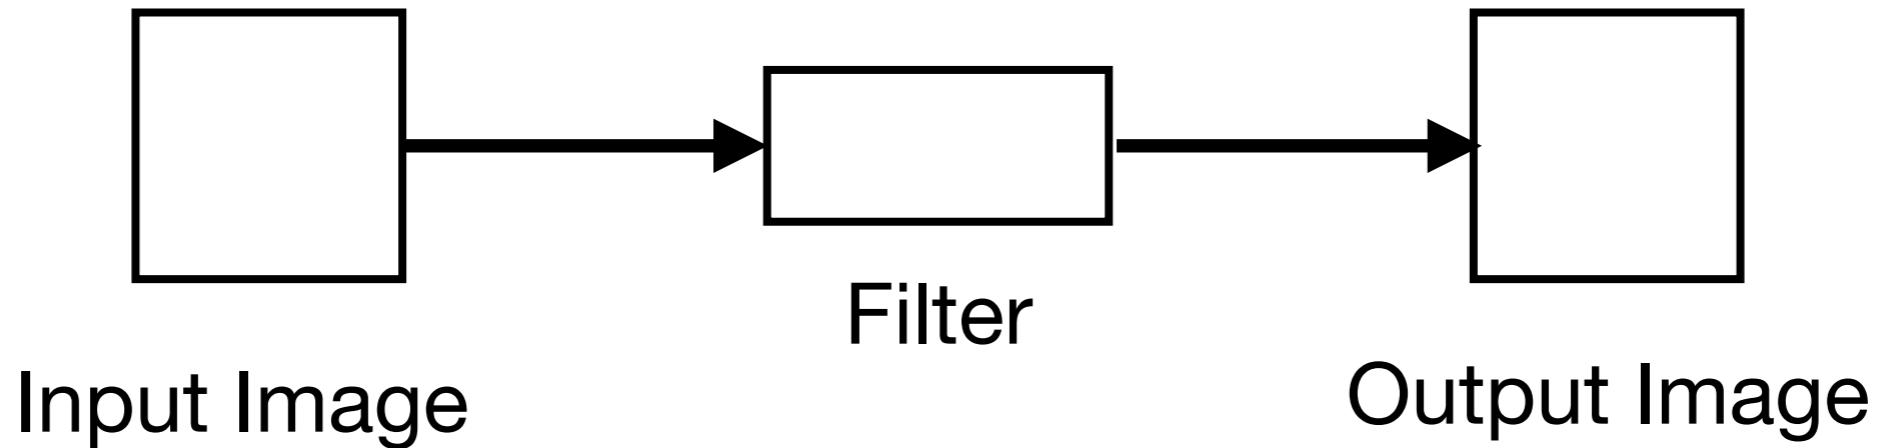
	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

Filtering



We want to remove unwanted sources of variation, and keep the information relevant for whatever task we need to solve

Linear Filtering



For a filter to be linear, it must satisfy two properties:

- $\text{filter}(\text{im}, \text{f1} + \text{f2}) = \text{filter}(\text{im}, \text{f1}) + \text{filter}(\text{im}, \text{f2})$
- $C * \text{filter}(\text{im}, \text{f1}) = \text{filter}(\text{im}, C * \text{f1})$

Convolution

$$F[x, y]$$

$$G[x, y]$$

*

$$\begin{array}{r} 1 \\ \hline 9 \end{array}$$

1

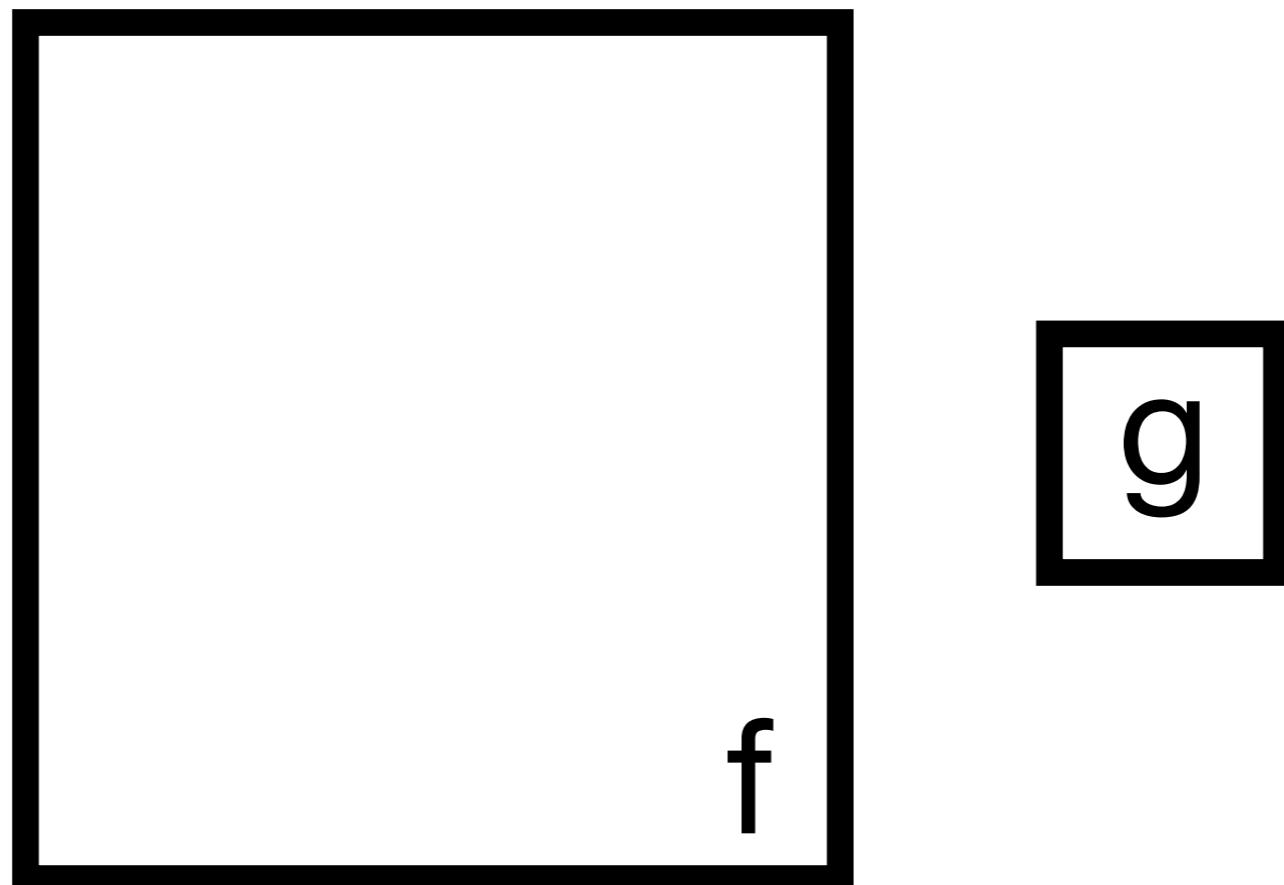
Convolution

- Let f be an image/function, and g is the kernel/filter
- The convolution is defined as:

$$(f * g)[x, y] = \sum_{i,j} f[x - i, y - j]g[i, j]$$

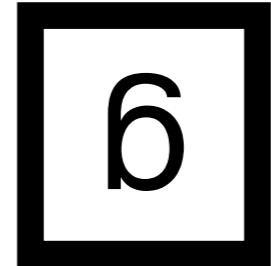
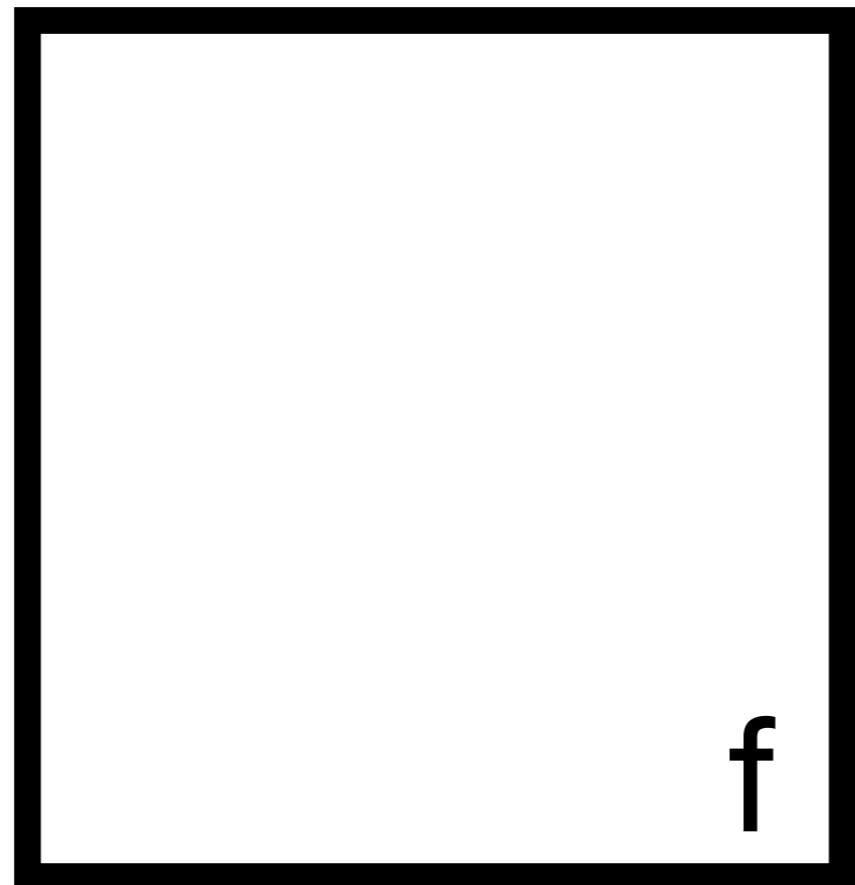
Convolution

$$(f * g)[x, y] = \sum_{i,j} f[x - i, y - j]g[i, j]$$



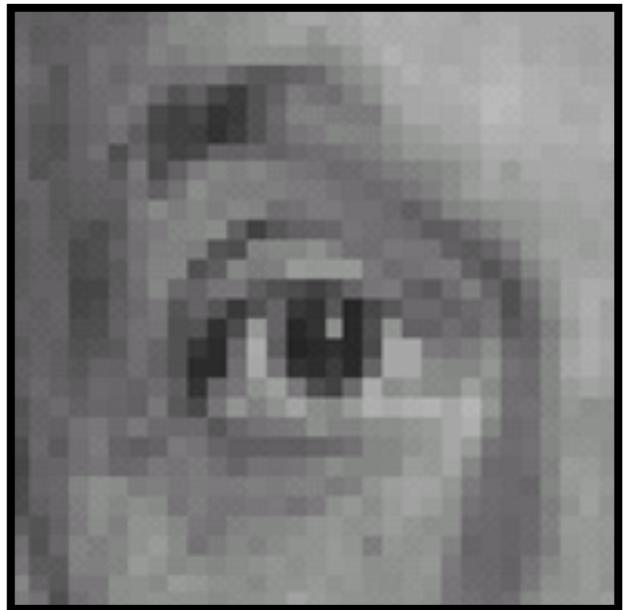
Convolution

$$(f * g)[x, y] = \sum_{i,j} f[x - i, y - j]g[i, j]$$



Flip LR, UD

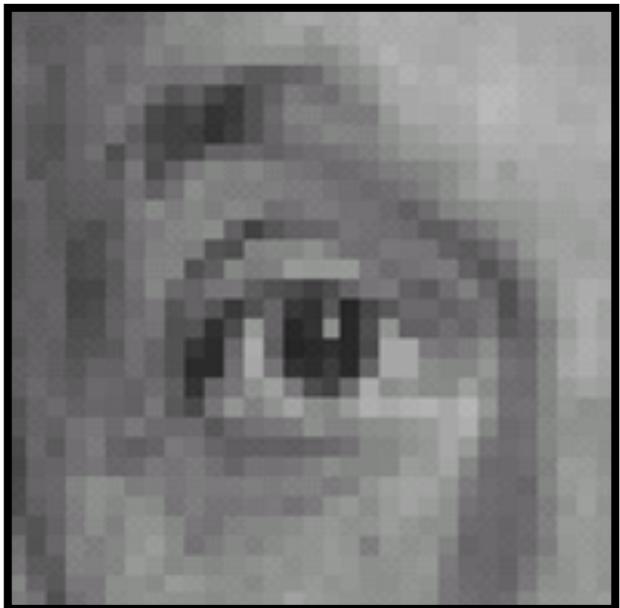
Convolution Practice



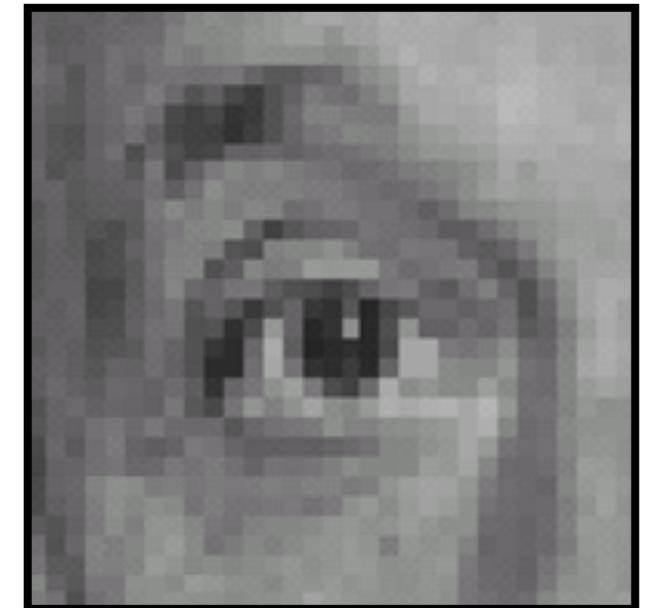
0	0	0
0	1	0
0	0	0

?

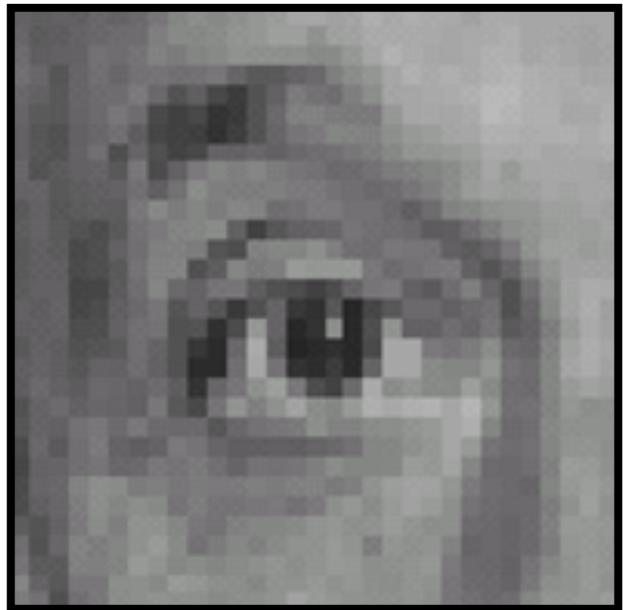
Convolution Practice



0	0	0
0	1	0
0	0	0



Convolution Practice

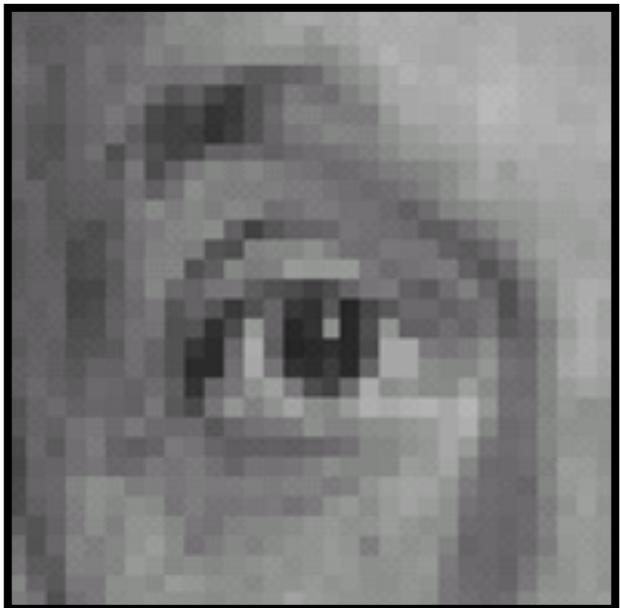


0	0	0
0	0	1
0	0	0

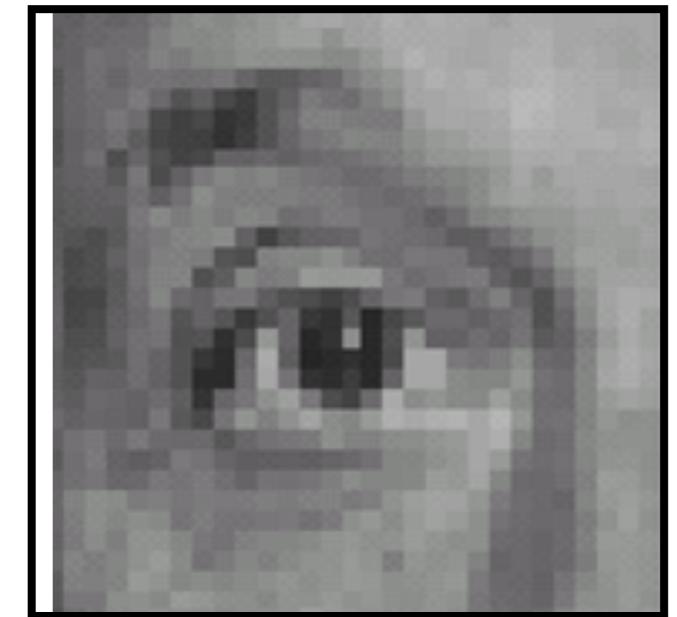
?

Convolution Practice

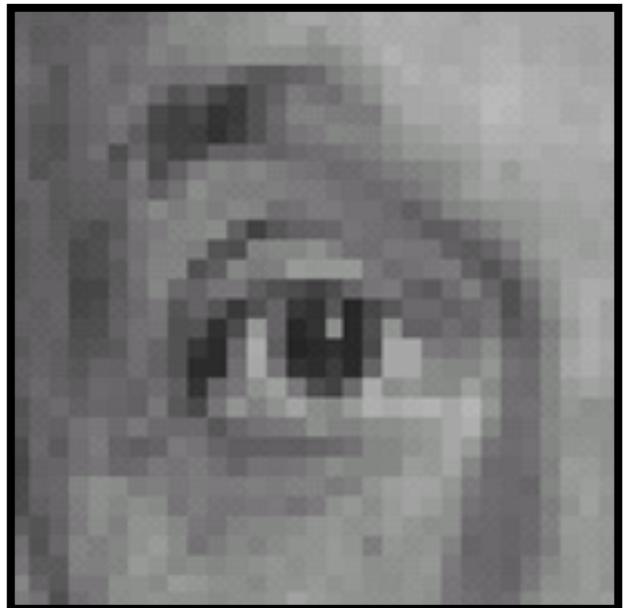
Translation Filter



0	0	0
0	0	1
0	0	0



Convolution Practice

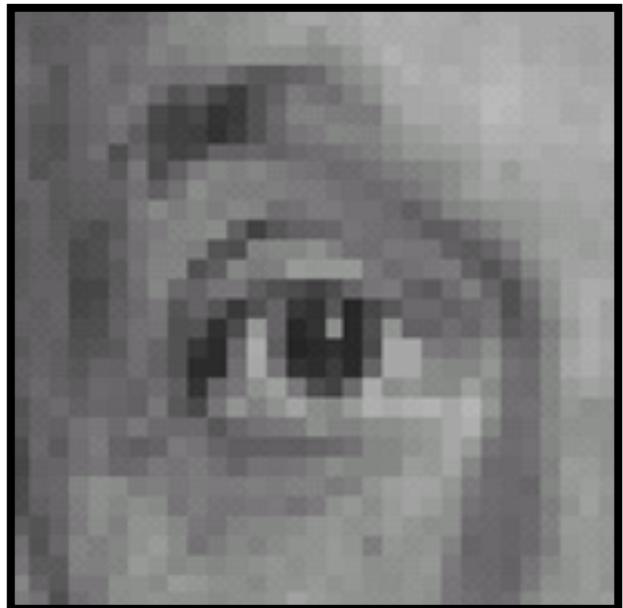


$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

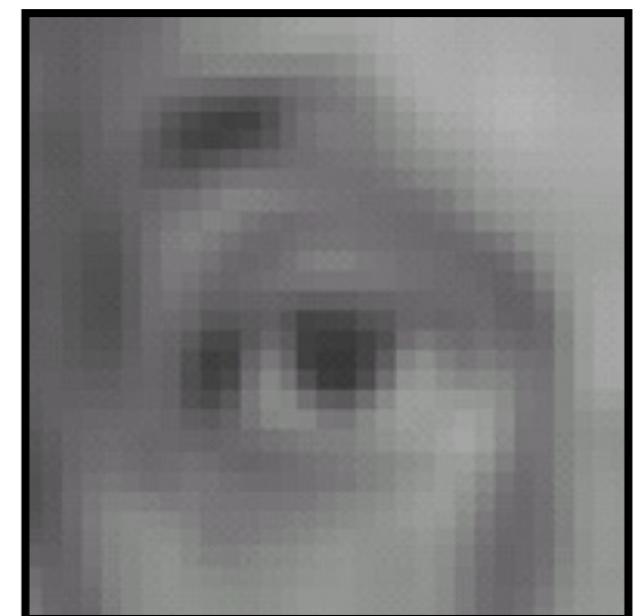
?

Convolution Practice

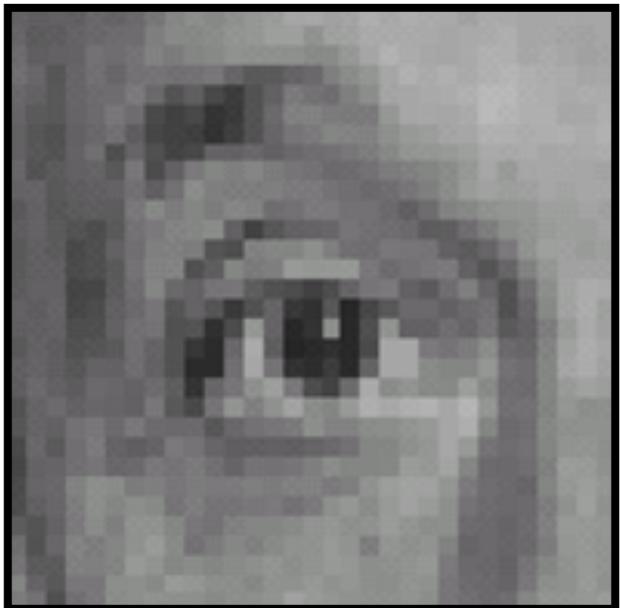
Blur Filter



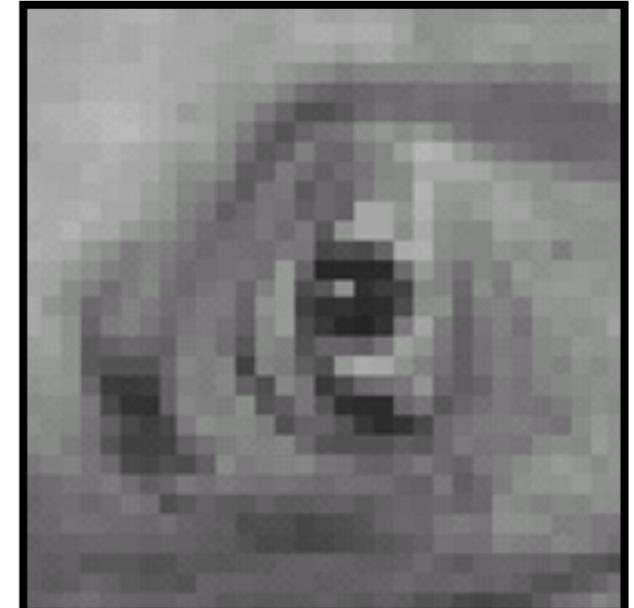
$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



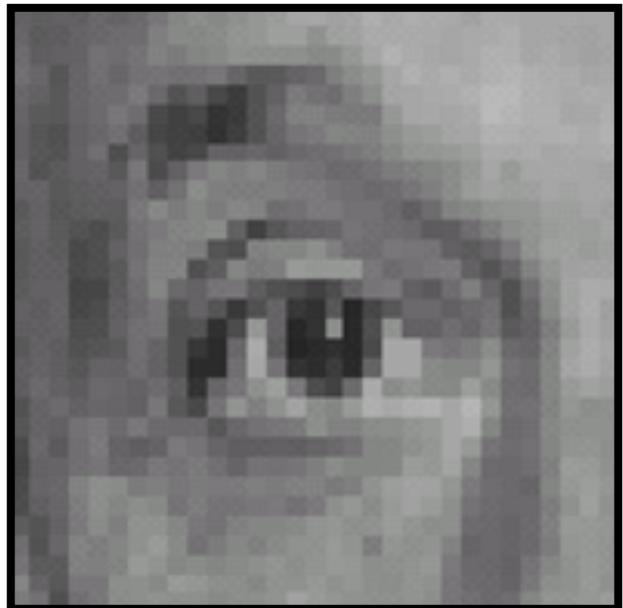
Convolution Practice



?



Convolution Practice



0	0	0
0	2	0
0	0	0

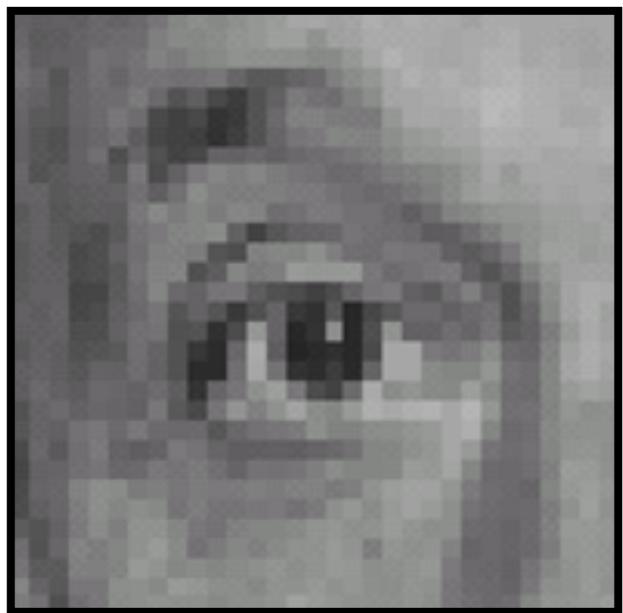
$$-\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

?

Convolution Practice

Sharpening Filter



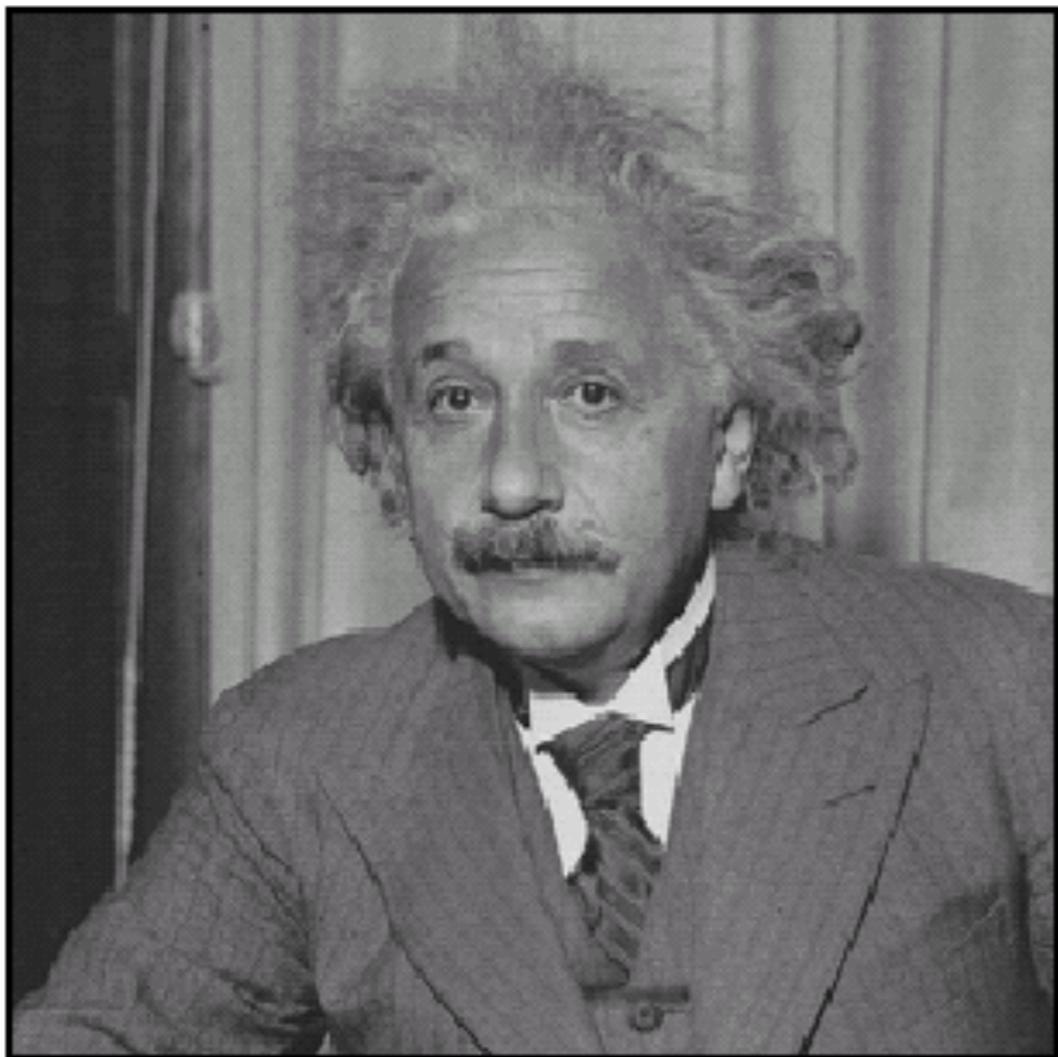
$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

$$-\frac{1}{9}$$

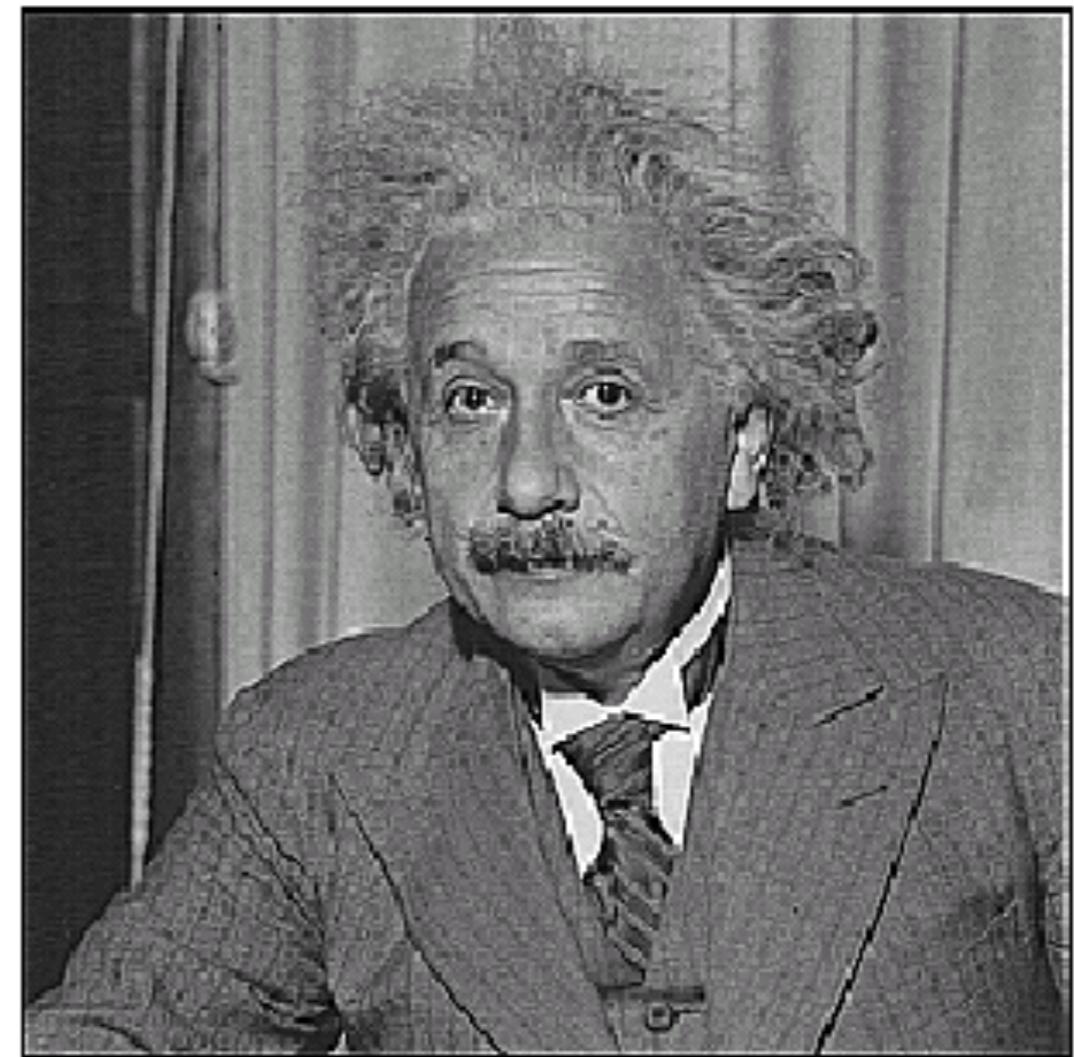
$$\begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



Sharpening

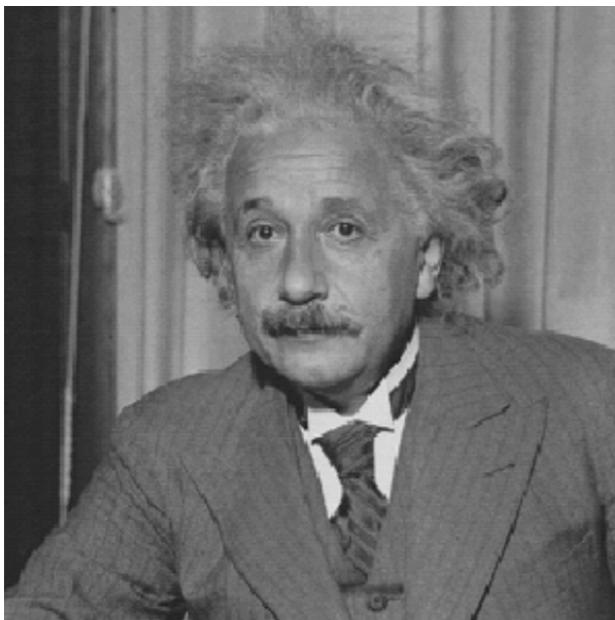


before

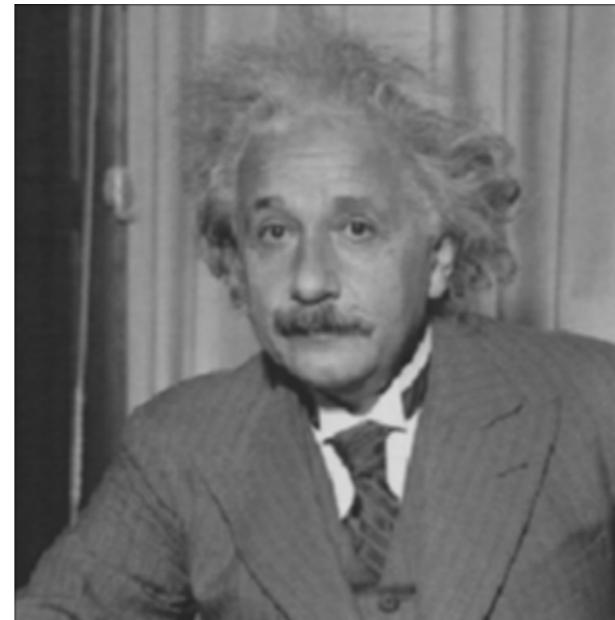


after

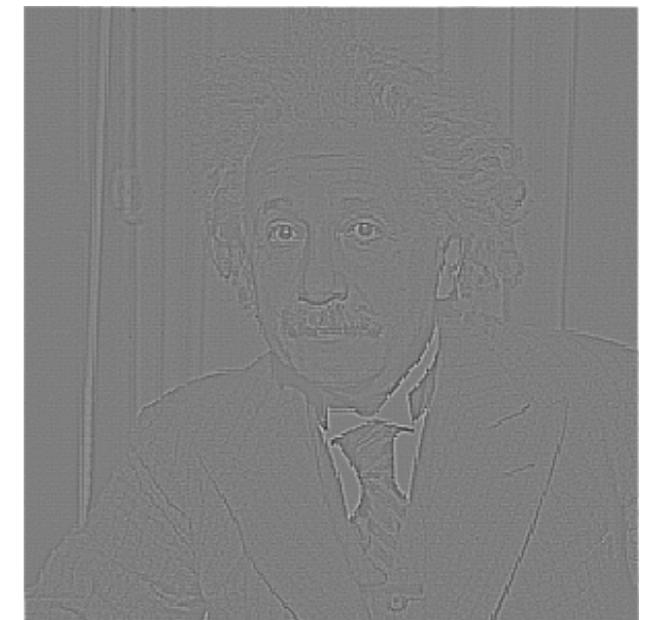
Sharpening



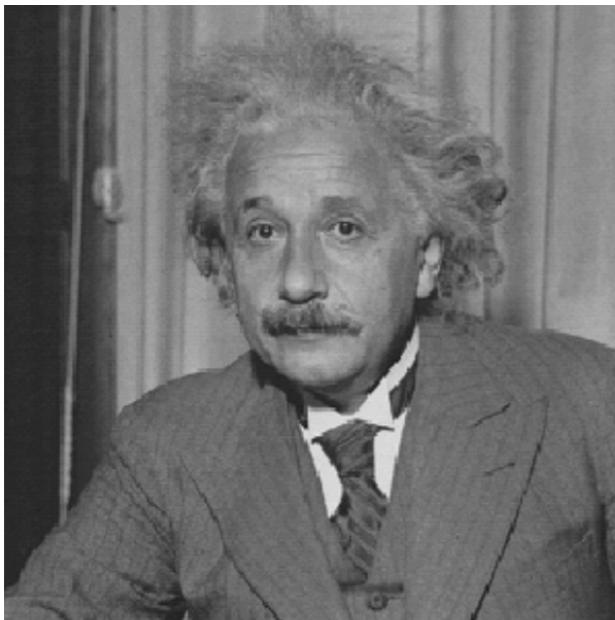
Image



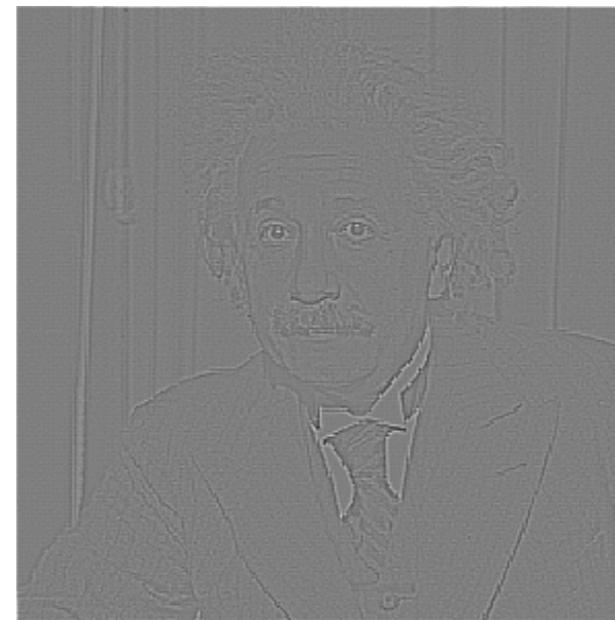
Blurred



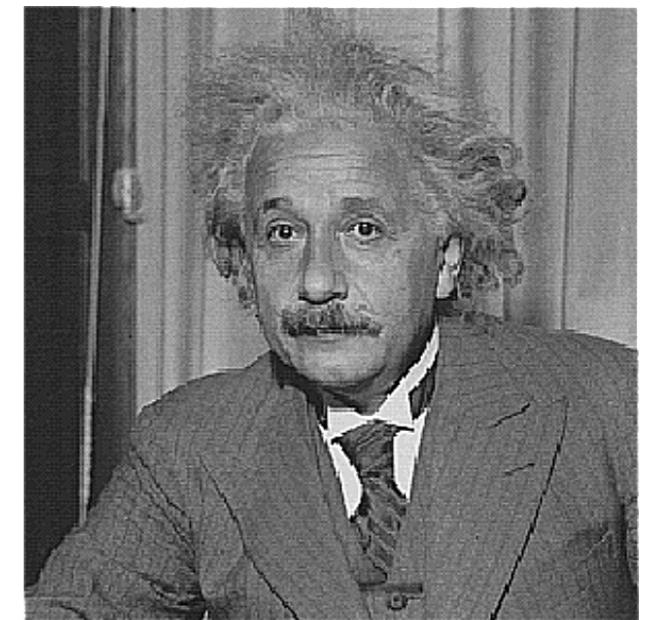
Detail



Image



Detail



Sharpened

Convolution Properties

Commutative:

$$F * H = H * F$$

Associative:

$$(F * H) * G = F * (H * G)$$

Distributive:

$$(F * G) + (H * G) = (F + H) * G$$

Convolution Properties

Scale:

$$\text{filter}(A * f) = A * \text{filter}(f)$$

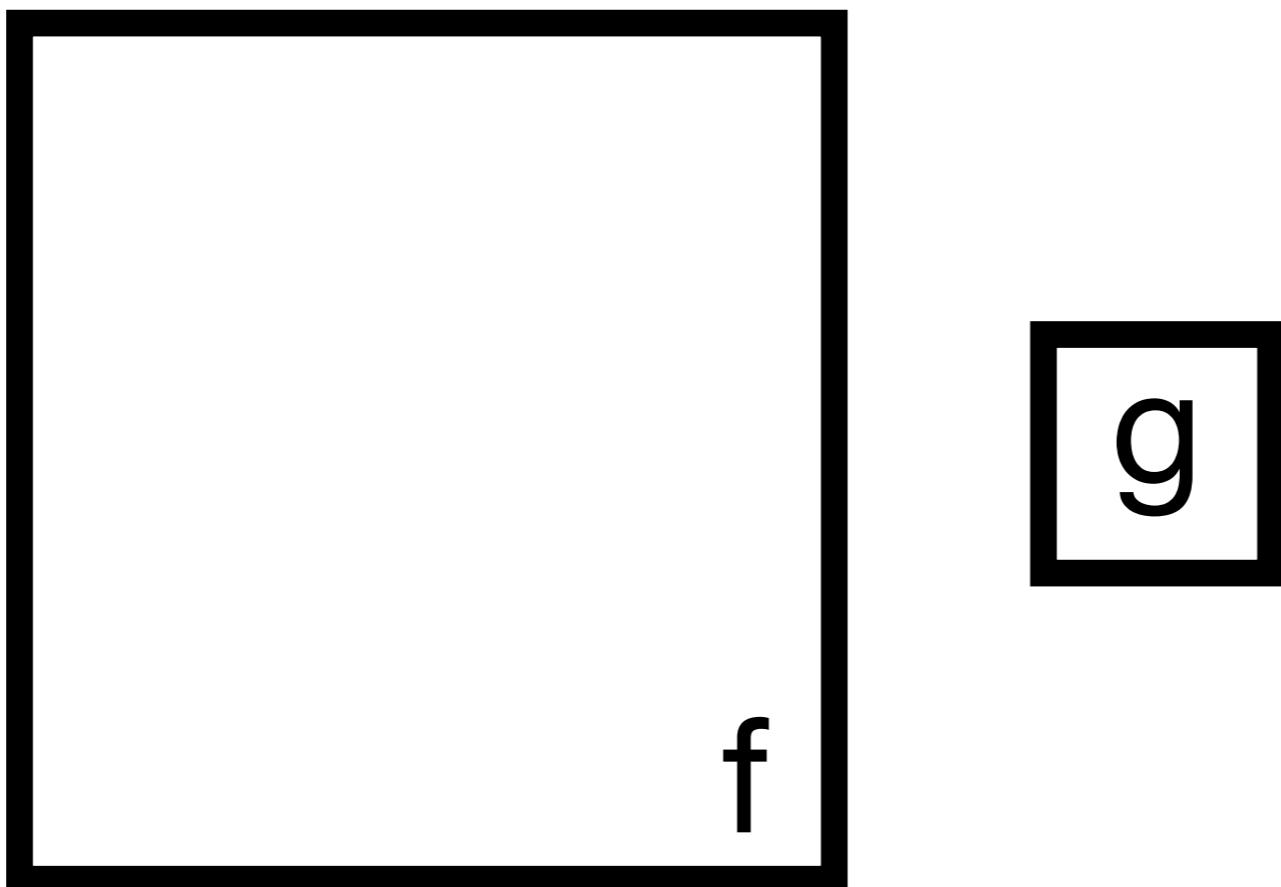
Shift Invariance:

$$\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$$

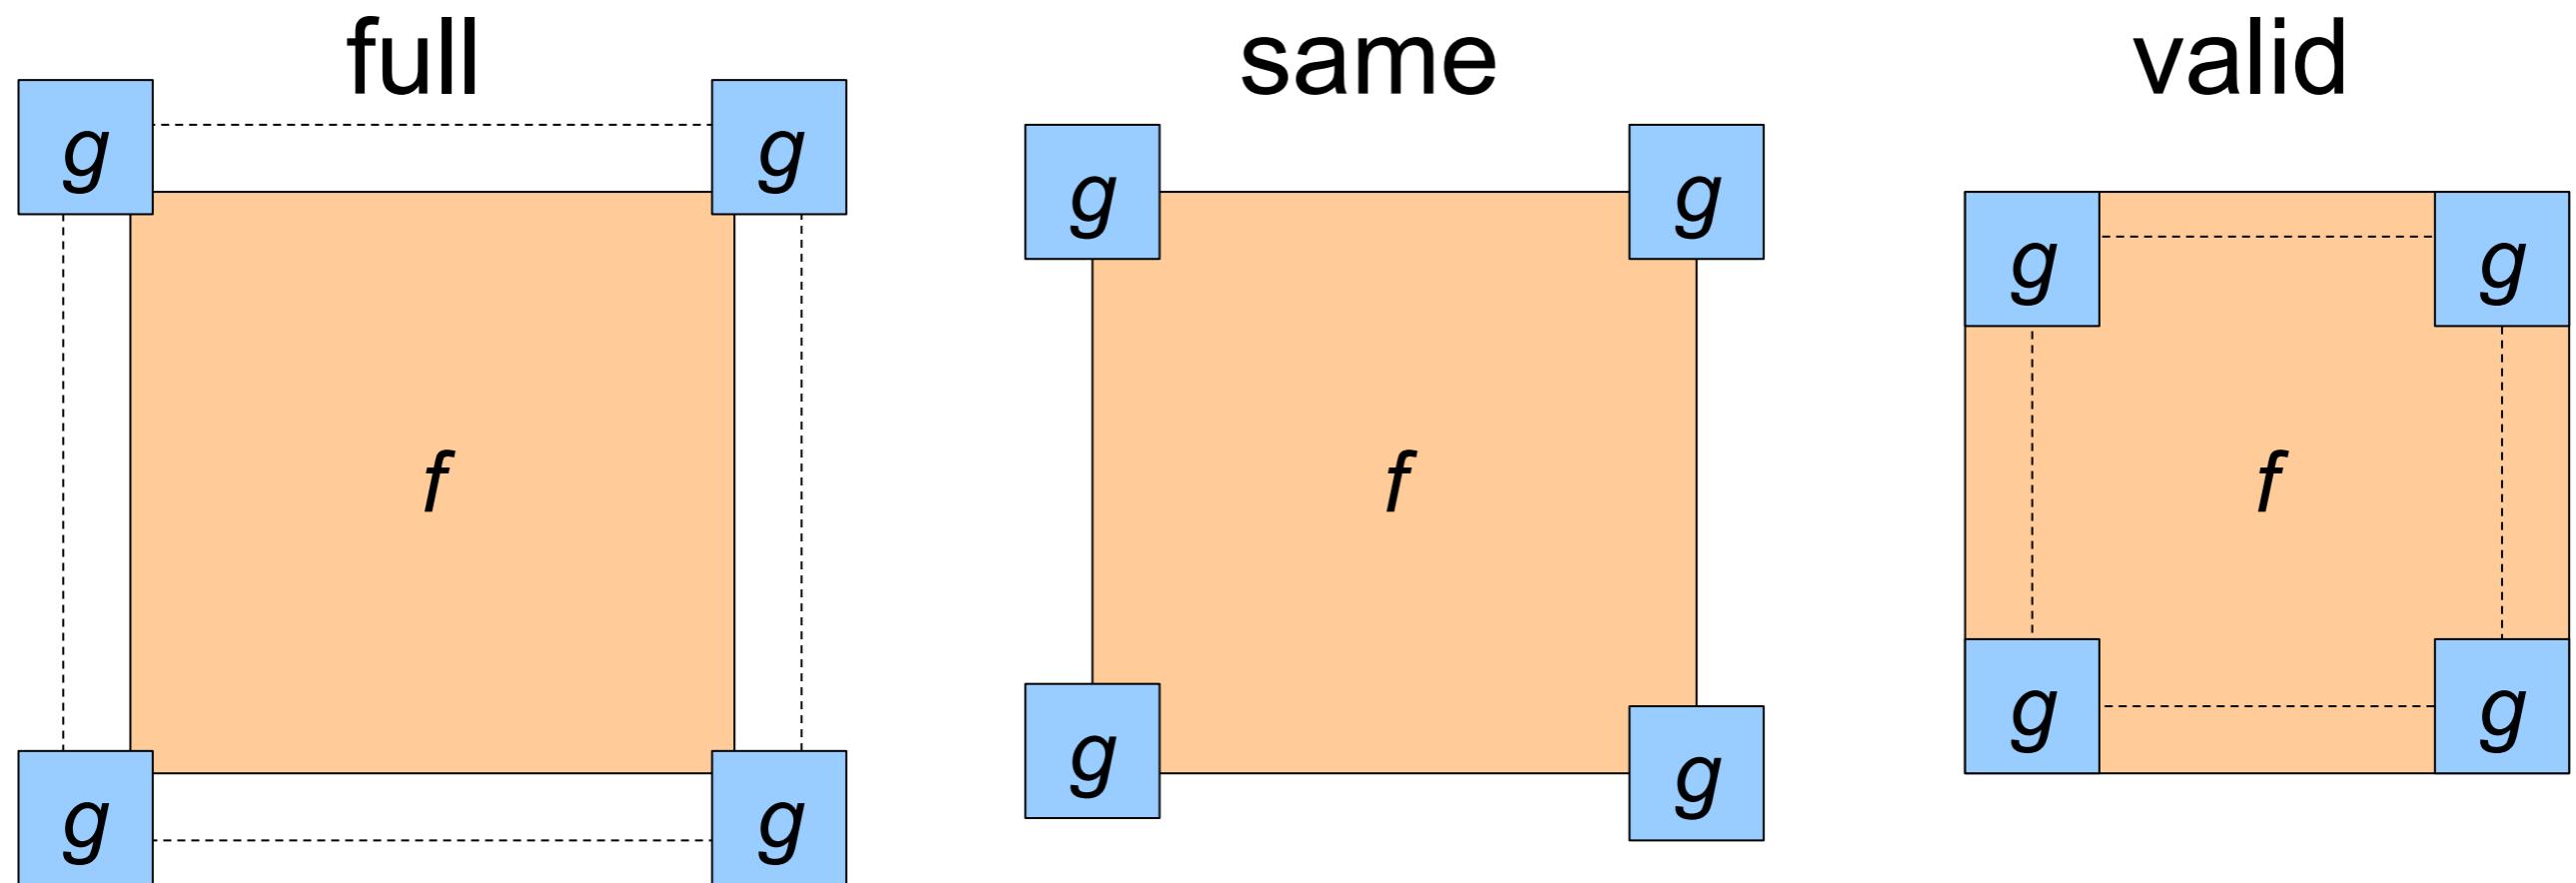
Cross-Correlation

- Conceptually simpler, but not as nice properties:

$$(f * g)[x, y] = \sum_{i,j} f[x + i, y + j]g[i, j]$$



Boundary Issues



Border Padding



Zero Pad



Circular



Replicate

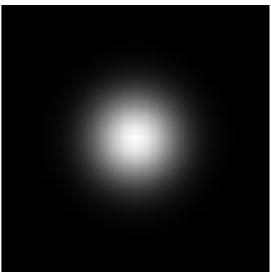


Symmetric

Box Filter

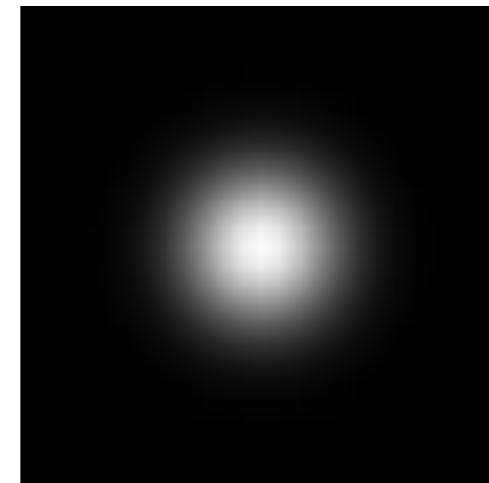
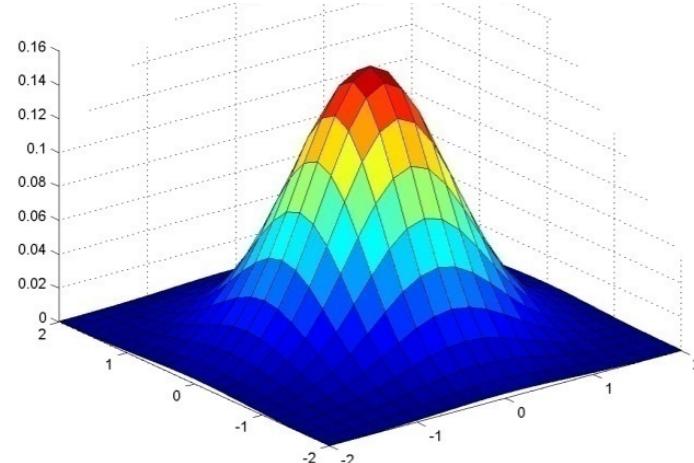


Gaussian Filter



Gaussian Filter

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



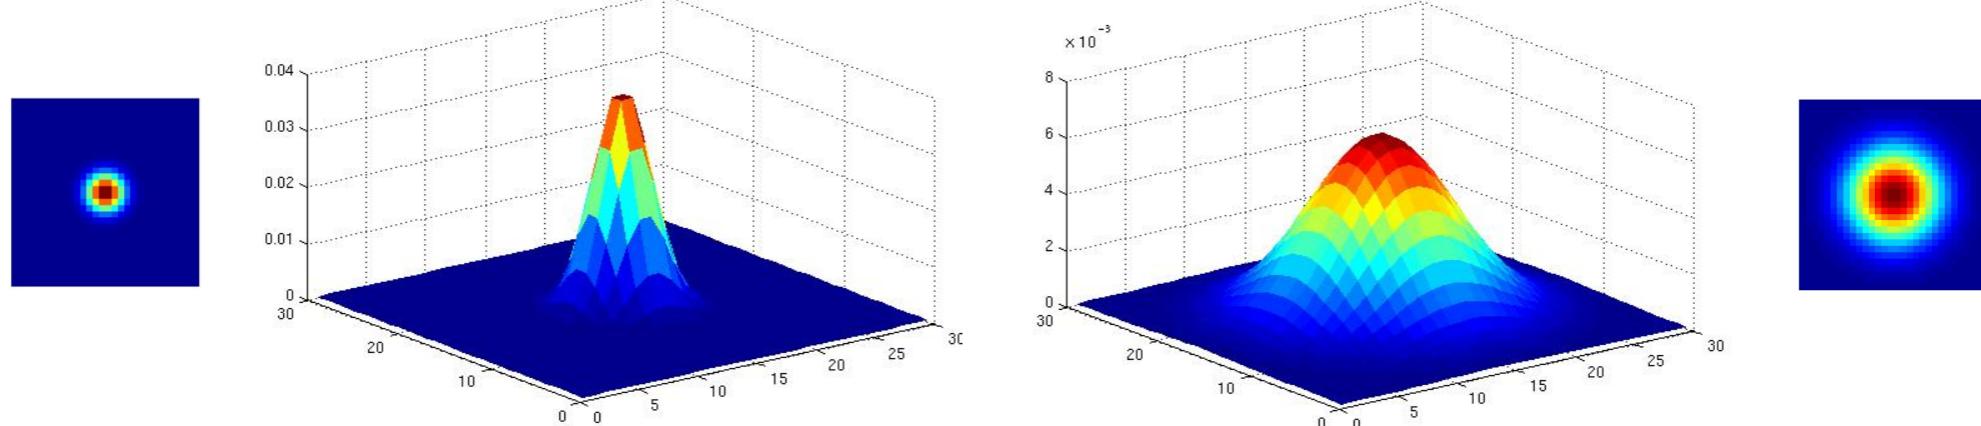
0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$5 \times 5, \sigma = 1$

Constant factor at front makes volume sum to unity

Standard Deviation

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



$\sigma = 2$ with
30 x 30
kernel

$\sigma = 5$ with
30 x 30
kernel

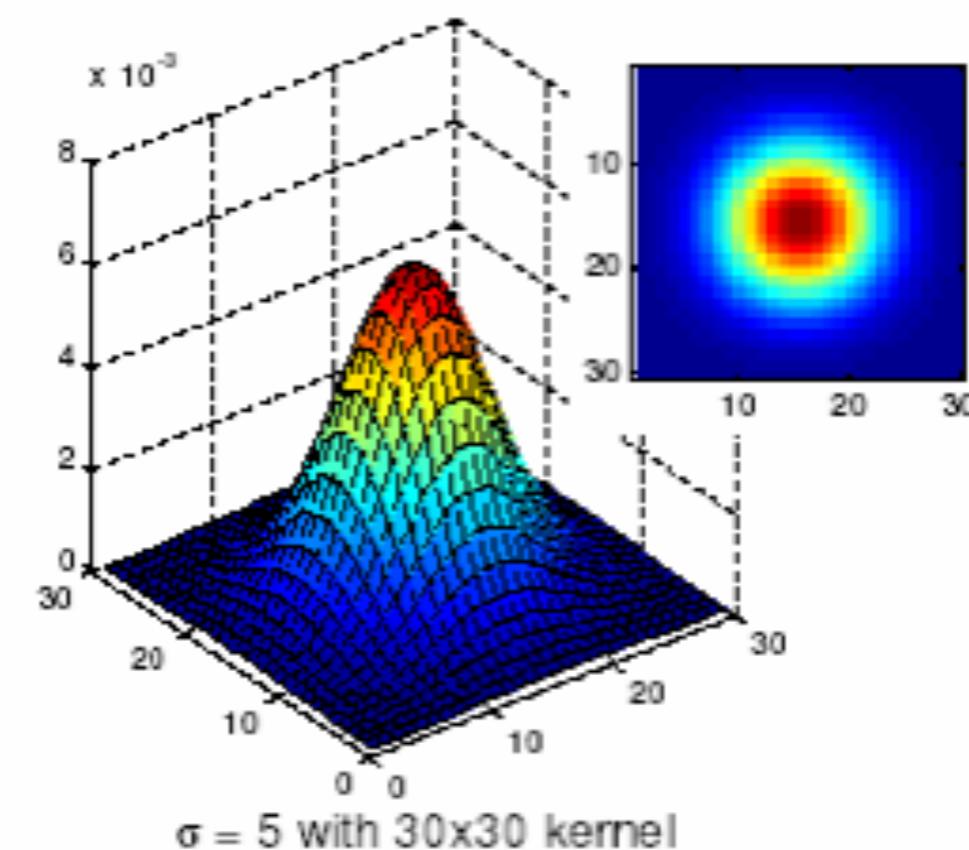
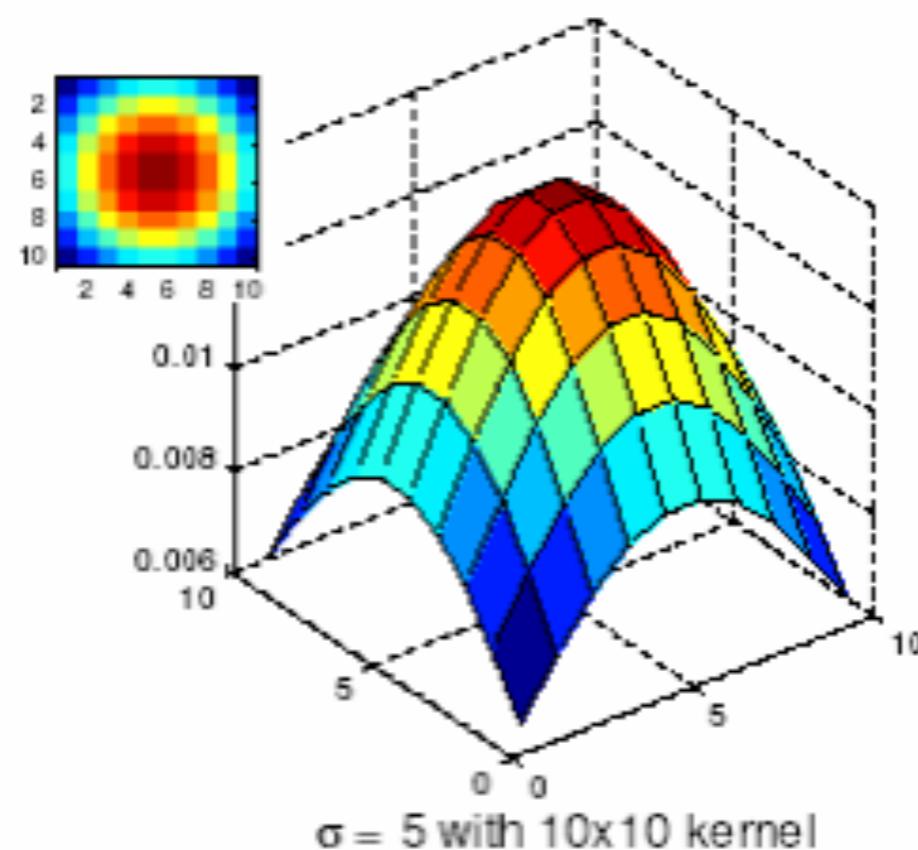
Standard deviation σ : determines extent of smoothing

Changing Sigma



Kernel Width

The Gaussian function has infinite support, but discrete filters use finite kernels



Rule of thumb: set filter half-width to about 3σ

Complexity

What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?

Complexity

What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?

$$O(n^2 m^2)$$

Separable Convolution

Two dimensional Gaussian is product of two Gaussians:

$$\begin{aligned} G(x, y) &= \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y^2}{2\sigma^2}\right) \right) \end{aligned}$$

Take advantage of associativity:

$$f * \boxed{\text{ }} = f * \boxed{\text{ }} * \boxed{\text{ }}$$

Complexity

What is the complexity of filtering an
 $n \times n$ image with an $m \times m$ kernel?

$$O(n^2 m^2)$$

What if kernel is separable?

Complexity

What is the complexity of filtering an
 $n \times n$ image with an $m \times m$ kernel?

$$O(n^2 m^2)$$

What if kernel is separable?

$$O(n^2 m)$$

Denoising

Additive Gaussian Noise



Gaussian Filter (sigma=1)



What's wrong?

Salt and Pepper Noise

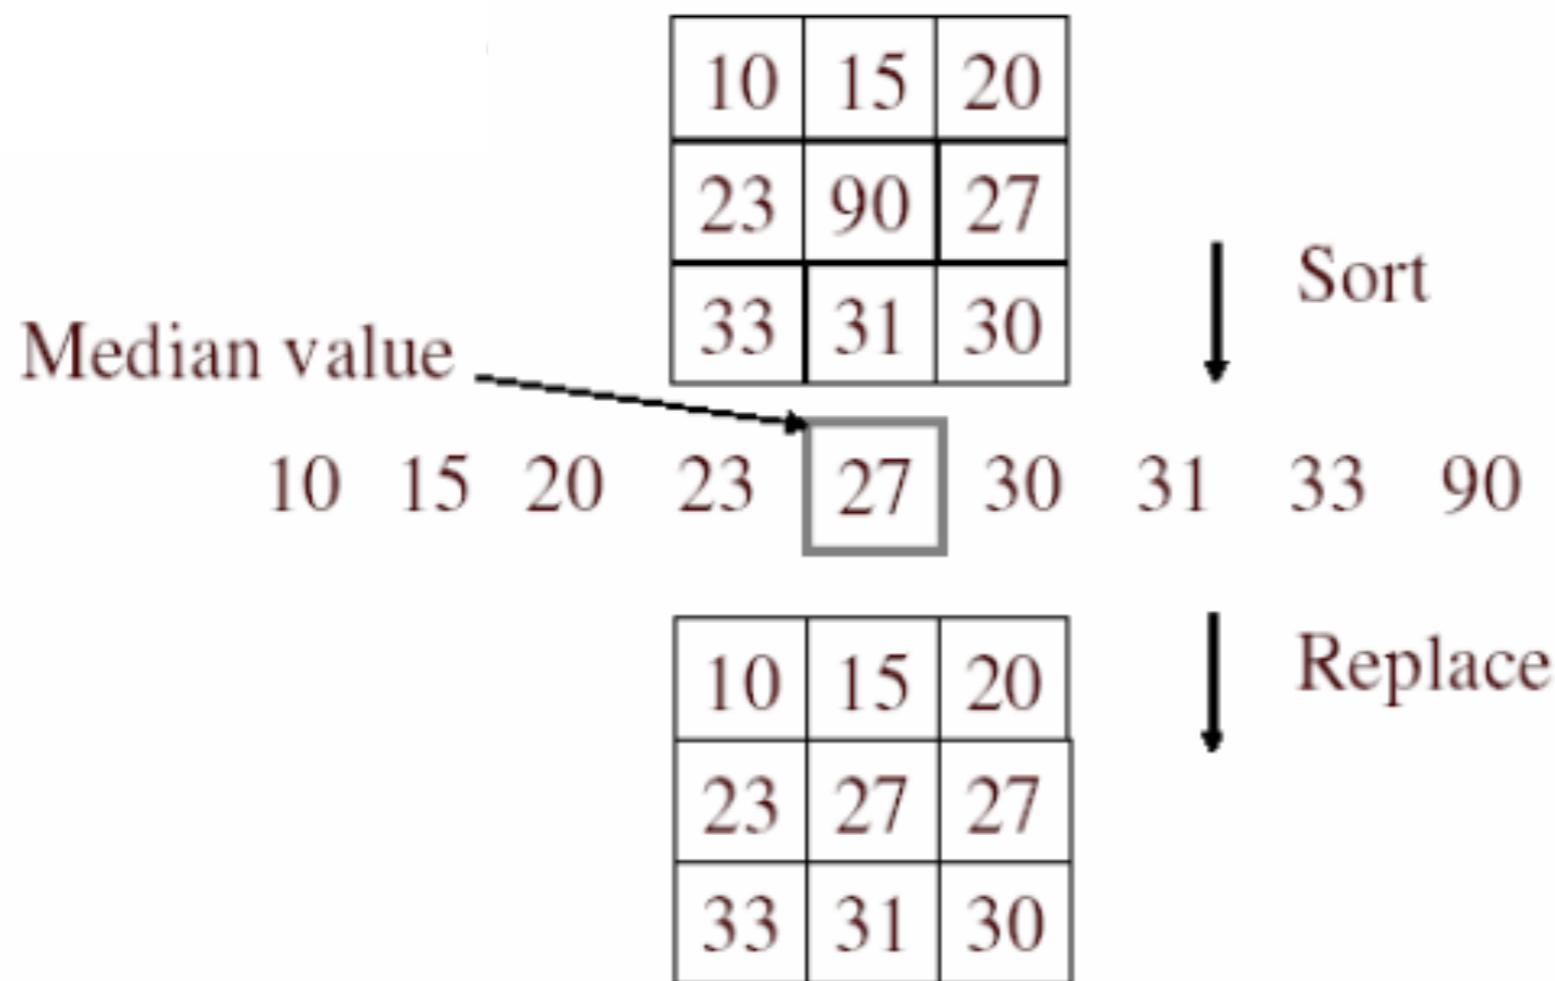


Gaussian Filter (sigma=1)



Median Filter

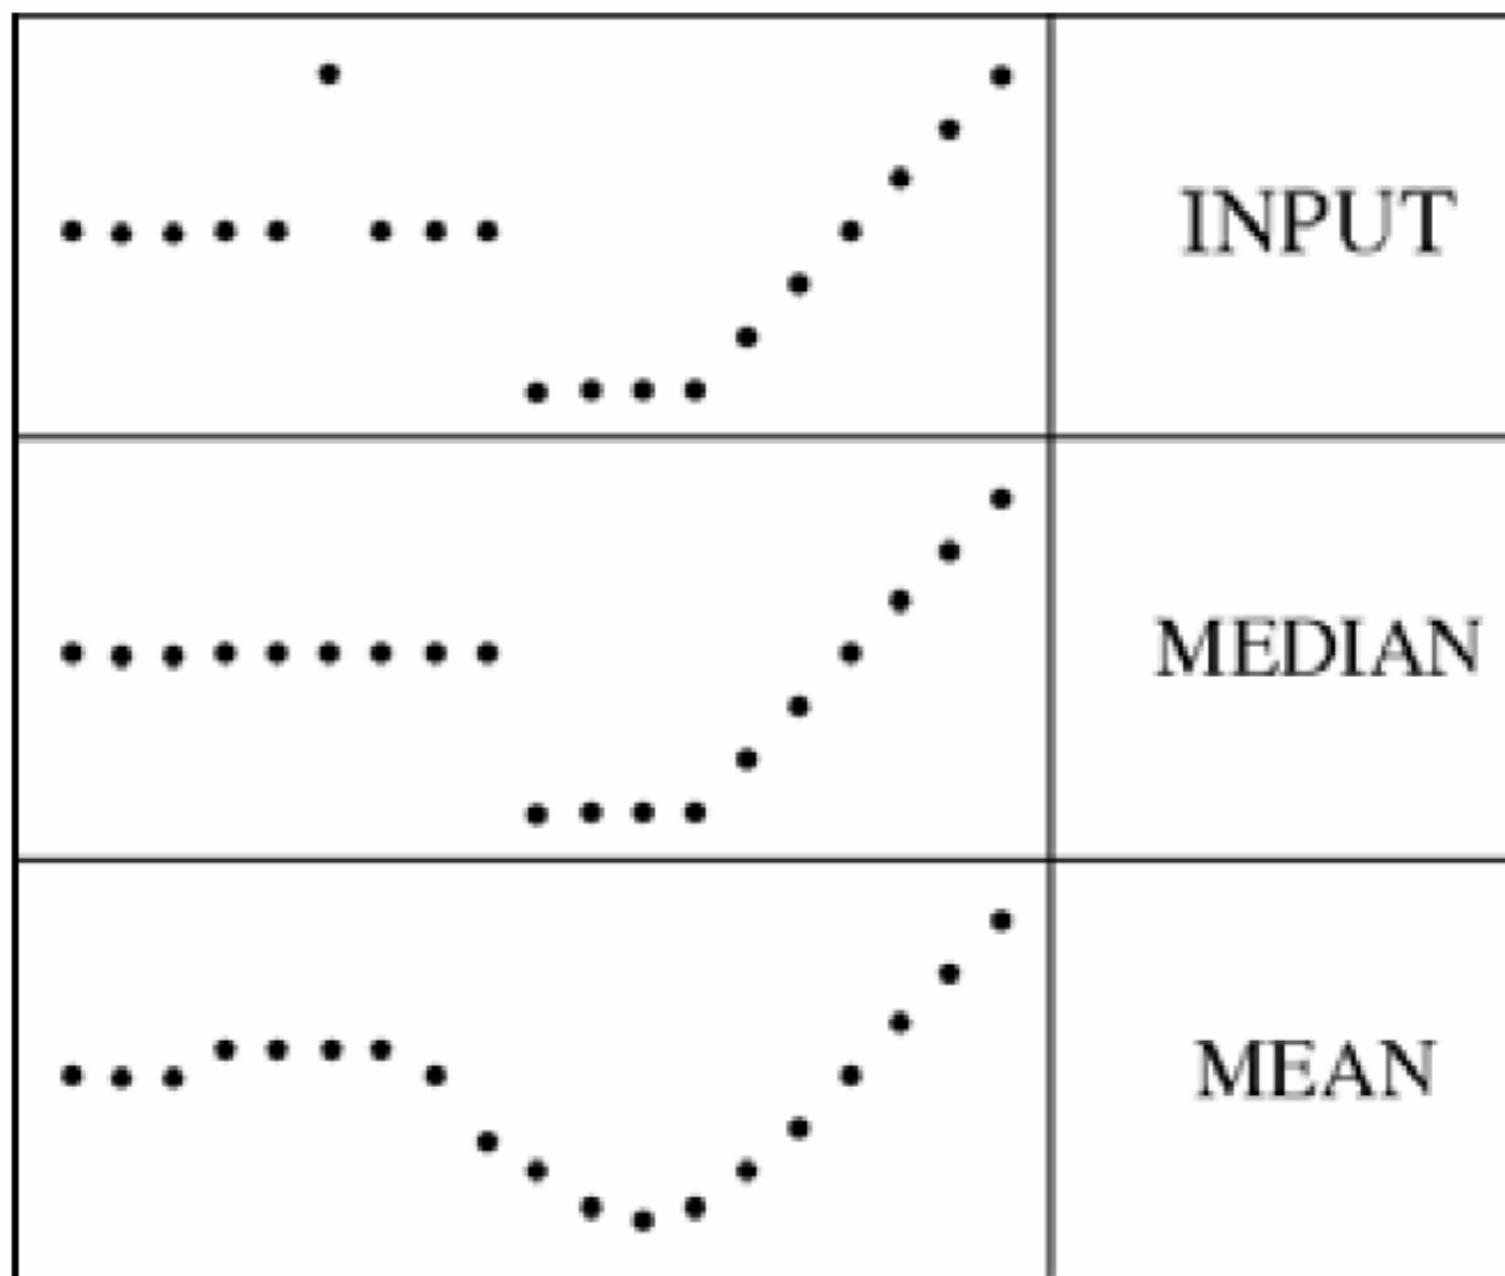
- A **median filter** operates over a window by selecting the median intensity in the window



Is median filtering linear?

Why use median?

filters have width 5 :



Median Filtering

Salt and Pepper Noise



Median Filter 3x3



Median Filtering

Median 3x3



Median 5x5



Median 9x9



Image Gradients

Image Gradients



How does intensity change as you move left to right?

How do you take the ***derivative*** of an image?

First Derivative



$$* [-1, 1] =$$



$$* [-1, 1]^T =$$



Second Derivative



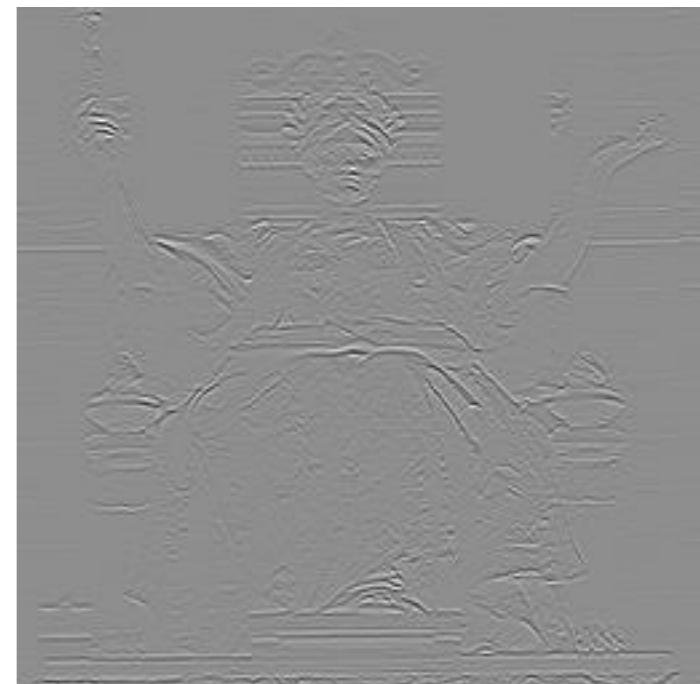
$$* [-1, 1] =$$



$$\frac{\partial^2 I}{\partial x^2}$$



$$* [-1, 1]^T =$$

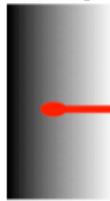


$$\frac{\partial^2 I}{\partial y^2}$$

Image Gradients

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

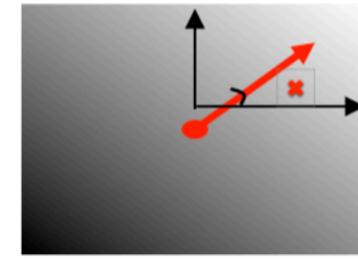
It points in the direction of most rapid change in intensity



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient direction is given by:

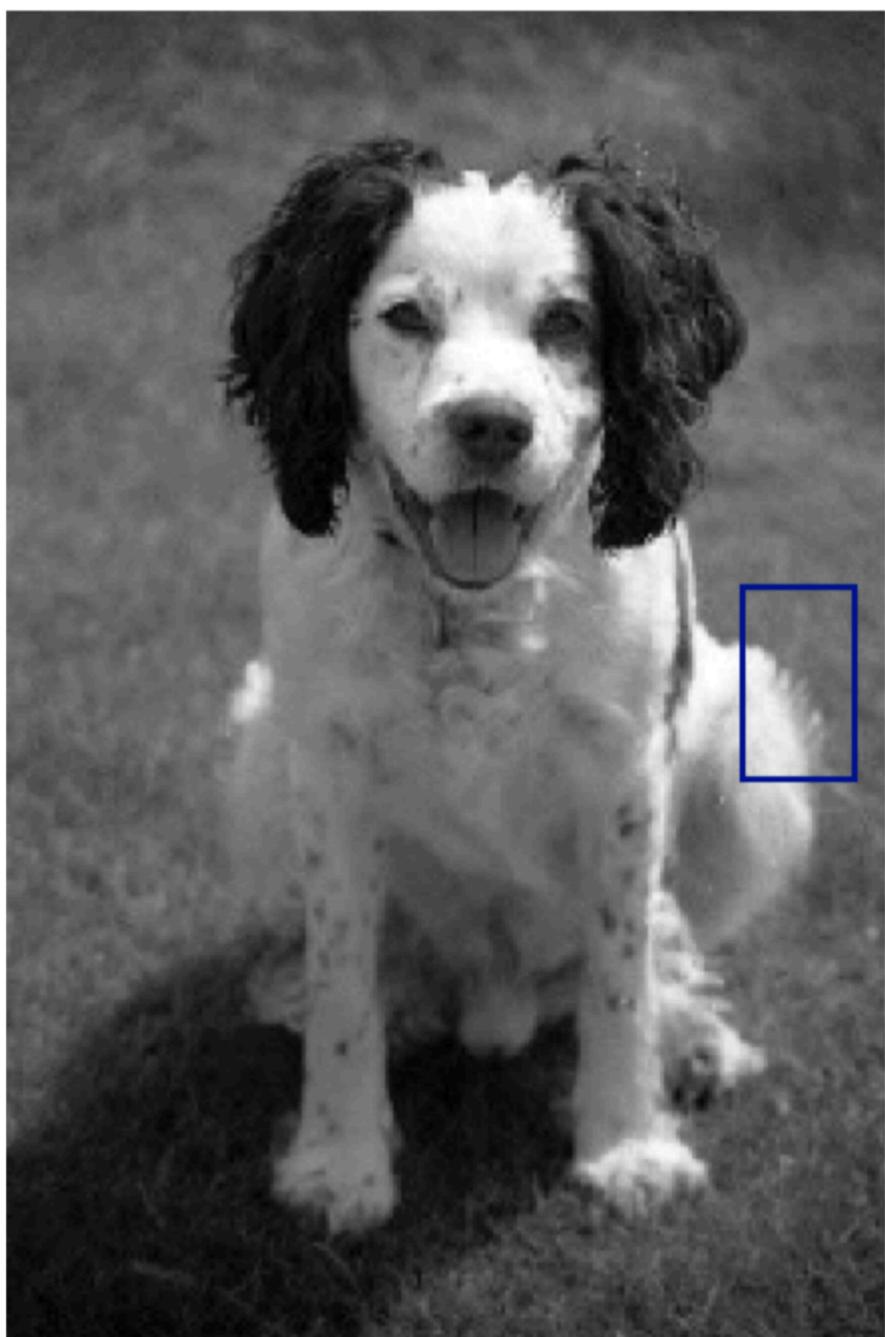
$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- ◆ how does this relate to the direction of the edge?

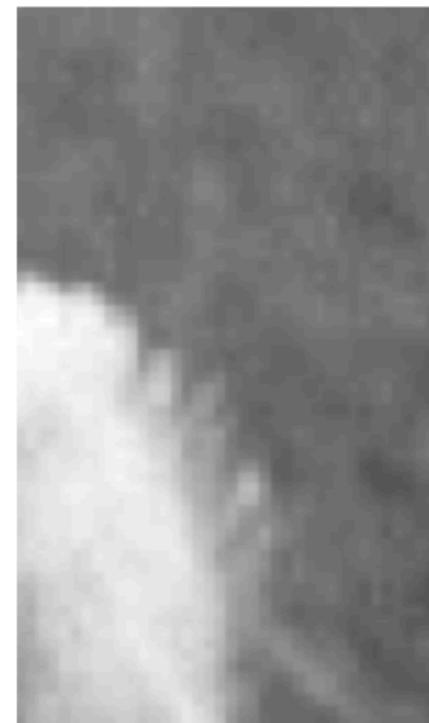
The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

What causes an edge?



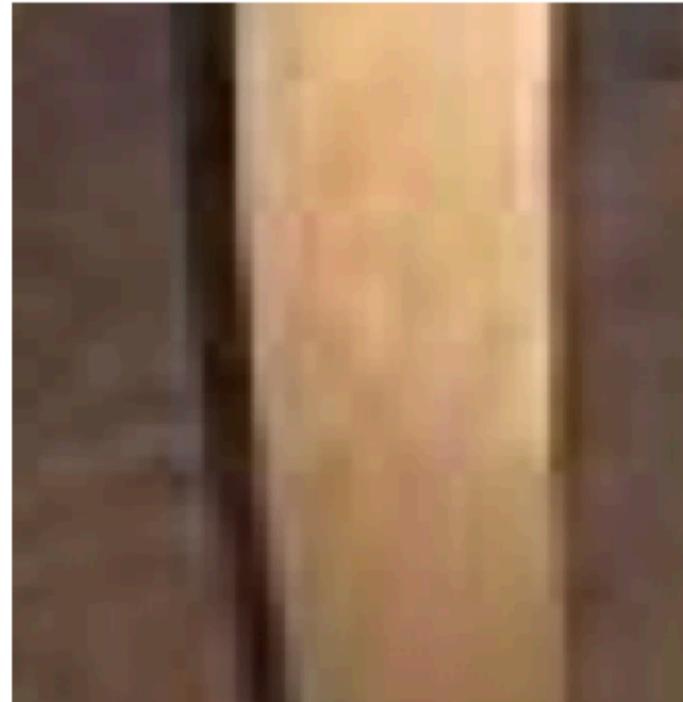
Object Boundaries



What causes an edge?

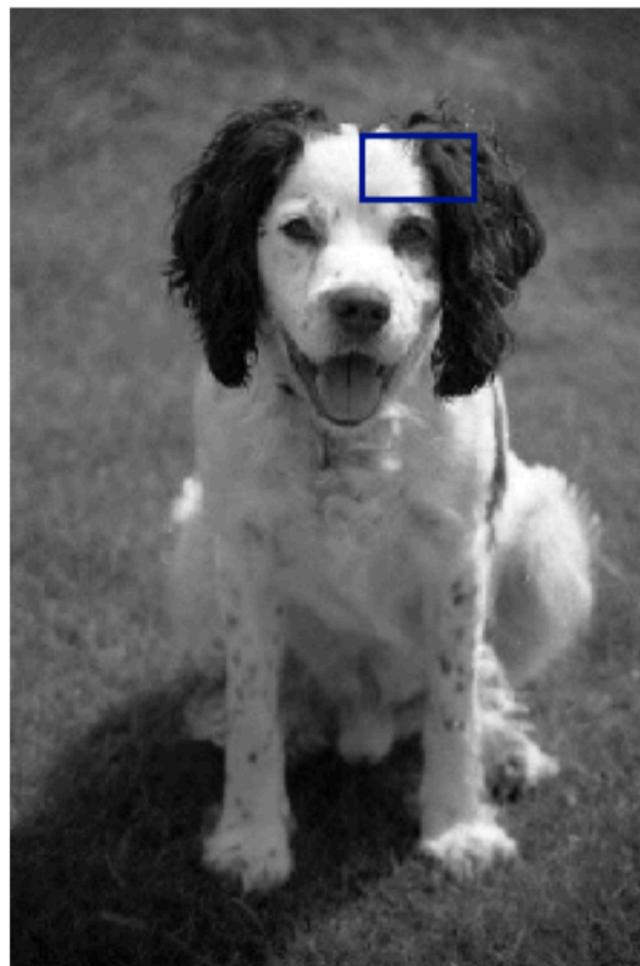


Surface normal discontinuities



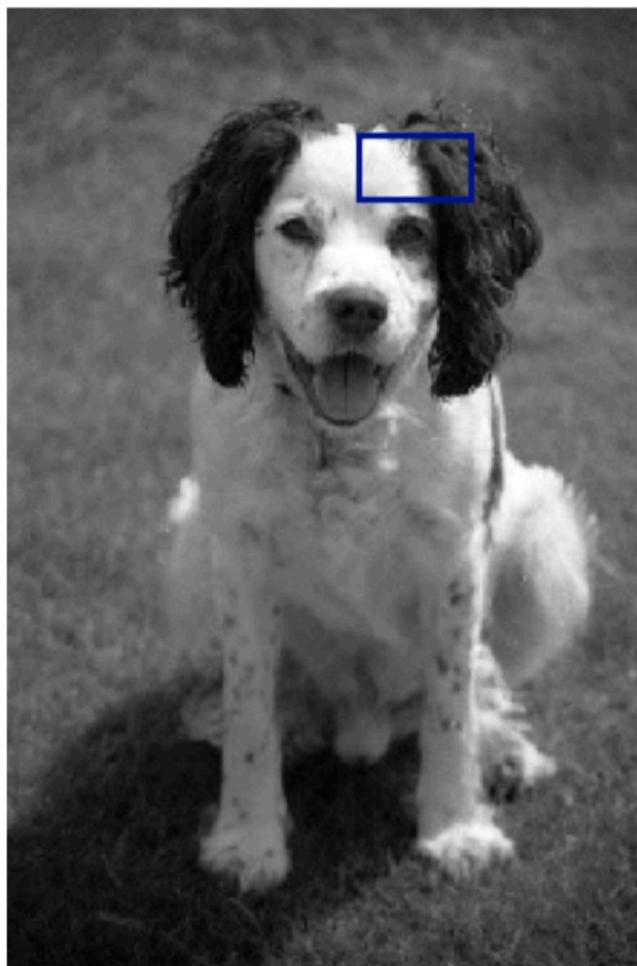
What causes an edge?

Boundaries of material properties



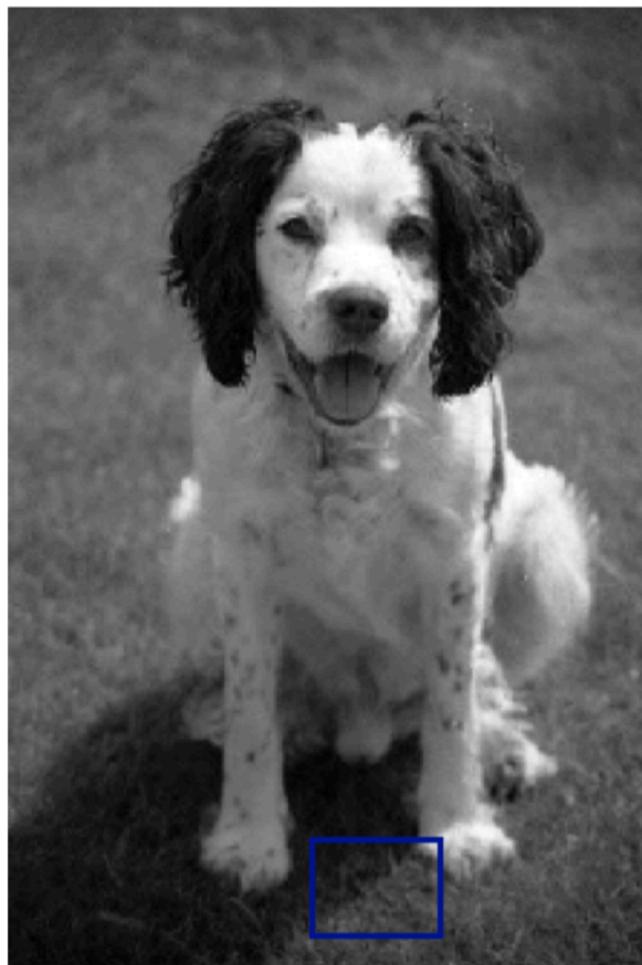
What causes an edge?

Boundaries of material properties



What causes an edge?

Boundaries of lighting



Edge Types



Step



Ridge

Which of these do you suppose
a derivative filter detects best?

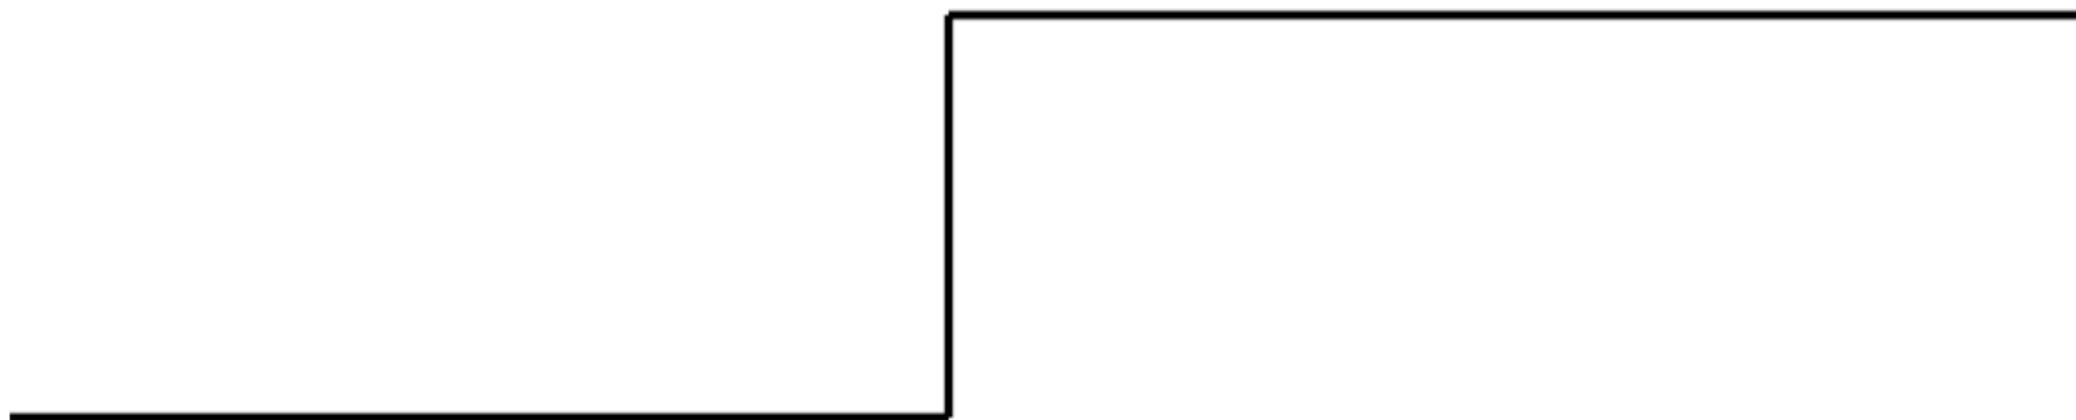


Roof

What is an edge?

Change is measured by derivative in 1D

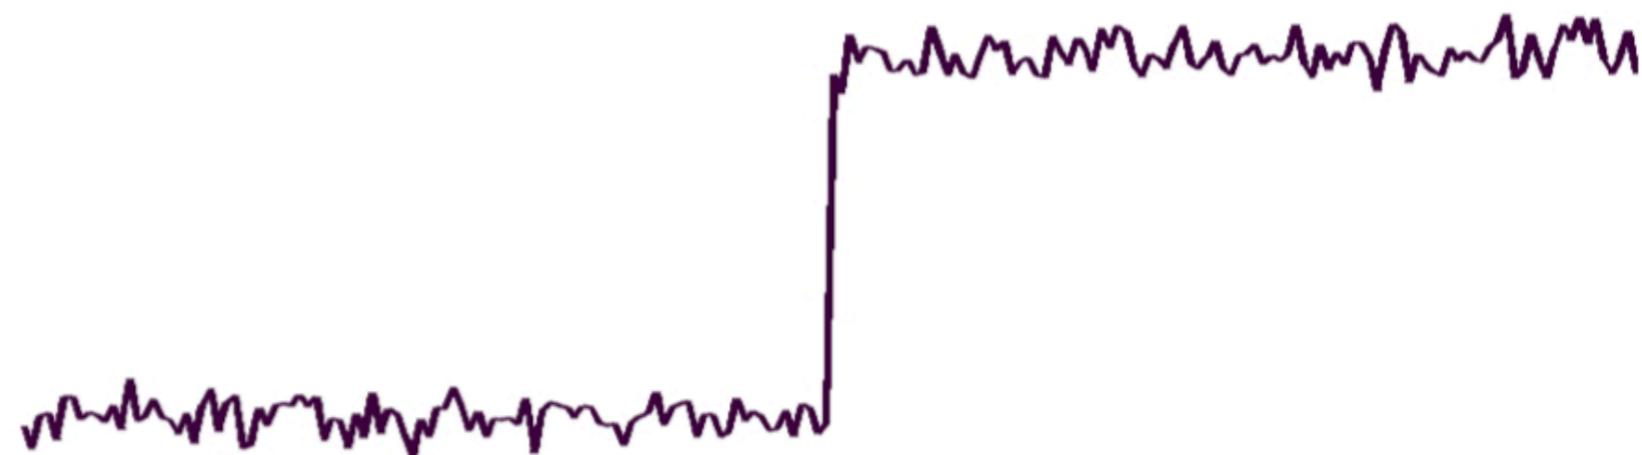
- Biggest change, derivative has maximum magnitude
- Or 2nd derivative is zero.



What about noise?

Derivative is high everywhere.

Must smooth before taking gradient.



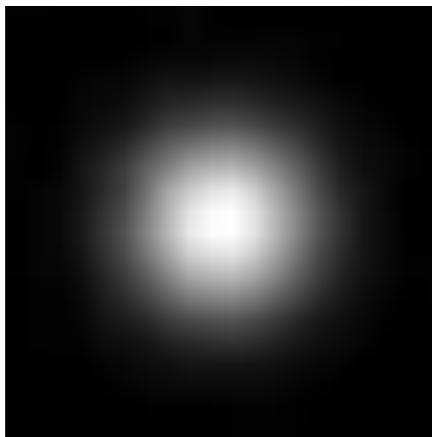
Handling Noise

- Filter with a Gaussian to smooth, then take gradients
- But, convolution is linear

Gaussian Filter

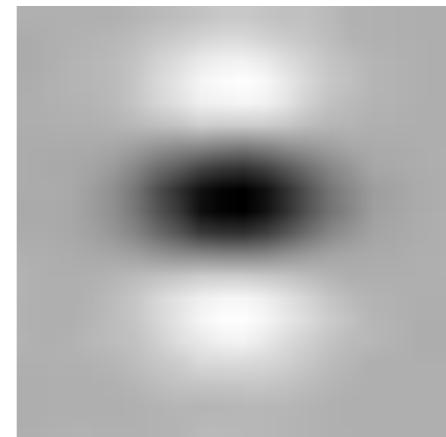
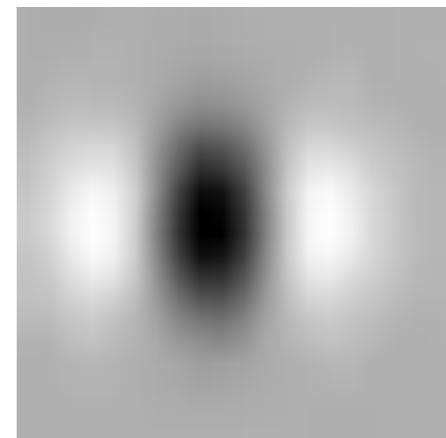


$$* [-1,1] * [-1,1] =$$



$$* [-1,1]^T * [-1,1]^T =$$

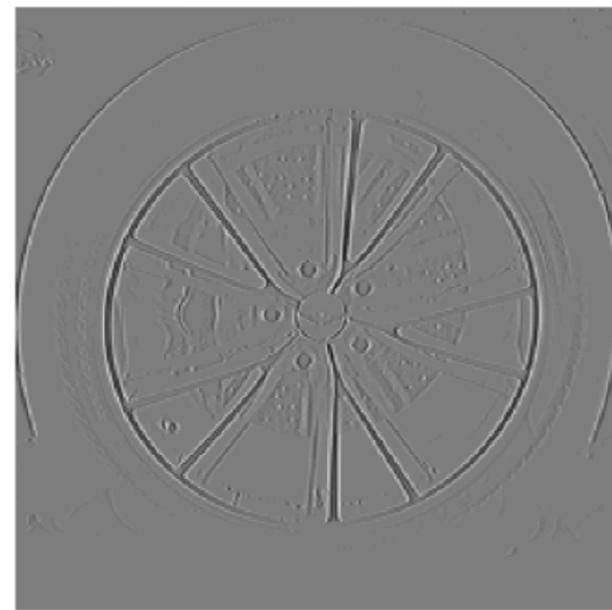
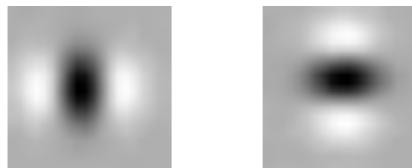
Laplacian Filter



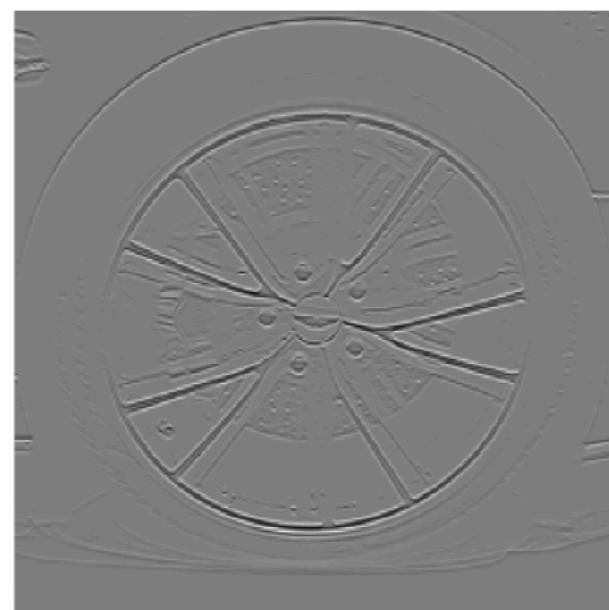
The Laplacian Filter

- Popularized by Marr and Hildreth in 1980 to locate boundaries between objects
- Defined as the sum of second order partial derivatives:

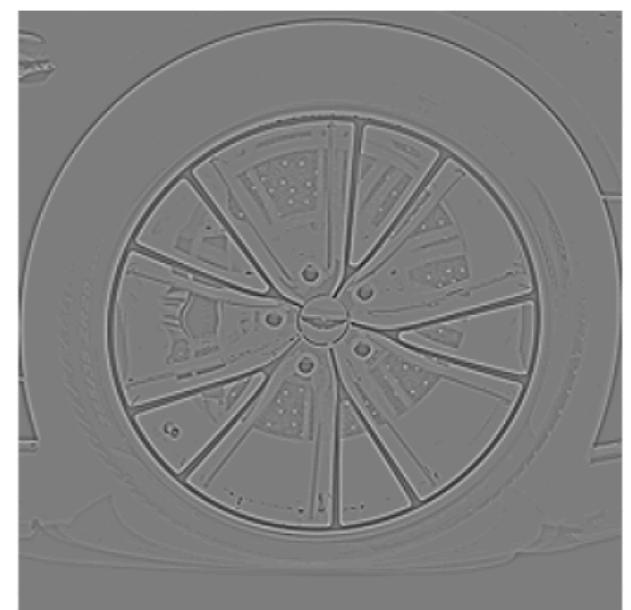
$$\nabla I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$



dx



dy

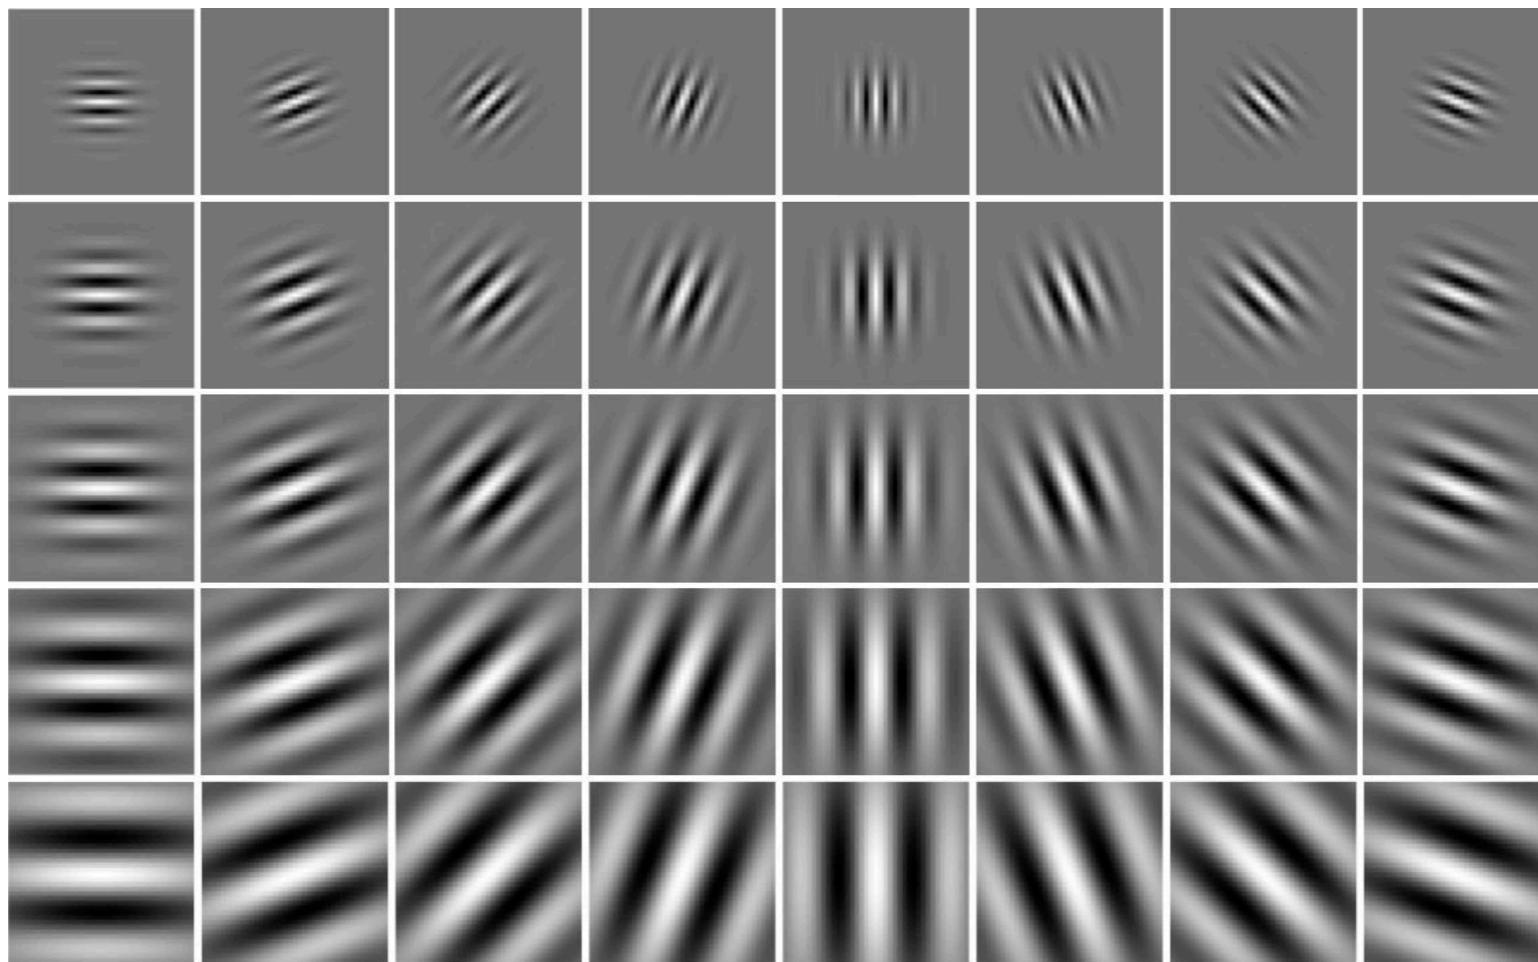


laplacian

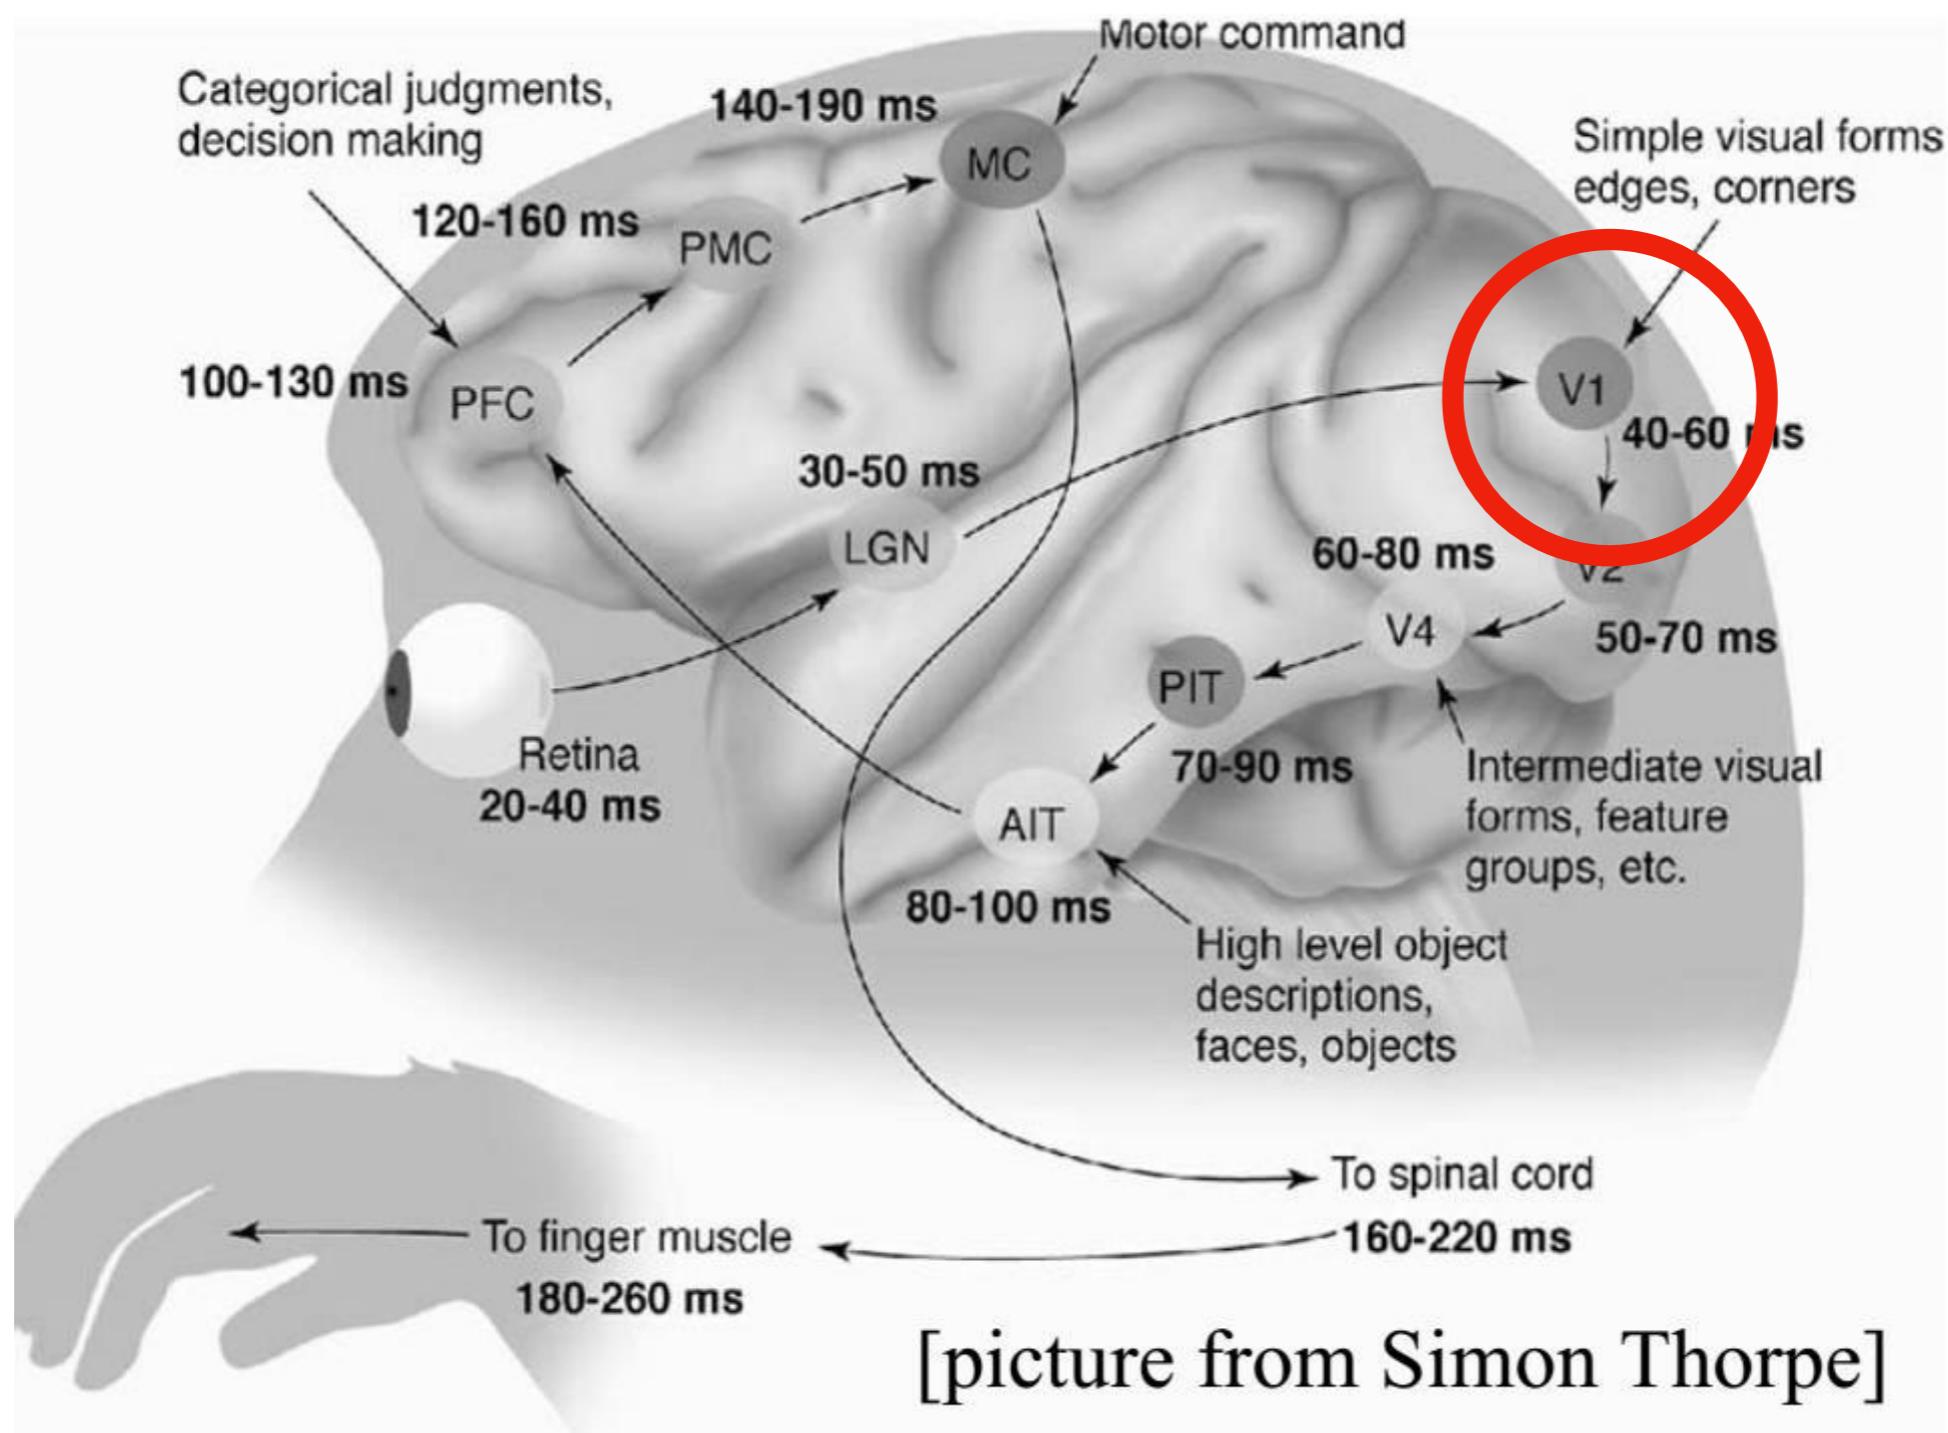
Aside: Gabor Filters

Cosine wave multiple by a Gaussian

$$\psi(x, y) = e^{-\frac{x^2 + y^2}{2\sigma^2}} \cos(2\pi\mu x)$$



Aside: Human Visual System



Aside: Cat Visual System

IEEE TRANSACTIONS ON ACOUSTICS, SPEECH, AND SIGNAL PROCESSING, VOL. 36, NO. 7, JULY 1988

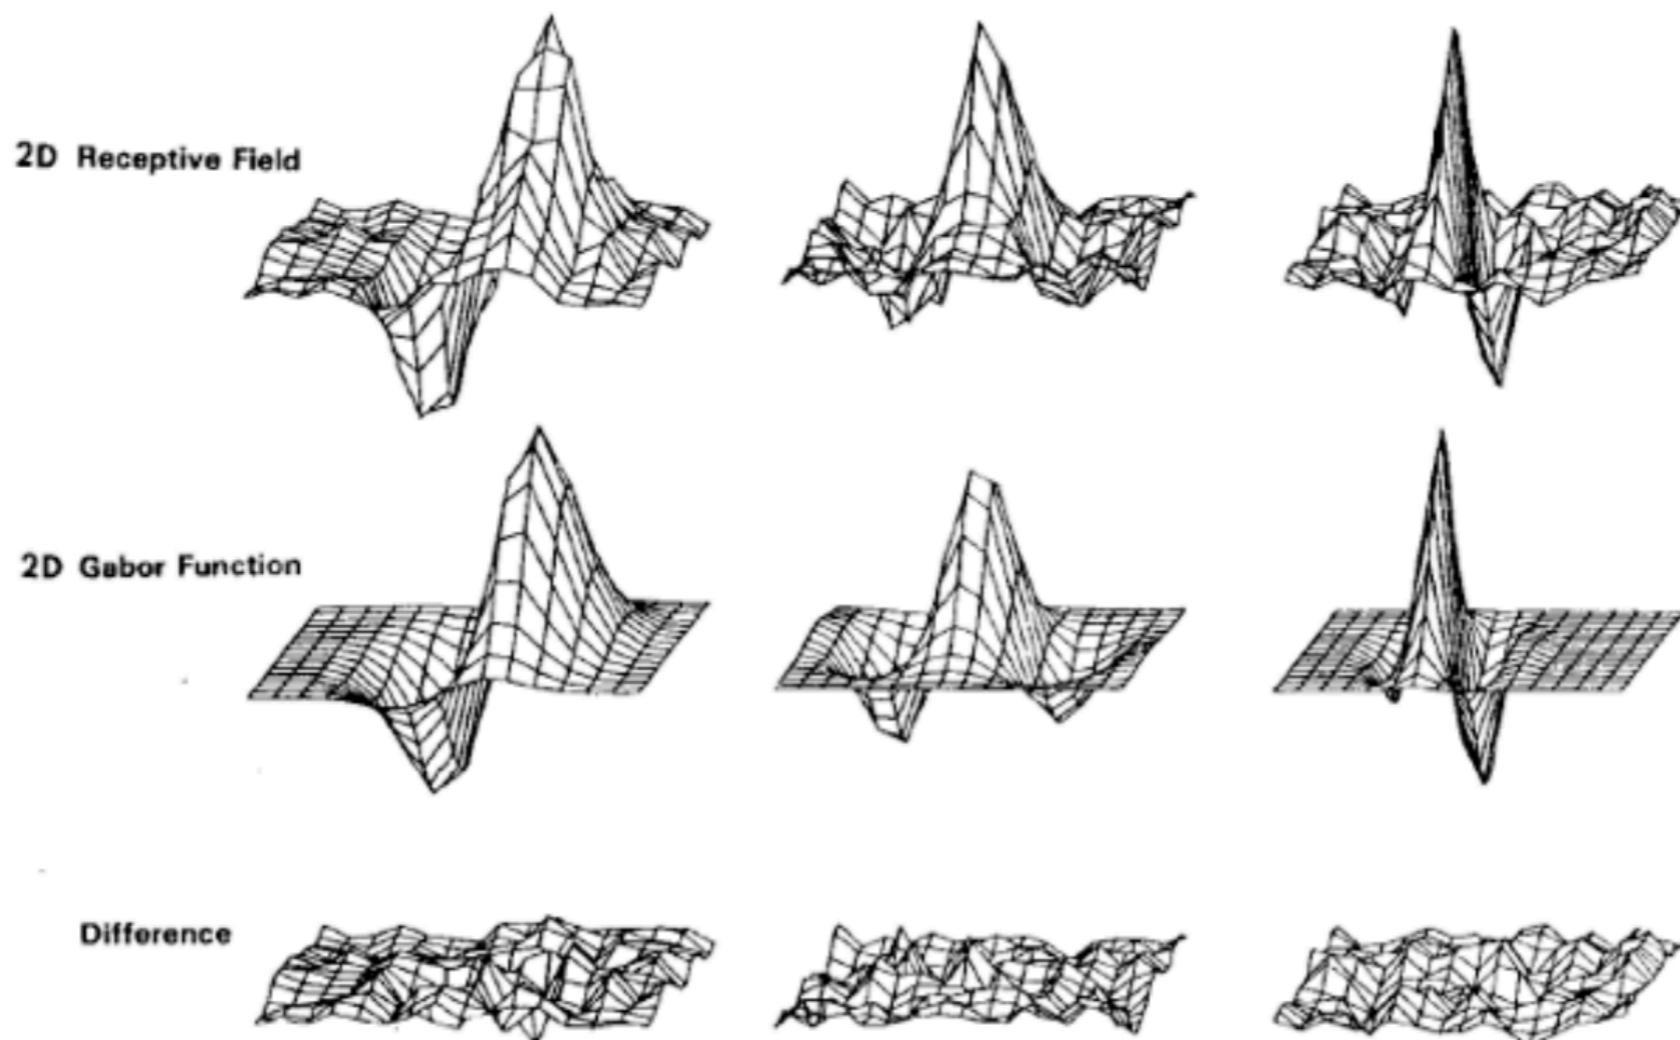
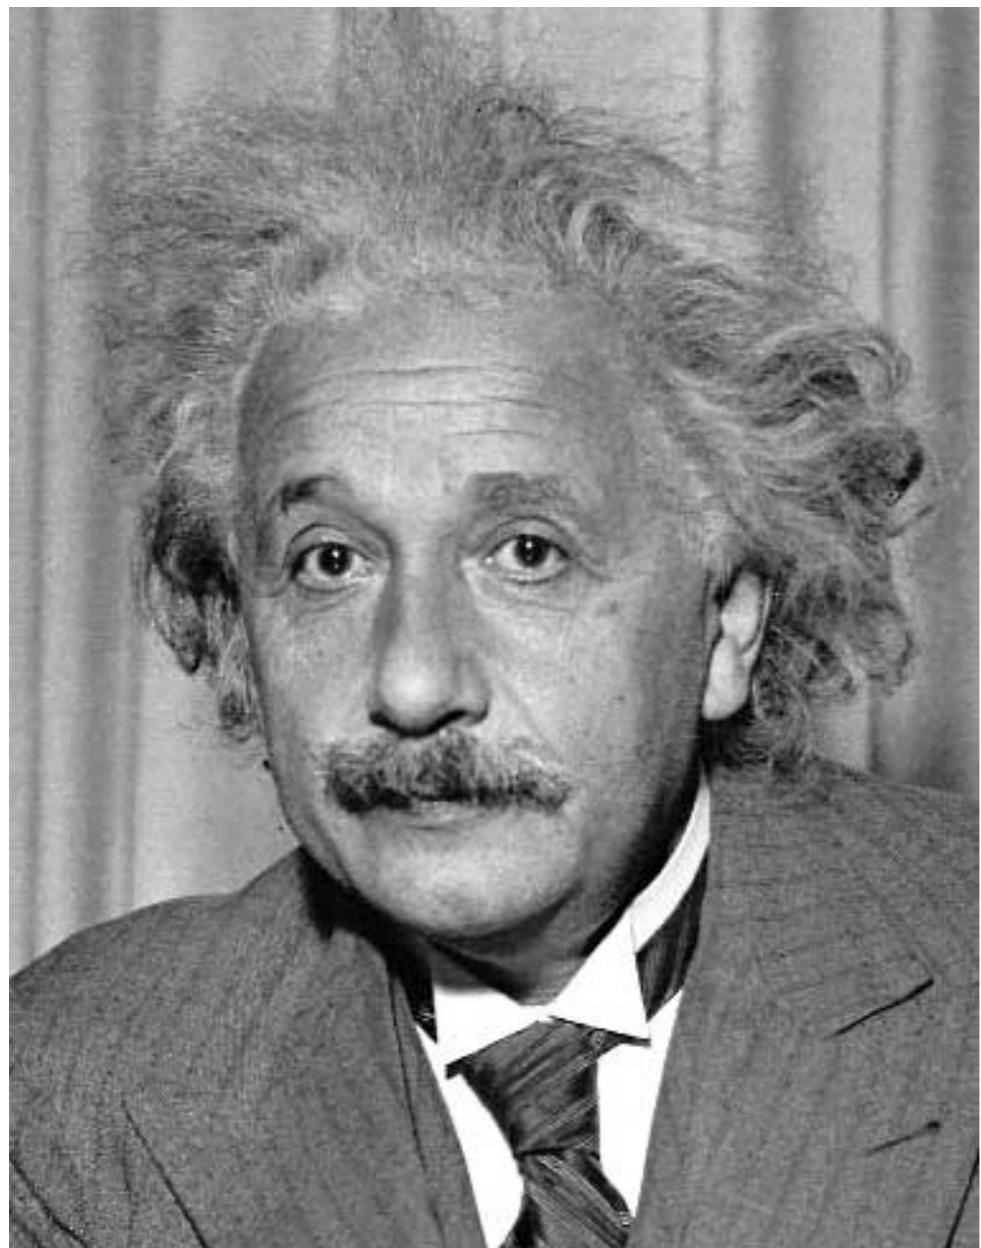


Fig. 5. Top row: illustrations of empirical 2-D receptive field profiles measured by J. P. Jones and L. A. Palmer (personal communication) in simple cells of the cat visual cortex. Middle row: best-fitting 2-D Gabor elementary function for each neuron, described by (10). Bottom row: residual error of the fit, indistinguishable from random error in the Chi-squared sense for 97 percent of the cells studied.

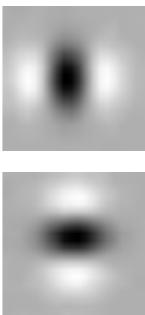
Detection

Finding Boundaries

f



g

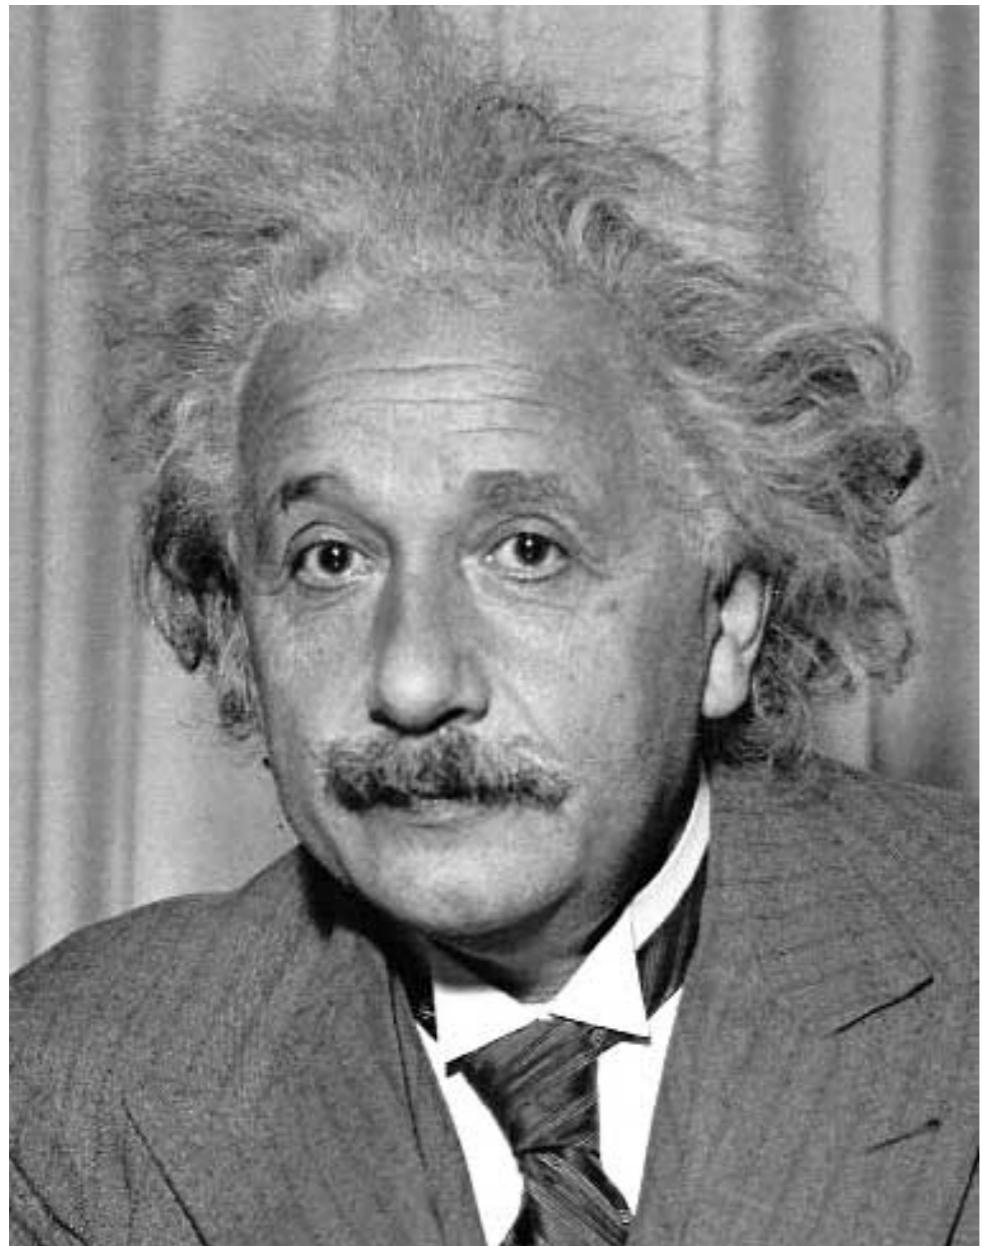


$$\frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} > \lambda$$



Finding Things

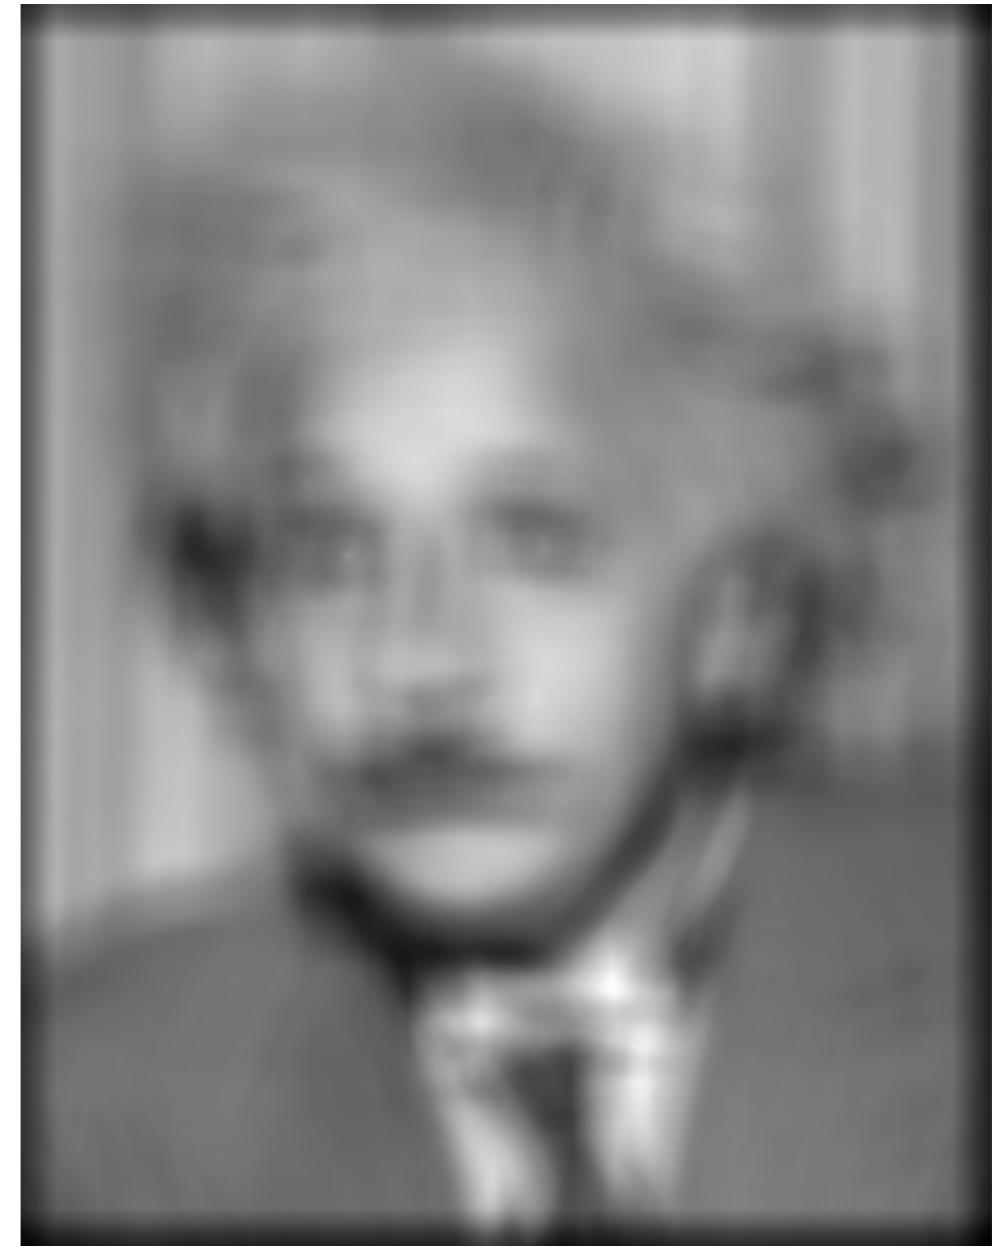
f



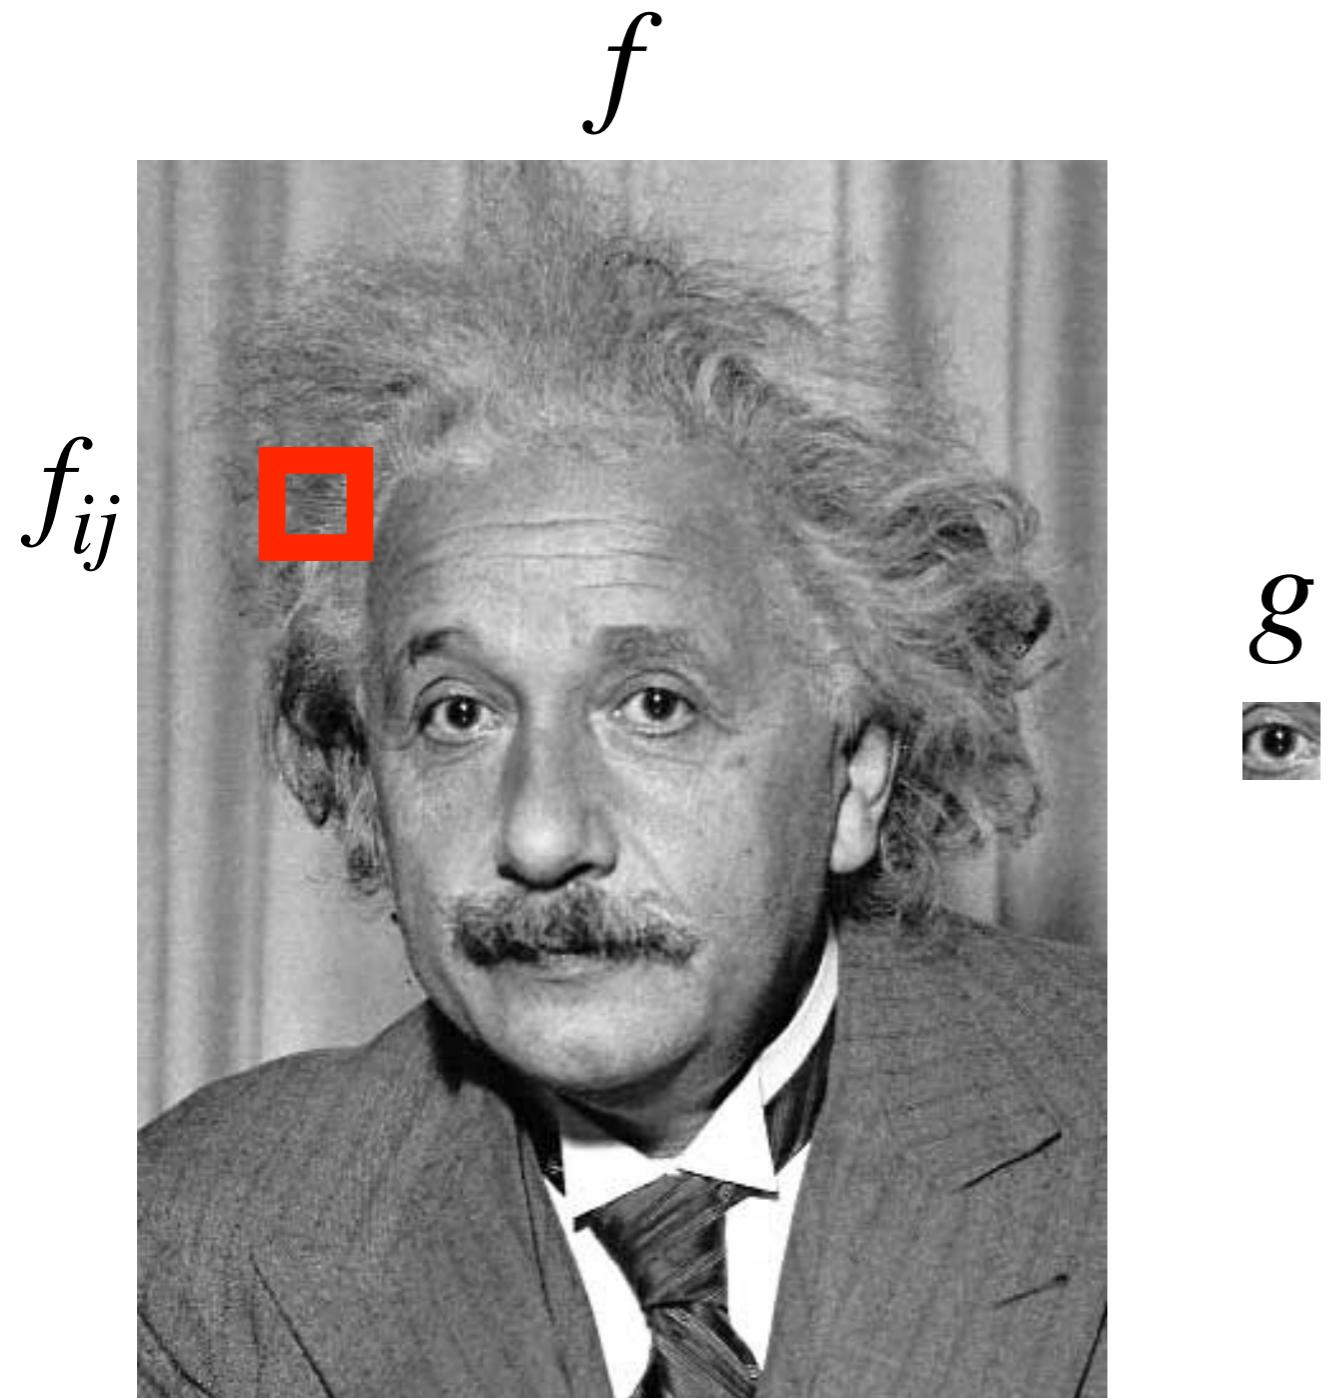
g



$f * g$

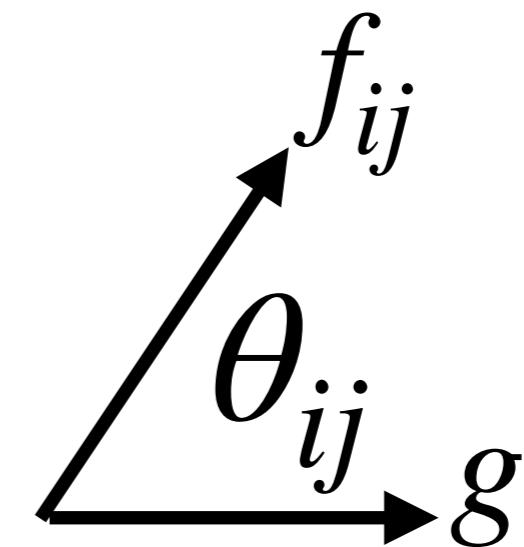


Detection by Filtering



Response for one window:

$$f_{ij}^T g = \|f_{ij}\| \|g\| \cos \theta_{ij}$$

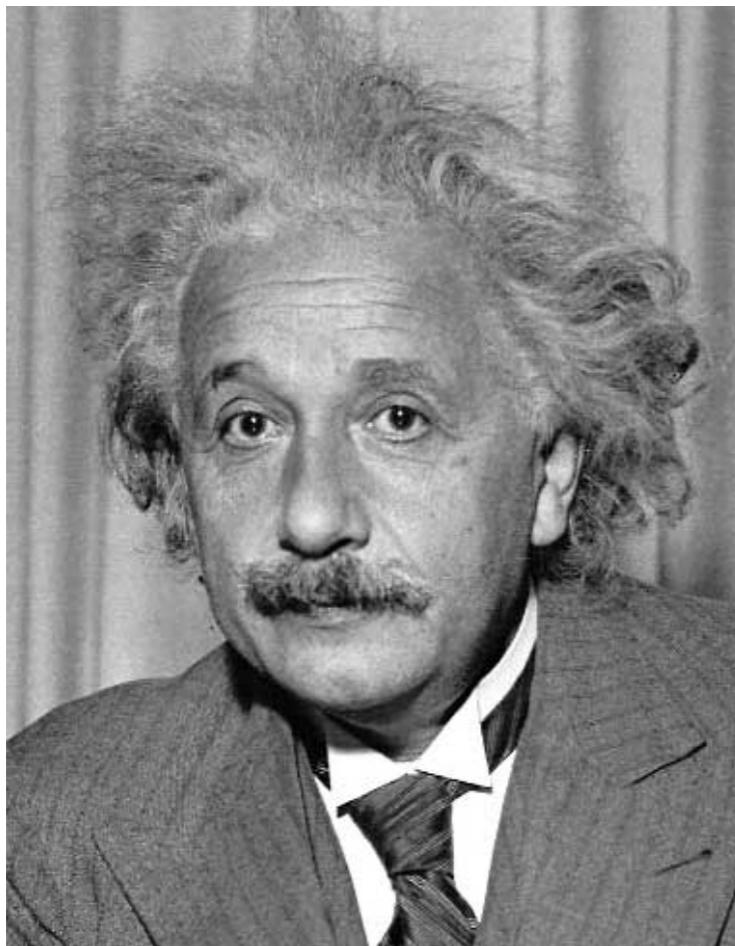


Detection by Filtering

Find the filter g



f

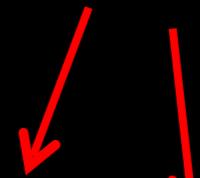


$f * (g - \bar{g})$

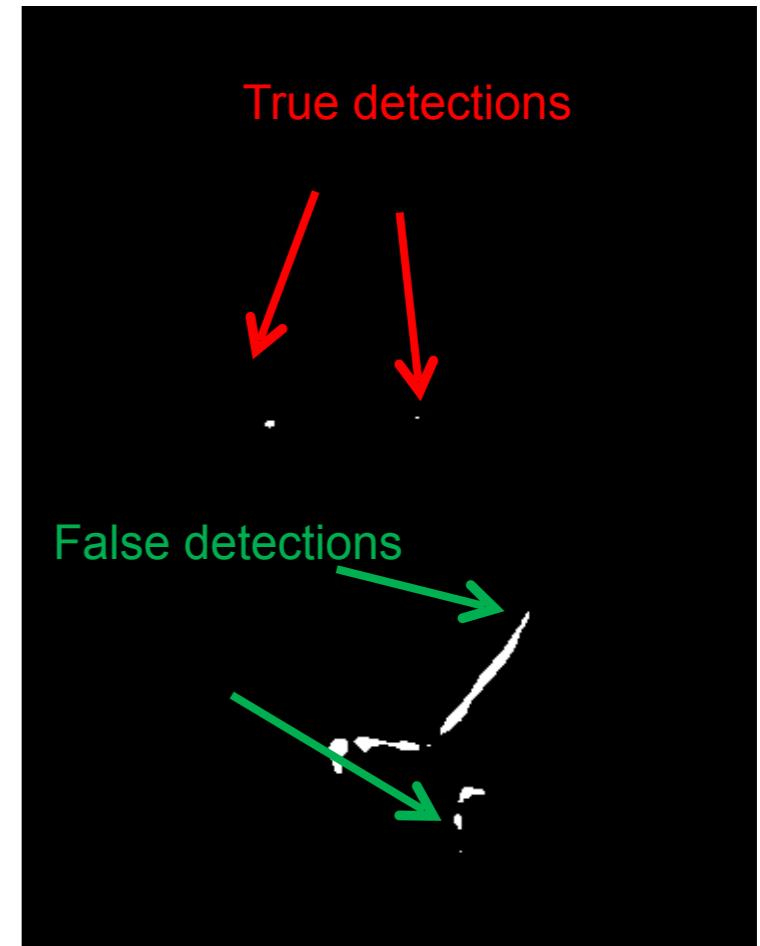
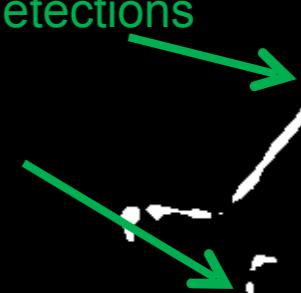


Filter Response

True detections



False detections

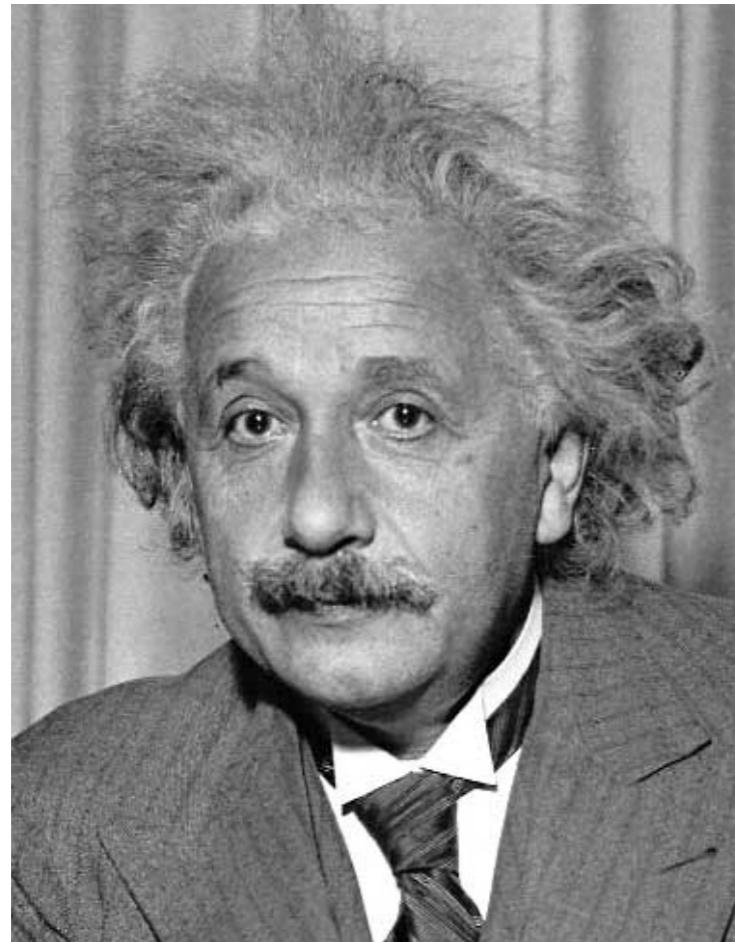


Thresholded

Sum of Squared Differences

$$\begin{aligned} SSD[i, j] &= \|f_{ij} - h\|_2^2 \\ &= (f_{ij} - h)^T (f_{ij} - h) \end{aligned}$$

How do you write
this as a linear filter?



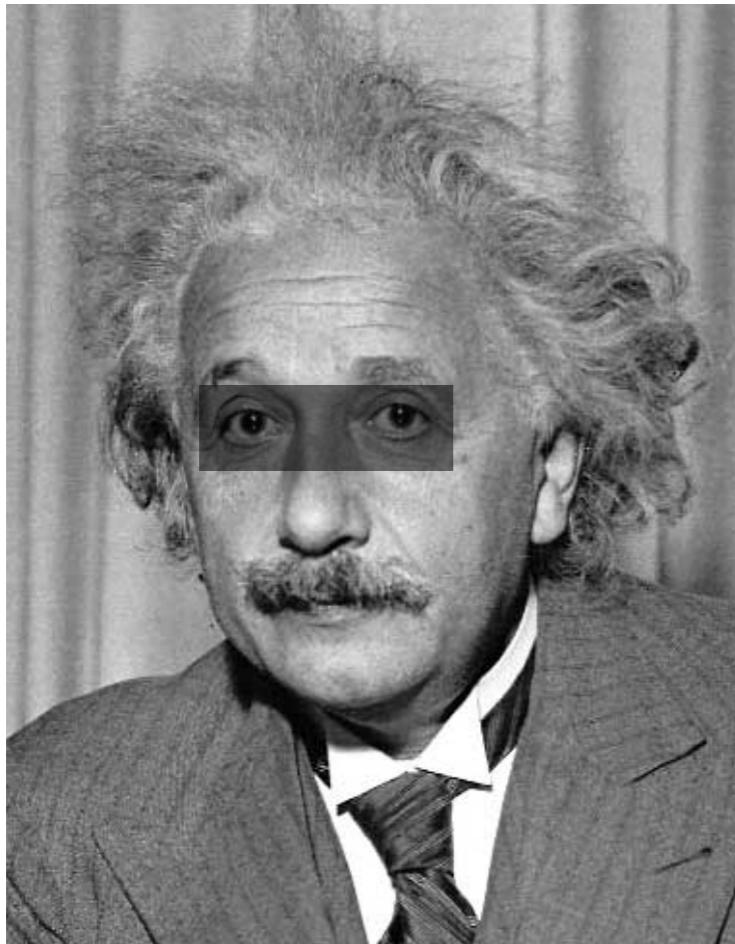
1-sqrt(SSD)



Thresholded

Sum of Squared Differences

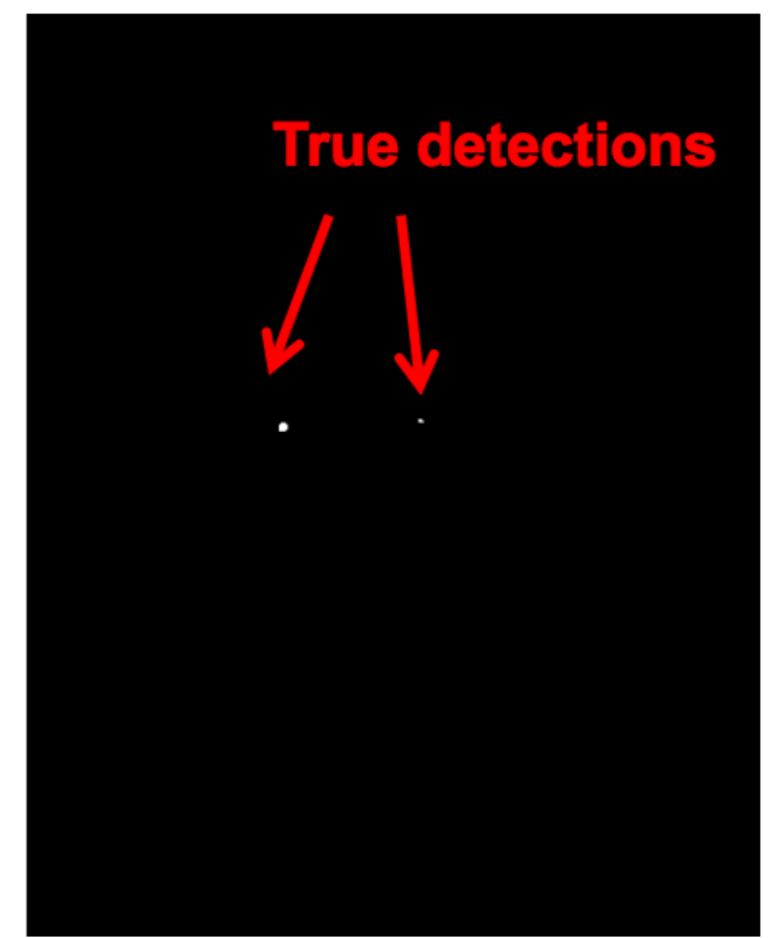
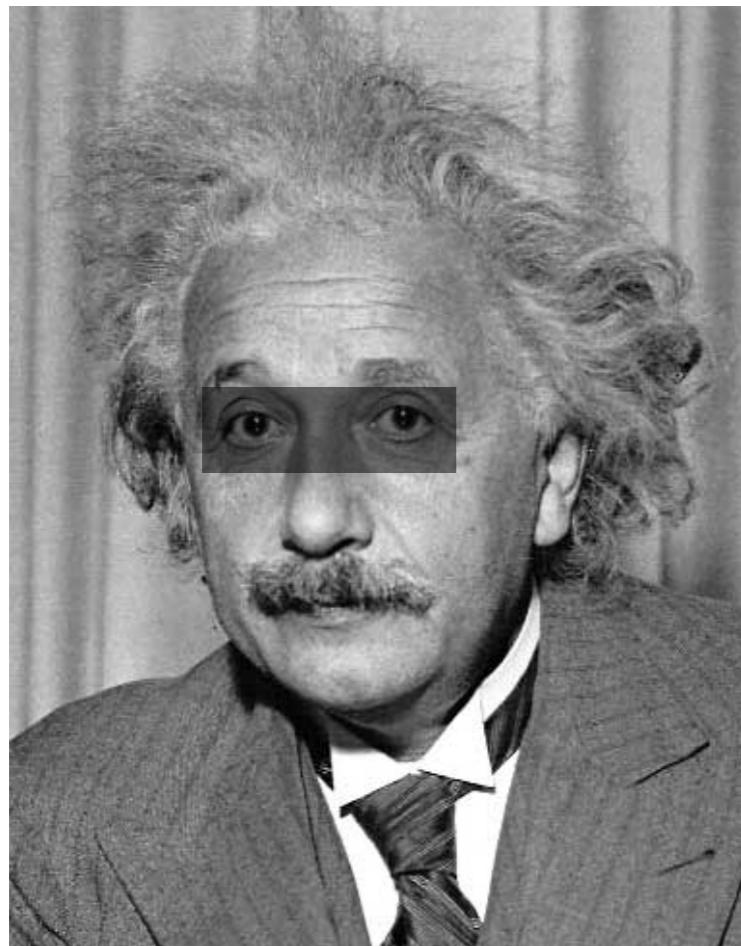
What does SSD do here?



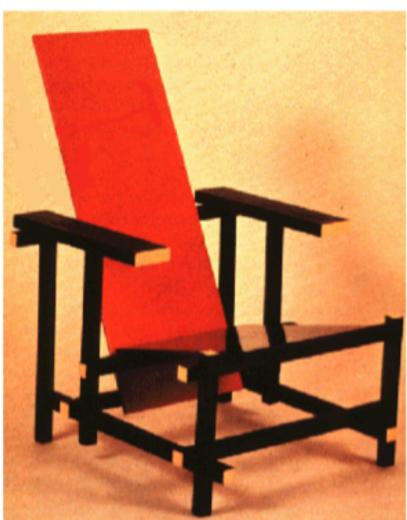
$1 - \text{sqrt(SSD)}$

Normalized Cross Correlation

$$NCC[i, j] = \frac{f_{ij}^T h}{\|f_{ij}\| \|h\|} = \cos \theta_{ij}$$

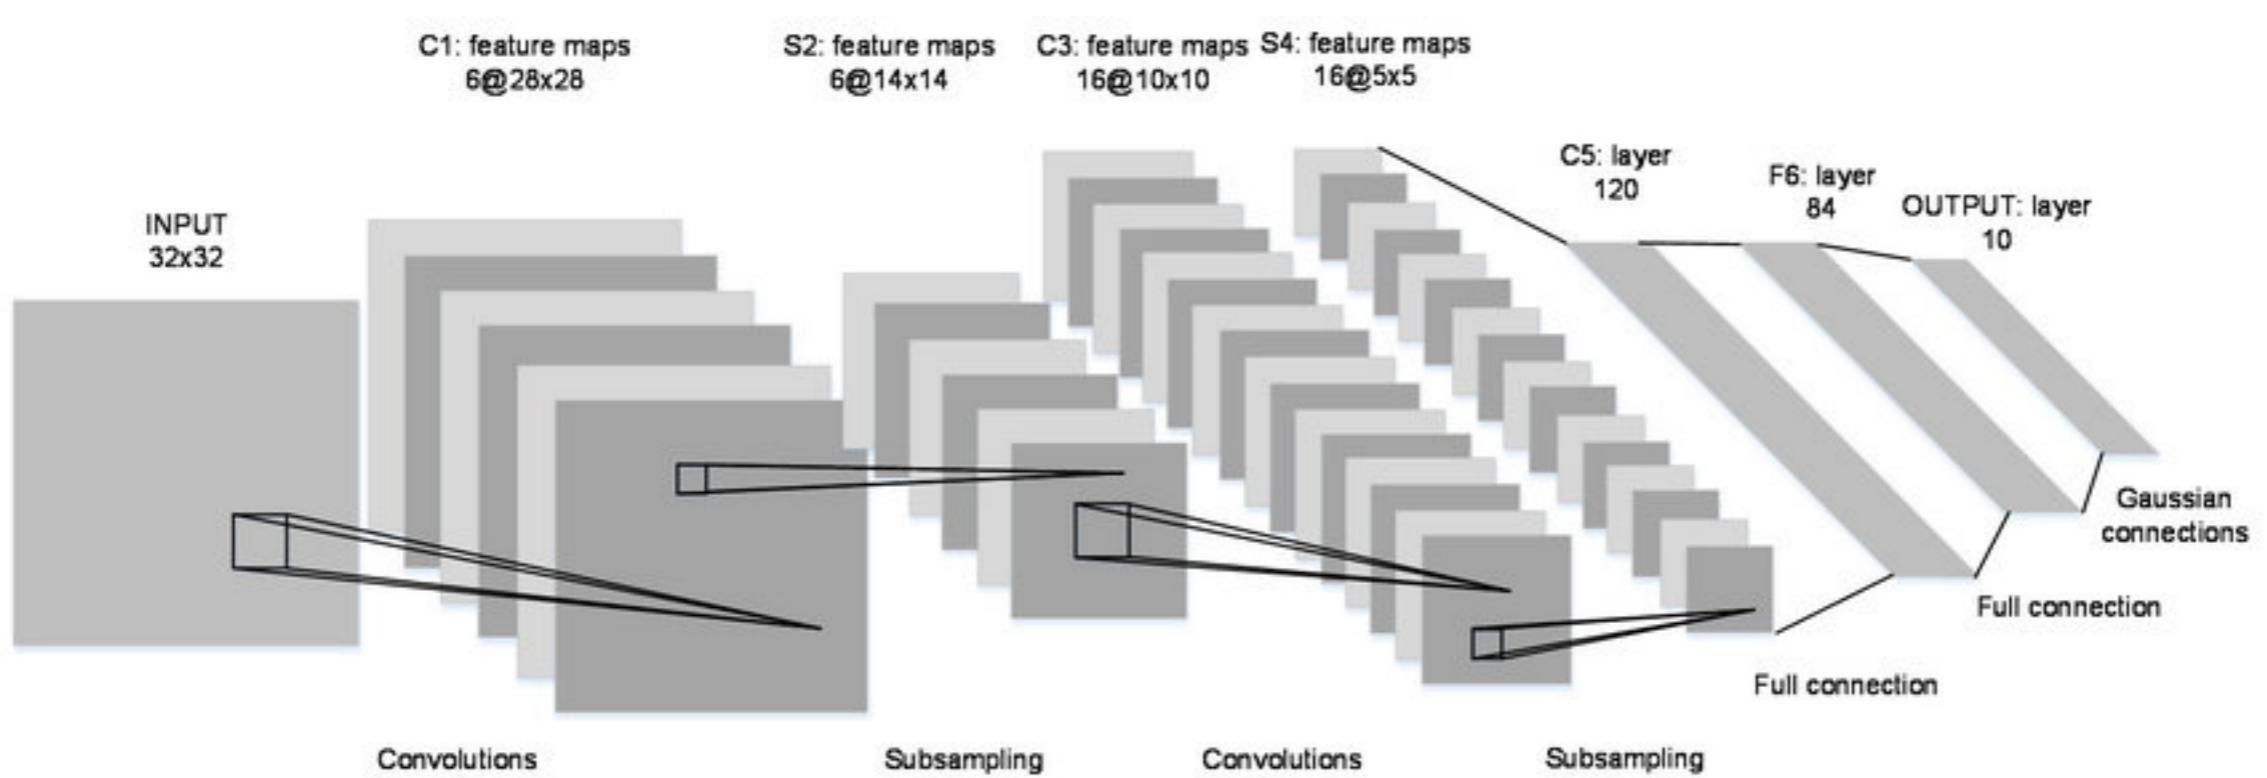


Intra-class variance



Convolutional Networks

Convolution is building blocks for modern object recognition systems



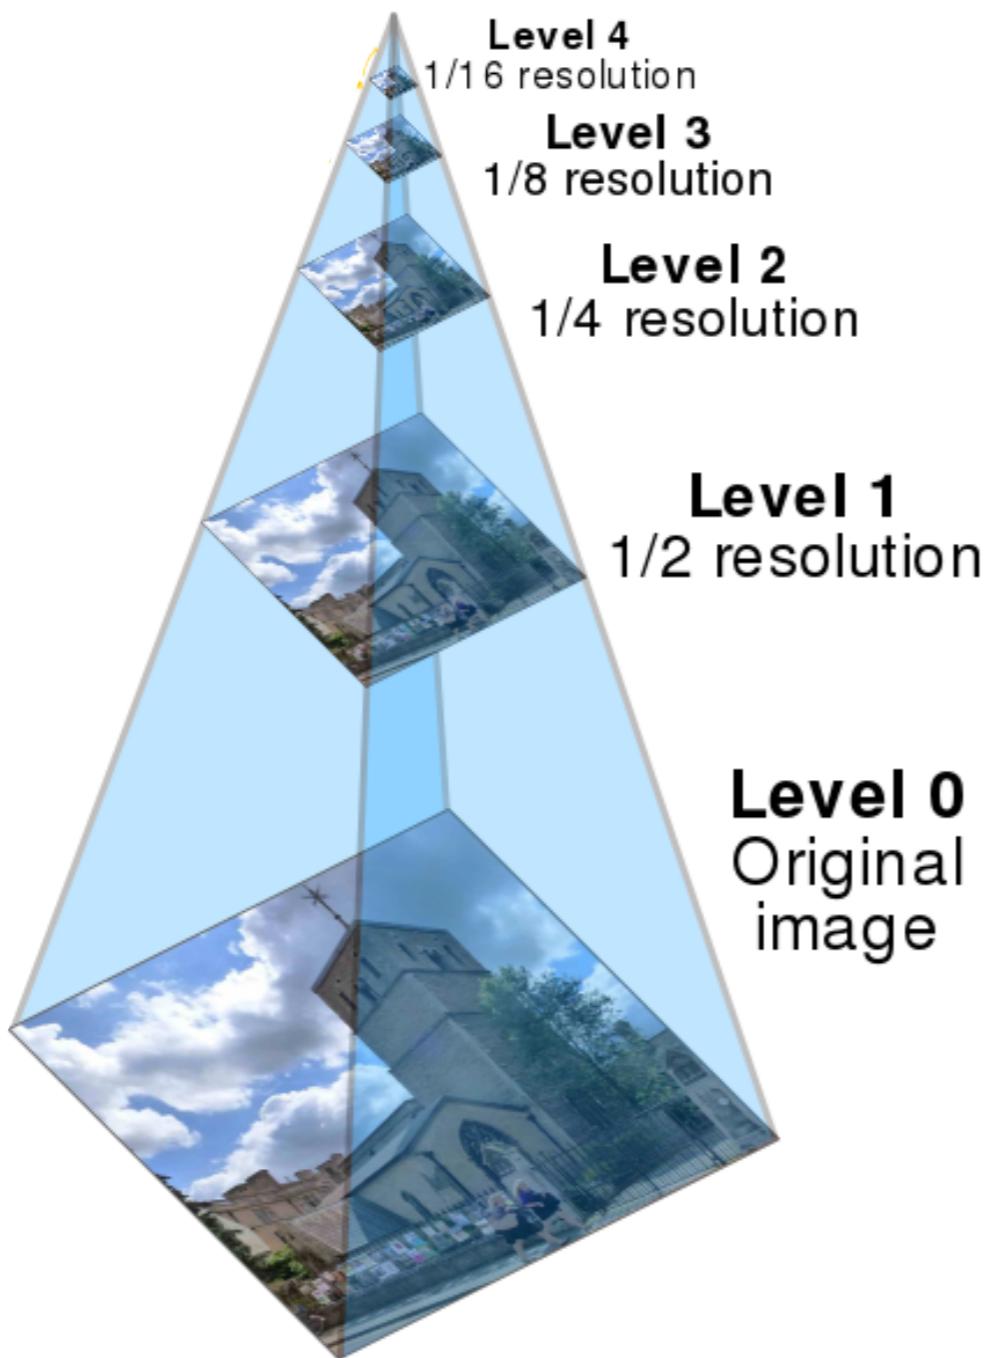
LeNet5

Pyramids



Scale

Image Pyramids



- Recursively resize image by a factor of two
- Called pyramid because it looks like a pyramid
- Invariance to scale by running operation over each level of the pyramid

How to resize images?

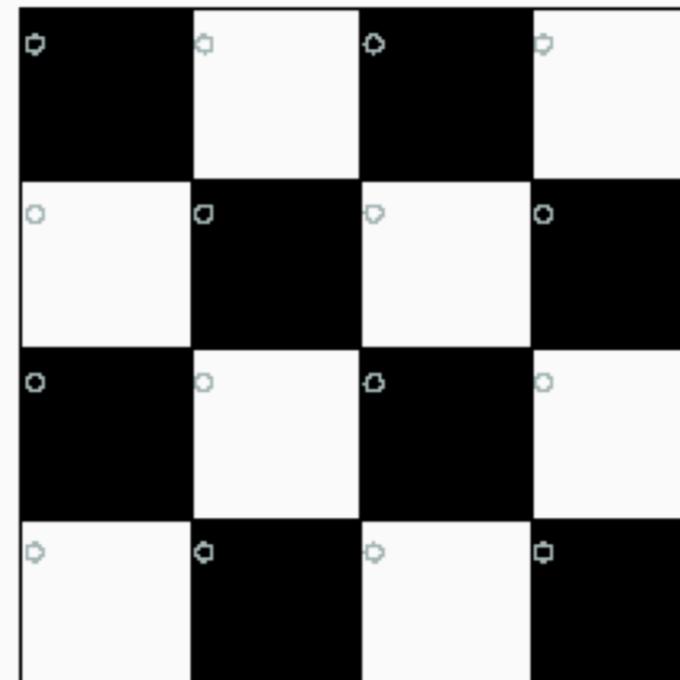
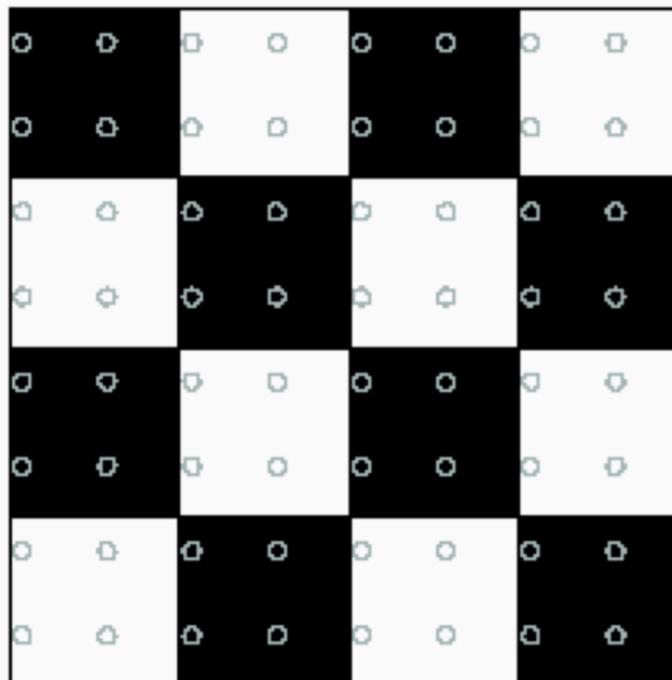


Skip every
other pixel



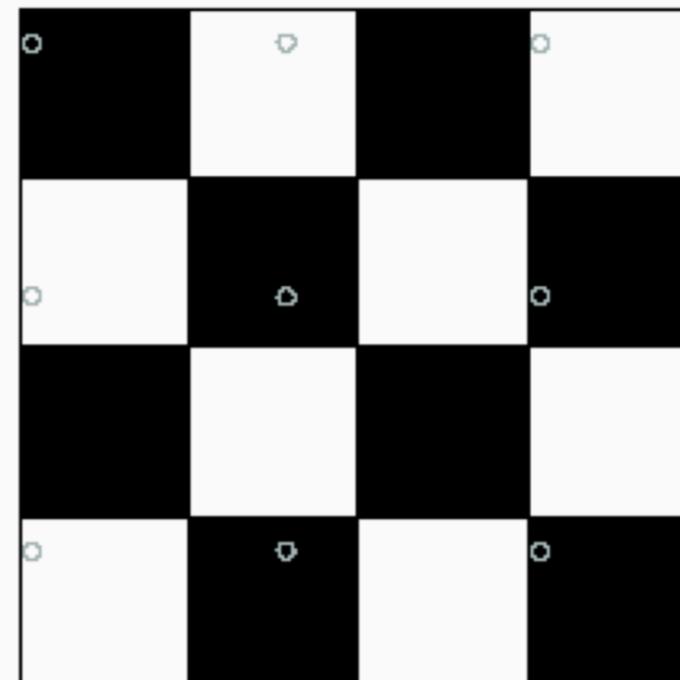
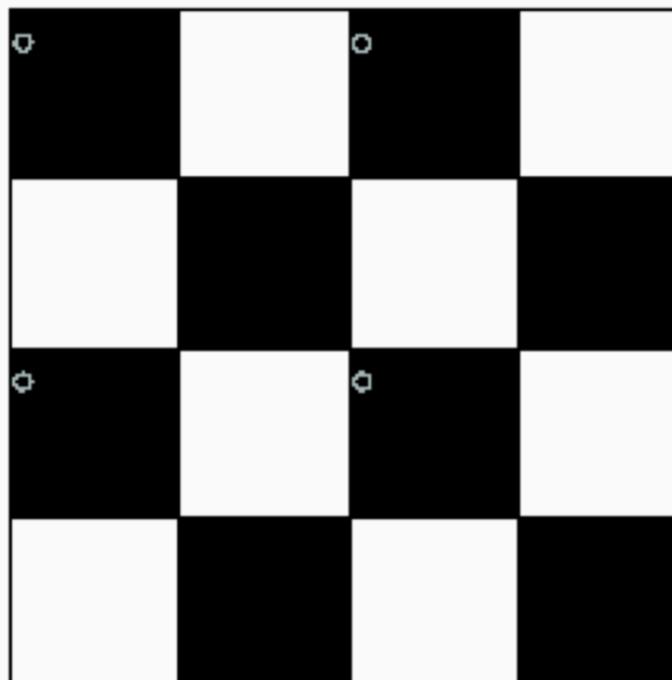
Why does this look bad?

Aliasing



Good sampling:

- Sample often or,
- Sample wisely



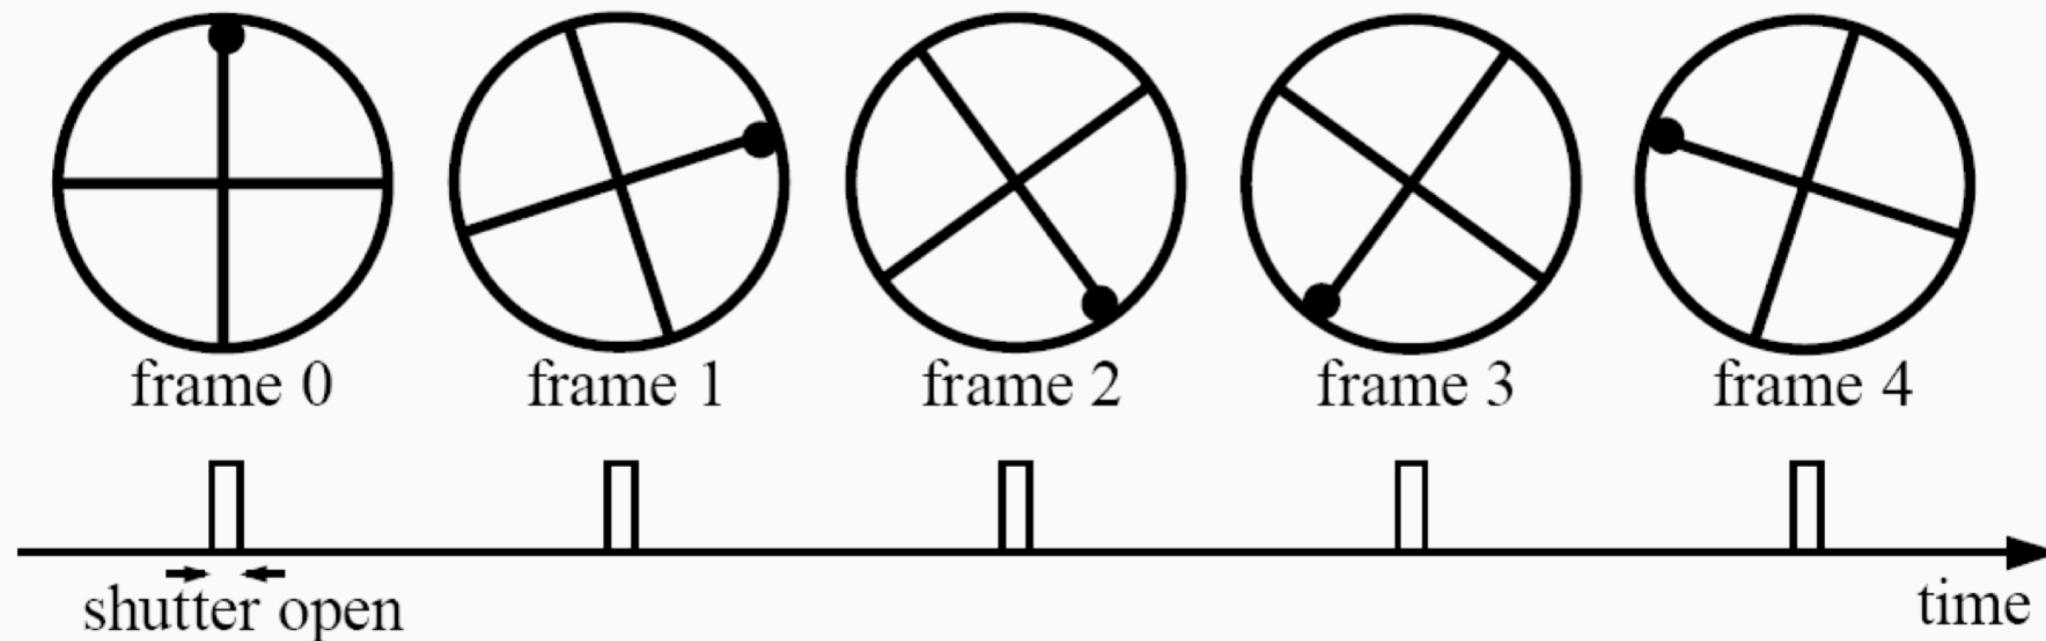
Bad sampling:

- see aliasing in action!

Aliasing

Imagine a spoked wheel moving to the right (rotating clockwise).
Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = 1/30 sec. for video, 1/24 sec. for film):



Without dot, wheel appears to be rotating slowly backwards!
(counterclockwise)



Gaussian Pyramids

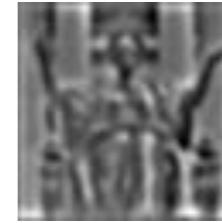
1. Convolve with Gaussian filter
2. Subsample every other pixel
3. Repeat



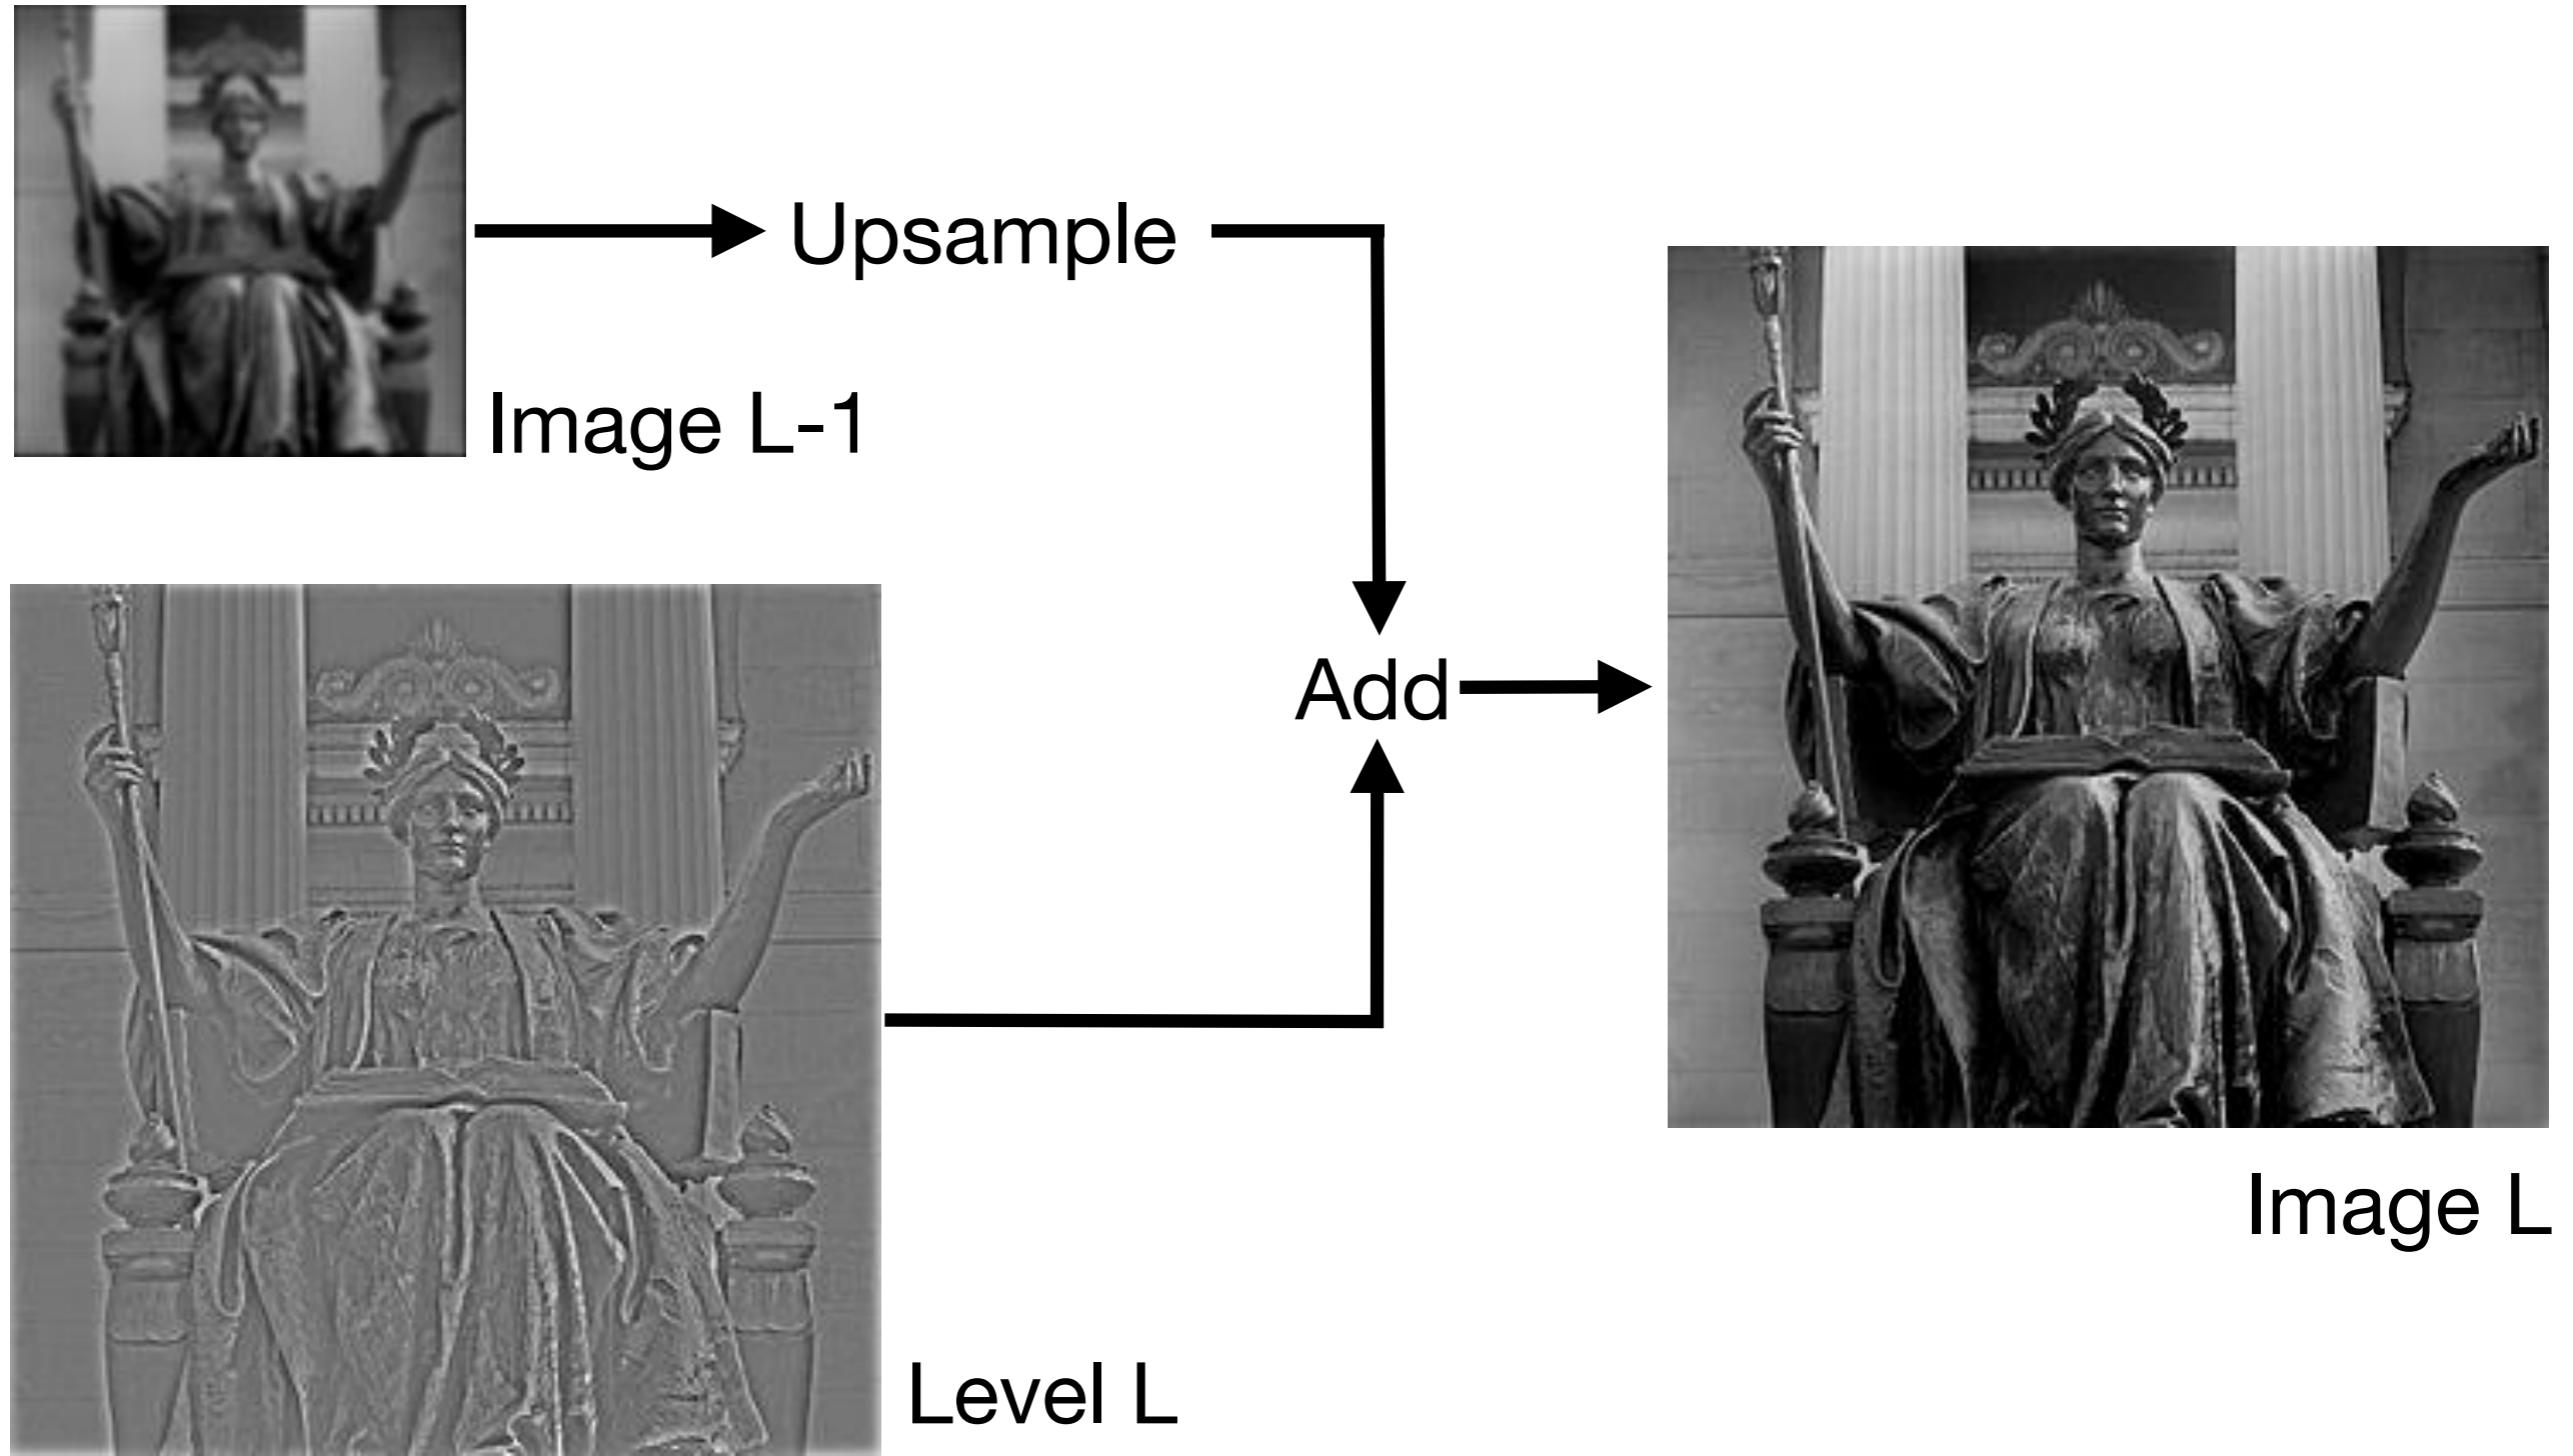
Laplacian Pyramids

1. Convolve with Laplacian filter
2. Subsample every other pixel
3. Repeat

Store downsampled image,
not gradients



Recovering Image



Laplacian Pyramids

Applications:

- Compression
- Incremental transmission

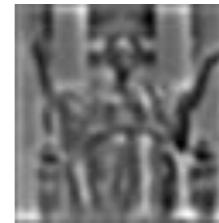
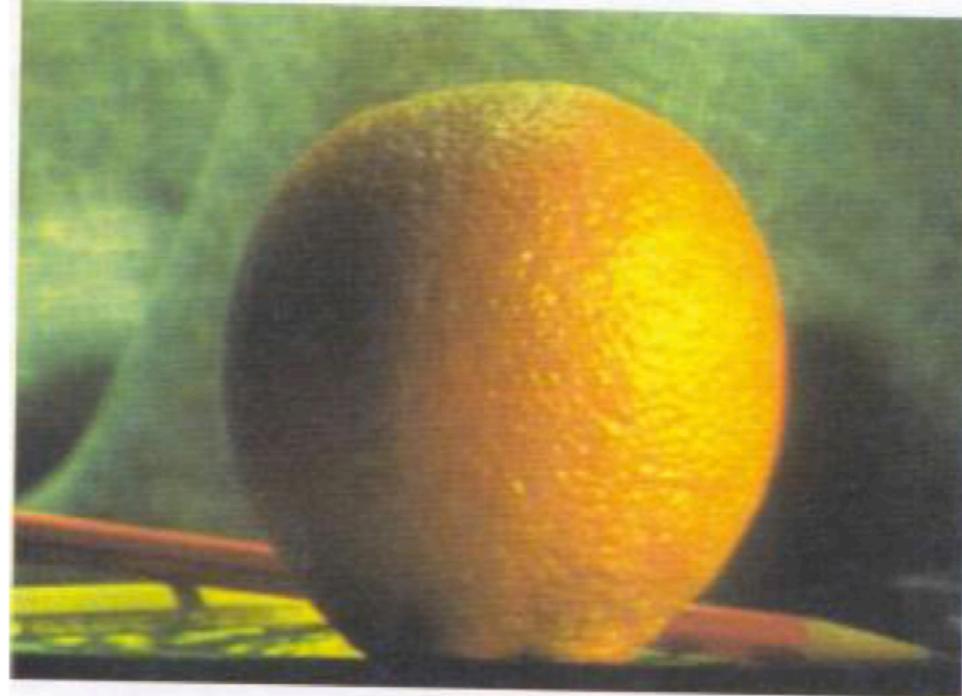
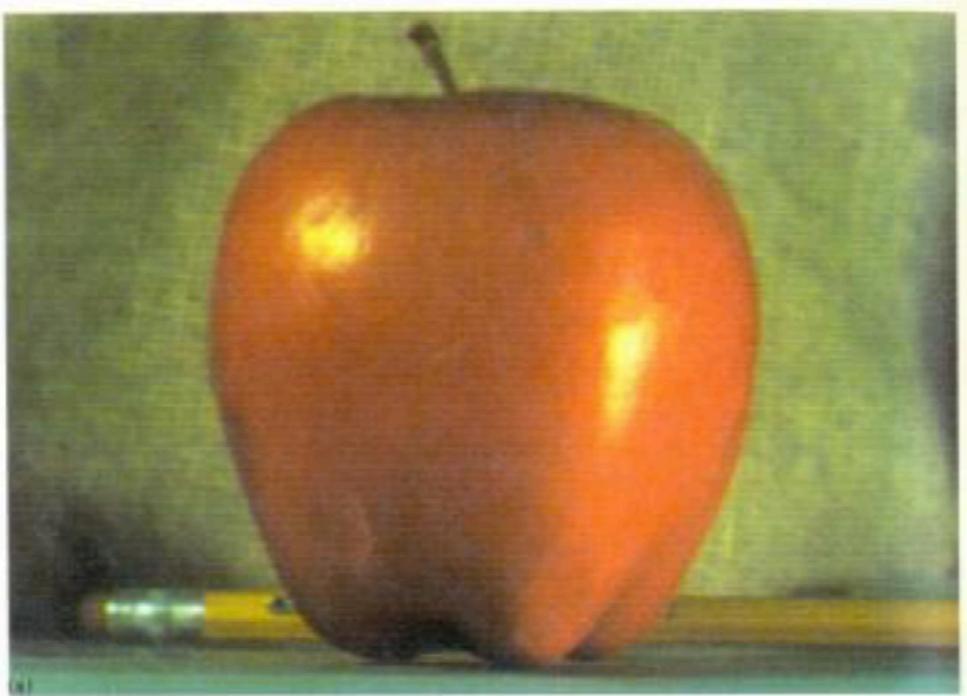


Image Blending



(d)

(h)

(i)

Image Blending

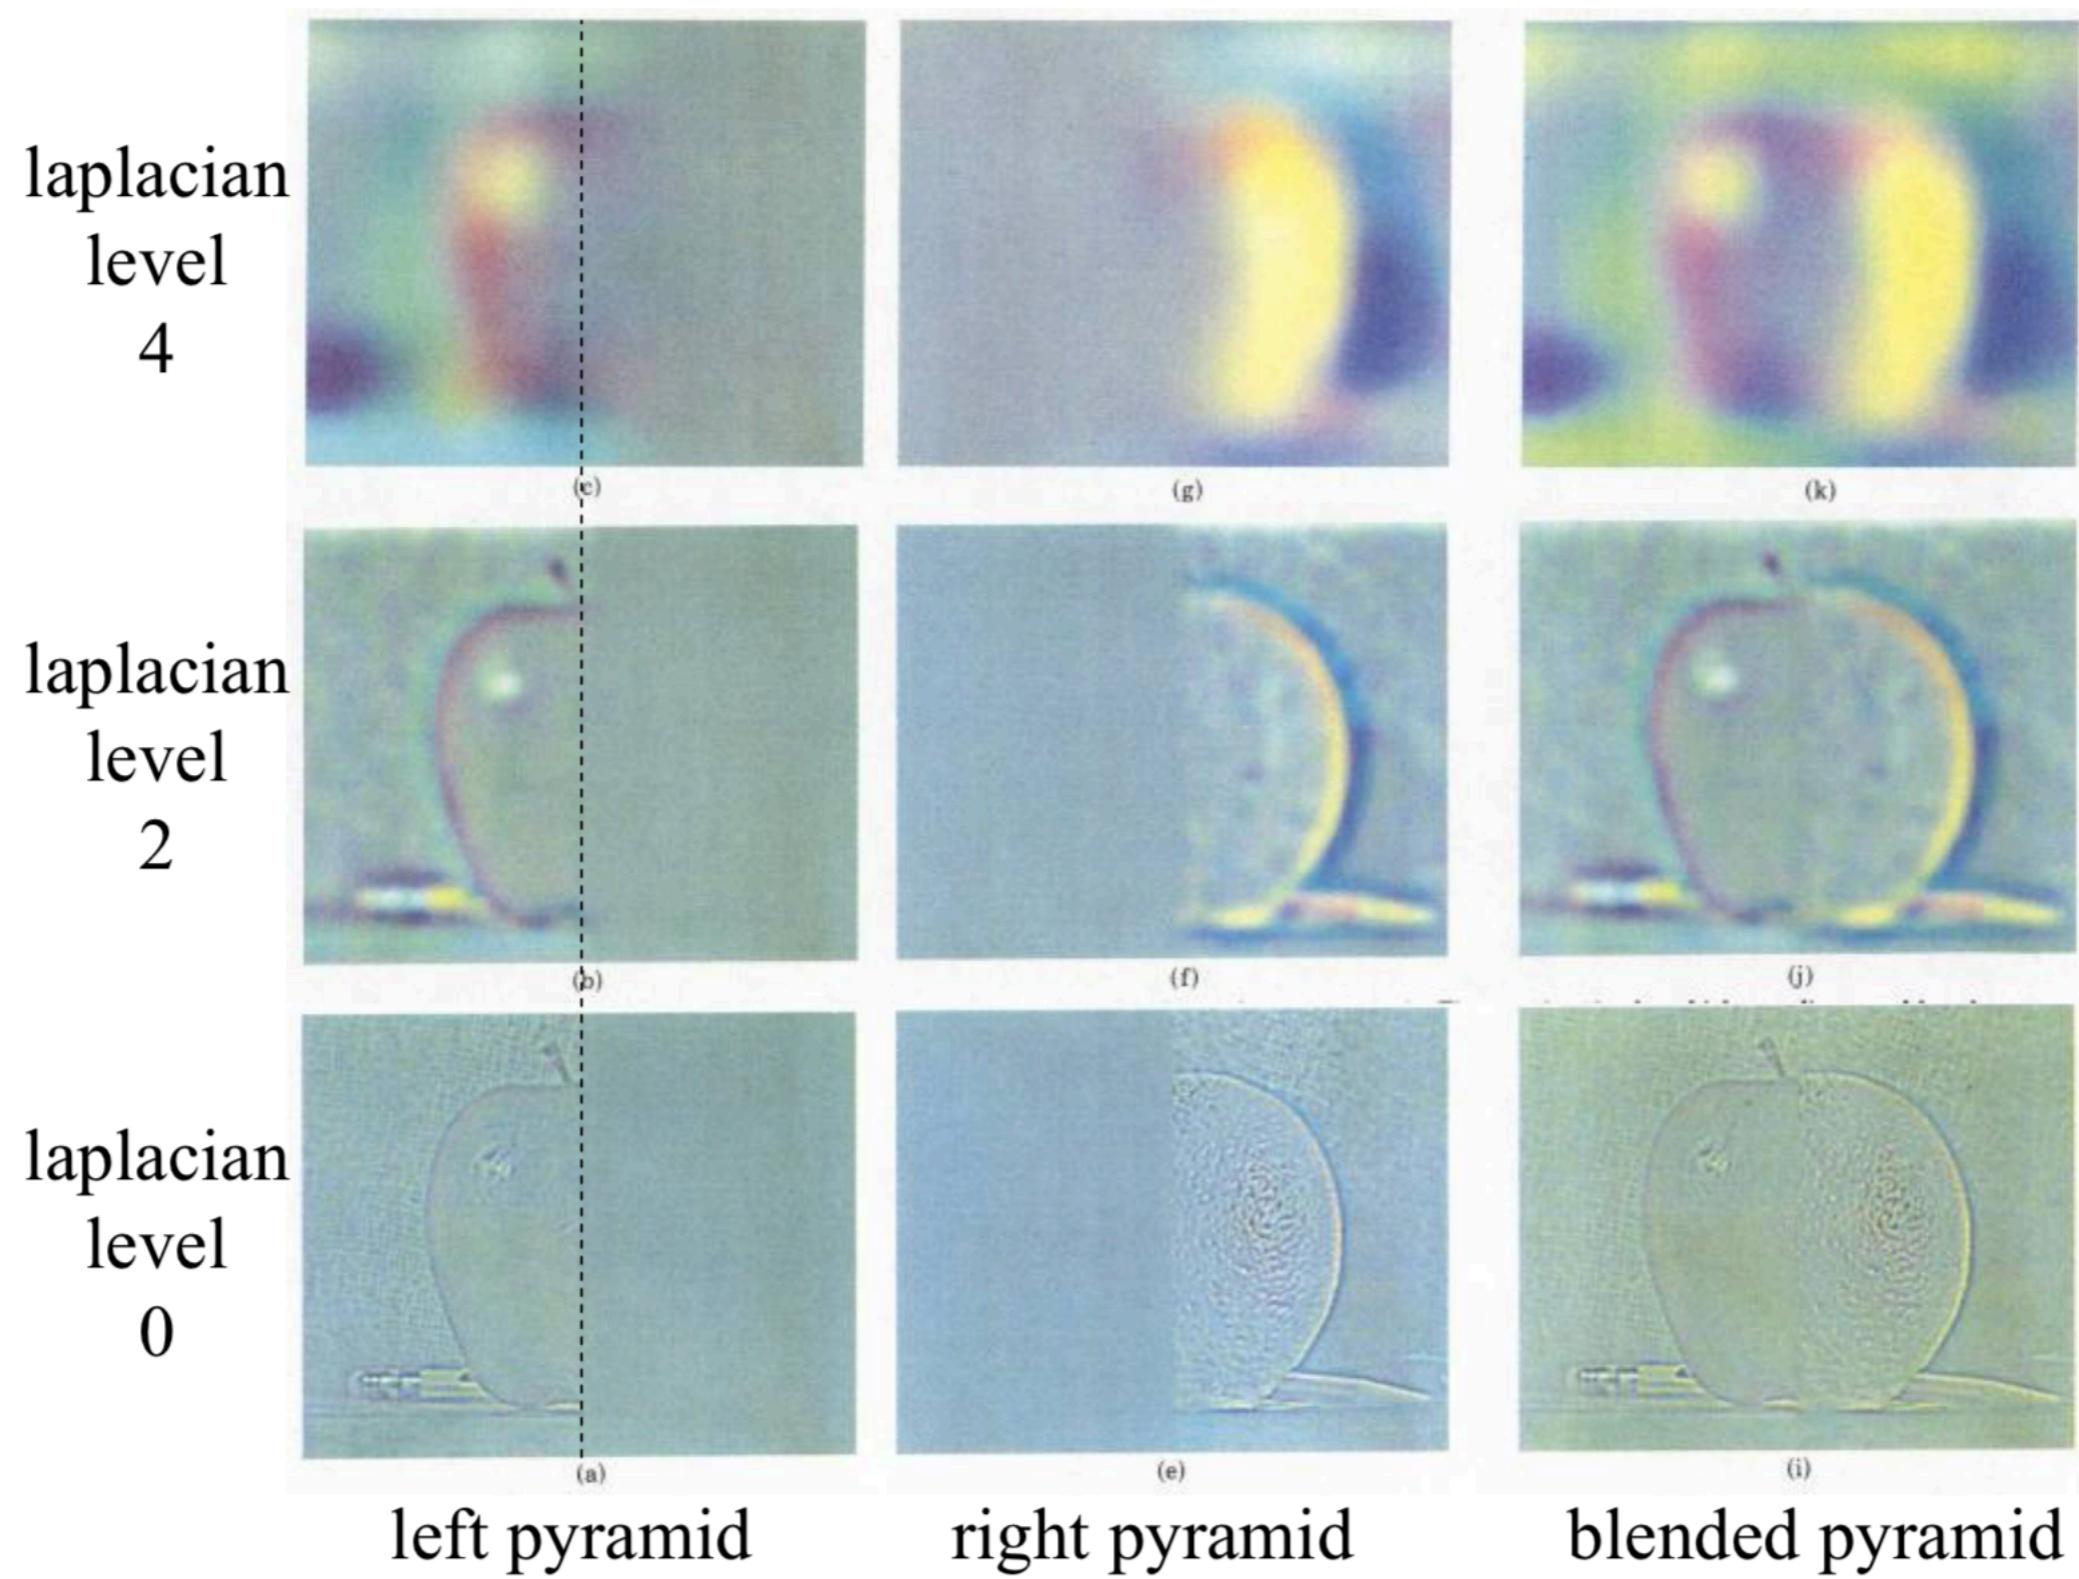


Image Blending

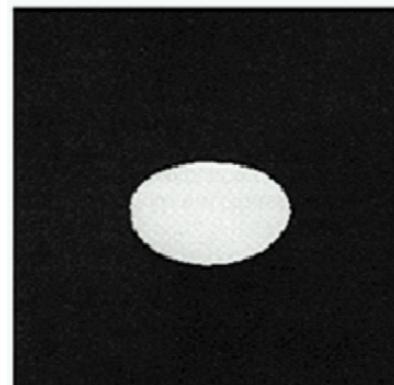
Image A



Image B



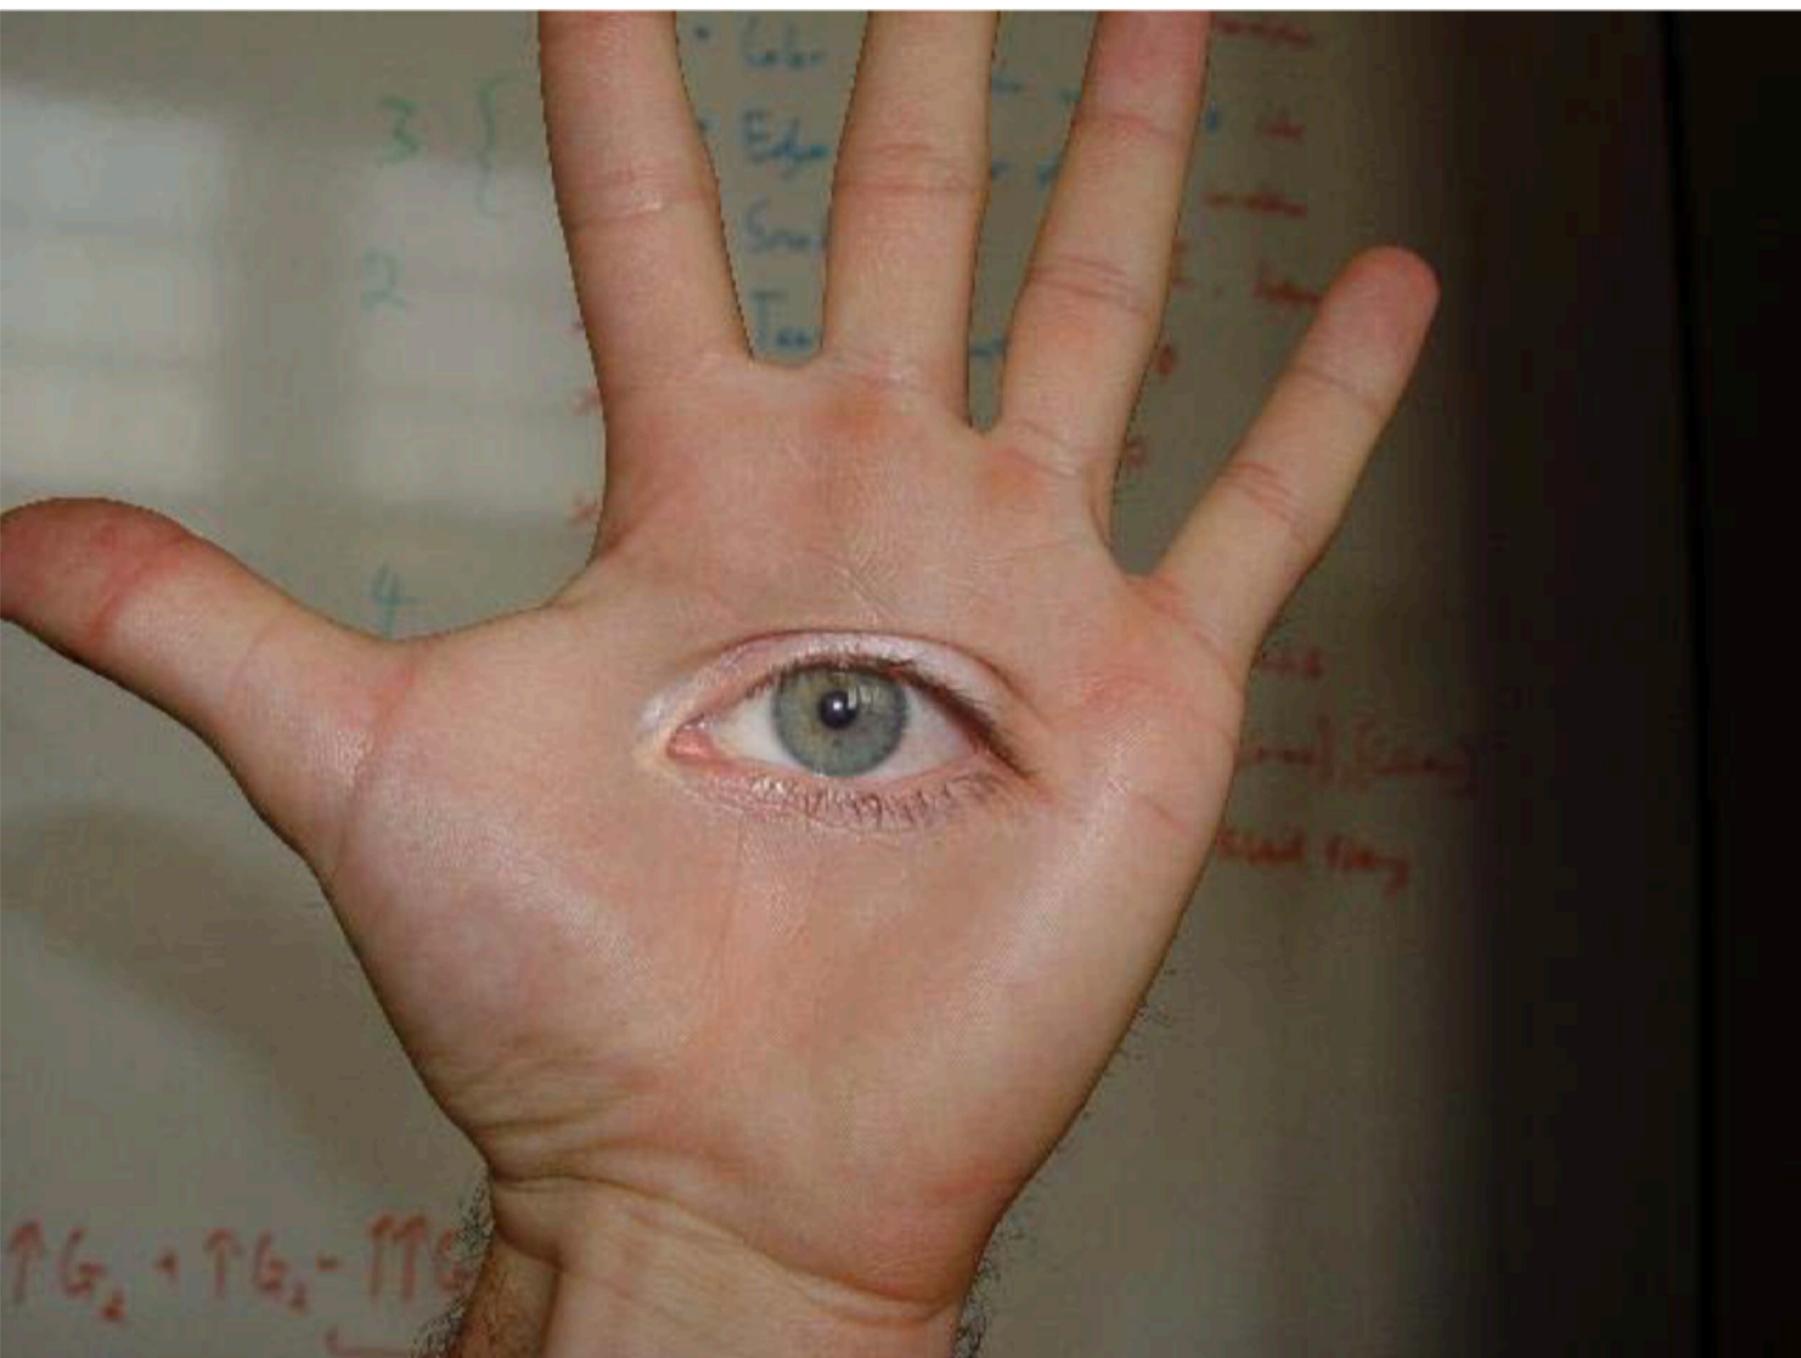
Region R



General Approach:

1. Build Laplacian pyramids LA and LB from images A and B
2. Build a Gaussian pyramid GR from selected region R
3. Form a combined pyramid LS from LA and LB using nodes of GR as weights:
 - $LS(i,j) = GR(i,j) * LA(i,j) + (1-GR(i,j)) * LB(i,j)$
4. Collapse the LS pyramid to get the final blended image

Image Blending



© prof. dmartin

Next Class: Repetition

